



DC – AI – Routing – Switching – Security – Wireless – Automation

JCNR for Equinix Metal

An HPE Juniper Networking Techpost

Vivek Shenoy

2024-01-31

HPE



Contents

1	JCNR for Equinix Metal	4
1.1	Problem Statement	4
1.1.1	Technology Centric Quote	4
1.2	Introduction	4
1.3	High-Level Architecture of an Equinix Metal Node with JCNR as Kubernetes CNI	5
1.4	JCNR Equinix Solution Architecture:	6
1.5	Solution validation	7
1.5.1	Status of Equinix metal servers	7
1.5.2	JCNR verification	7
1.5.3	Validation of end-end connectivity	9
1.6	Conclusion	11
1.7	Useful links	11
1.8	Glossary	12
1.9	Acknowledgements	12

List of Figures

1	JCNR Equinix Architecture	5
2	Solution topology and architecture	6
3	Status	7

List of Tables

1 JCNR for Equinix Metal

Vivek Shenoy - 01/31/2024

JCNR brings a lot of value by providing seamless connectivity between workloads across locations, public cloud boundaries, and workload form-factor, by providing full router functionality.

Author would like to acknowledge and thank Oleg Berzin and Vinod Nair (Juniper Networks) for their help in putting this solution together in such a short time.

1.1 Problem Statement

In this day and age of hybrid cloud, it is very common for applications for a given user or organization to span across different datacenter (on-prem or colo) locations, across multiple public cloud providers. Furthermore, they may exist in the form of bare metal, virtual machine or container form-factors. Often apps require seamless connectivity at both Layer 2 and Layer 3 that cuts across cloud boundaries, which is not easy to achieve. Moreover, the typical VPC networking provided by public cloud environments is use-case specific and poses a lot of restrictions as compared to traditional DC networking constructs. This is where JCNR brings a lot of value by providing seamless connectivity between workloads across locations, public cloud boundaries, and workload form-factor, by providing full router functionality.

1.1.1 Technology Centric Quote

“Juniper Cloud Native Router (JCNR) is a new breed of Cloud-native Network Functions (CNFs) that have been designed to integrate directly and organically with Kubernetes and deliver a rich set of advanced multi-tenant networking features as well as very-high packet processing and forwarding performance using minimal CPU resources. These characteristics make JCNR applicable not only to the Telco/5G use cases but to a wide variety of enterprise scenarios including hybrid multi-cloud networking. With Equinix Metal and Fabric we can demonstrate that customers can deploy and configure performant hybrid multi-cloud networks using JCNR software on Equinix infrastructure to enable multi-metro, multi-cloud private low latency connectivity for their applications in the matter of minutes.”

Oleg Berzin - Senior Distinguished Engineer, Technology and Architecture, Office of the CTO, Equinix.

1.2 Introduction

Equinix Metal is a bare metal infrastructure, which is based on dedicated physical servers, that center on time-to-market, overhead reductions, cloud-like consumption, and integration with other Equinix services in addition to providing high-speed connectivity to public cloud infrastructure using Equinix's high-speed global network infrastructure.

Juniper Cloud-Native Router (JCNR) is a Kubernetes container-based software router that combines the Junos cRPD control and management plane and the vRouter forwarding plane. Telcos are deploying JCNR in their production environments for use-cases such as a cloud-native Cell Site Router (CSR) in the case of 5G Distributed Radio Access Network (D-RAN) and to provide advanced network connectivity and segmentation for 5G Core workloads hosted on public cloud environments.

In addition to the 5G use cases, JCNR is an attractive solution for Enterprises to achieve multi-site, multi-cloud networking for application workloads which this article focuses on.

1.3 High-Level Architecture of an Equinix Metal Node with JCNR as Kubernetes CNI

The diagram below shows the high-level architecture of JCNR running on top of an Equinix Metal server that is orchestrated on the Equinix Metal infrastructure. Equinix Metal infrastructure along with Equinix's global high-speed backbone provides the required fabric connectivity (L3) between the various JCNR nodes to achieve seamless connectivity between the applications across clouds.

As shown in the diagram below the key JCNR components consists of the following:




Figure 1: JCNR Equinix Architecture

- Juniper Containerized Route Processing Daemon or cRPD is the brain of JCNR that runs as a stateful-set on each of the Kubernetes nodes providing advanced networking capabilities for Kubernetes application pods and which talks to the JCNR forwarding plane (vRouter) using GRPC APIs to program configuration data, control-plane state information and to exchange operational state with the vRouter forwarding plane.
- CRPD, just like Junos running on a physical Juniper router, provides support for most of the Layer2/Layer3 networking protocols and features such as OSPF, ISIS, BGP, MPLS L2/L3 VPNs, EVPN-VXLAN, Segment-Routing, ACLs, TWAMP, QOS, etc. In addition, there are a variety of options to manage/automate the JCNR configuration using helm-charts, Netconf-SSH, CLI, etc.
- JCNR forwarding plane (vRouter) is a high-performance, DPDK-based forwarding plane developed by Juniper Networks that has direct access to the physical and virtual functions (VF/PF) on the fabric interfaces for high-performance traffic forwarding. In addition, pods can be attached to the JCNR

through Kubernetes Network Attachment Definition (NAD) construct using vhost (DPDK application-based pod) or veth (kernel mode) interfaces.

- JCNR vRouter forwarding plane also provides a native telemetry interface for fetching rich telemetry and metrics information that can be ingested into any standard 3rd party telemetry importer such as Prometheus.
- To provide seamless integration of JCNR for Kubernetes workloads, Juniper provides its own JCNR CNI which sits below the Multus meta plugin. The job of the JCNR CNI is to keep a watch on the Kubernetes configuration changes and translate them to equivalent JCNR configuration constructs as and when required.

Note: For in-depth architecture of JCNR please refer to the Junos documentation and other JCNR blog URLs provided at the end of this article in the references section.

1.4 JCNR Equinix Solution Architecture:

This section explains in detail the solution architecture and the use cases that were tested. At a high level, the JCNR Equinix solution aligns with the Equinix Network Edge architecture, a reference architecture from Equinix for connecting multiple locations and clouds. JCNR is deployed in a DIY fashion and plays the role of vGW in this architecture. We will not cover the details of JCNR installation in this post as a detailed installation guide can be found under JCNR Technical documentation for various OS types and environments.




Figure 2: Solution topology and architecture

Use case and test details:

- In the Dallas region, 2 JCNR instances are running on 2 separate Equinix metal nodes, and in the Sunnyvale region we have 1 JCNR node. These bare metal servers act as independent Kubernetes clusters for hosting JCNR and other application pods. This is depicted in the diagram as shown above.

- Each of these servers is equipped with a dual-port Intel XXV710 25G NIC with the first 25G port connected to Equinix global backbone TOR switches for high-speed inter-POP connectivity. The second port is used for intra-POP connectivity in the Dallas region and in the Sunnyvale region connection towards public cloud AWS-VPC/Azure-VNET.
- As indicated in the diagram above, the first 25G port on each of the servers is configured as an untagged port with a single SRIOV-VF allocated to JCNR, whereas multiple SRIOV-VFs are configured on the second 25G port, and are connected to bare metal servers and towards public cloud VPC/VNET interconnects (using different VLANs).
- In addition, each metal node has an application pod/CNF attached to JCNR using a Kubernetes native construct known as Network Attachment Definition (NAD).
- All the above-mentioned workloads have their JCNR interfaces part of Layer3 VRF vpic-cust1 which is configured on the three JCNR instances. The inter-metro communication between workload endpoints is achieved using the BGP-based L3VPN MPLSoUDP overlay tunnels that are constructed between the three JCNR nodes.
- One use case not mentioned in the above diagram is internet connectivity for the above workloads. But that can be achieved very easily by service chaining JCNR with cSRX (which provides Firewall and NAT functions) and breaking out locally on the local metal server for internet access. That's a topic that deserves its own blog!

1.5 Solution validation

Note: For the sake of brevity, JCNR Junos configuration and the cli capture from the Linux host and JCNR are provided in the blog repo whose link is provided at the end. Only key verification commands are covered in this document in the next section.

1.5.1 Status of Equinix metal servers

	Hostname	Server	IPV4 Address	Type	Status	OS	Metro	Tags
LAN Servers								
☐	c3-small-x86-xrd-1-L...	c3.small.x86	145.40.102.59	On Demand	Running	CentOS	DA	>_ →
☐	c3-small-x86-xrd-1-L...	c3.small.x86	147.28.155.9	On Demand	Running	CentOS	DA	>_ →
☐	c3-small-x86-xrd-3-L...	c3.small.x86	139.178.89.233	On Demand	Running	CentOS	SV	>_ →
☐	forte-edge-cloud-1	c3.small.x86	139.178.70.181	On Demand	Running	CentOS	SV	>_ →
SONiC-VPP Nodes								
☐	octo-test-sonic-vpp-...	c3.medium.x86	147.75.55.93	On Demand	Running	Ubuntu	DA	>_ →
☐	octo-test-sonic-vpp-...	c3.medium.x86	147.28.183.221	On Demand	Running	Ubuntu	DA	>_ →
☐	octo-test-sonic-vpp-...	c3.medium.x86	139.178.91.1	On Demand	Running	Ubuntu	SV	>_ →

Figure 3: Status

1.5.2 JCNR verification

Note: The JCNR verification o/p is shown on jcnr-3. For other nodes, the o/p will be similar.

```

1 [root@octo-test-jcnr-3 ~]# kubectl get pod -A --field-selector metadata.namespace!=kube-system
2 NAMESPACE NAME READY STATUS RESTARTS AGE
3 contrail-deploy contrail-k8s-deployer-584f55bdc-fcvb4 1/1 Running 0 14d
4 contrail contrail-vrouter-masters-dc6n6 3/3 Running 0 14d
5 jcnr kube-crpd-worker-sts-0 1/1 Running 0 14d
6 jcnr syslog-ng-f55rn 1/1 Running 0 14d
7 kube-flannel kube-flannel-ds-bfk75 1/1 Running 1 (14d ago) 14d
8 vpic-cust1 vpic-kernel-pod-3 1/1 Running 0 14d
9
10 [root@octo-test-jcnr-3 ~]#
11 [root@octo-test-jcnr-3 ~]# kubectl exec -n jcnr -it kube-crpd-worker-sts-0 -- cli
12 Defaulted container "kube-crpd-worker" out of: kube-crpd-worker, jcnr-crpd-config (init),
    install-cni (init)
13 root@octo-test-jcnr-3> show version
14 Hostname: octo-test-jcnr-3
15 Model: cRPD
16 Junos: 23.2R1.14
17 cRPD package version : 23.2R1.14 built by builder on 2023-06-22 13:51:13 UTC
18
19 root@octo-test-jcnr-3> show interfaces terse | grep "(cni)|(flannel)|(eth)|(enp)|(lo)"
20 cni0 UP 10.244.0.1/24 fe80::cc5:e0ff:fe73:5e9/64
21 enp65s0f0 UP
22 enp65s0f0v0 UNKNOWN 10.67.121.135/31 fe80::9034:865b:5119:e7e1/64 fe80::bc58:efff:fed2:c549/64
23 enp65s0f0v1 UNKNOWN 139.178.88.65/31 fe80::5f16:bda5:e915:6e22/64 fe80::d0f5:25ff:fed2:6f9c/64
24 enp65s0f0v2 UP fe80::fc0a:40cc:a3f8:2b8/64
25 enp65s0f0v3 UP fe80::58e8:bd44:5758:7da7/64
26 enp65s0f1 UP fe80::42a6:b7ff:fe70:dae1/64
27 enp65s0f1v0 UNKNOWN fe80::938e:77c9:78cc:9ef9/64 fe80::f023:11ff:fe1b:3057/64
28 enp65s0f1v1 UNKNOWN 192.168.33.11/24 fe80::2010:5cff:fe28:9f2d/64
29 enp65s0f1v2 UNKNOWN 169.254.12.1/30 fe80::9830:96ff:fe66:a92e/64
30 enp65s0f1v3 UNKNOWN 169.254.12.5/30 fe80::6c01:97ff:fed6:5ba3/64
31 eth0 UNKNOWN 169.254.152.42/32 fe80::3cda:1aff:fee5:7d7d/64
32 flannel.1 UNKNOWN 10.244.0.0/32 fe80::e873:eeff:fe74:fd90/64
33 lo UNKNOWN 127.0.0.1/8 1.2.3.4/32 ::1/128
34 lo0.0 UNKNOWN fe80::38a1:81ff:fee8:87ed/64
35 veth370be25d@if4 UP fe80::4816:37ff:fee0:e29c/64
36 vethc9ef06f8@if4 UP fe80::8086:d7ff:fe94:ed0f/64
37 vethde65034f@if8 UP fe80::d475:1fff:fee7:c837/64
38
39 root@octo-test-jcnr-3> show bgp summary
40 Threading mode: BGP I/O
41 Default eBGP mode: advertise - accept, receive - accept
42 Groups: 4 Peers: 6 Down peers: 0
43 Table Tot Paths Act Paths Suppressed History Damp State Pending
44 bgp.l3vpn.0
45
46      7 7 0 0 0 0
47
48 inet.0
49      2 1 0 0 0 0
50 Peer AS InPkt OutPkt OutQ Flaps Last Up/Dwn State|#Active/Received/Accepted/Damped...
51 10.70.35.129 65000 461416 456029 0 0 20w4d 17:44:55 Establ
52   bgp.l3vpn.0: 3/3/3/0
53   vpic-cust1.inet.0: 3/3/3/0
54 10.70.35.135 65000 192350 190108 0 8 8w4d 8:05:34 Establ
55   bgp.l3vpn.0: 4/4/4/0
56   vpic-cust1.inet.0: 4/4/4/0
57 169.254.12.2 65029 73371 78541 0 1 3w3d 22:12:15 Establ
58   vpic-cust1.inet.0: 5/5/5/0
59 169.254.12.6 12076 423072 456046 0 0 20w4d 17:44:45 Establ
60   vpic-cust1.inet.0: 0/0/0/0
61 169.254.255.1 65530 489511 456023 0 0 20w4d 17:44:51 Establ
62   inet.0: 1/1/1/0
63 169.254.255.2 65530 489575 456022 0 0 20w4d 17:44:47 Establ
64   inet.0: 0/1/1/0
65
66 root@octo-test-jcnr-3>
67
68 root@octo-test-jcnr-3> show dynamic-tunnels database terse
69 *- Signal Tunnels #- PFE-down
70 Table: inet.3
71
72 Destination-network: 10.70.35.129/32
73 Destination Source Next-hop Type Status
74 10.70.35.129/32 10.67.121.135 0x563bc26e3f9c nhid 0 UDP Up
75
76 Destination-network: 10.70.35.135/32
77 Destination Source Next-hop Type Status
78 10.70.35.135/32 10.67.121.135 0x563bc4d1661c nhid 0 UDP Up

```

```

77 10.70.35.135/32 10.67.121.135 0x563bc4d16cdc nhid 0 UDP Up
78 10.70.35.135/32 10.67.121.135 0x563bc4d16a9c nhid 0 UDP Up
79
80 root@octo-test-jcnr-3>
81
82 root@octo-test-jcnr-3> show route table vpic-cust1.inet.0
83
84 vpic-cust1.inet.0: 20 destinations, 20 routes (20 active, 0 holddown, 0 hidden)
85 + = Active Route, - = Last Active, * = Both
86
87 169.254.12.0/30 *[Direct/0] 20w6d 01:48:24
88 > via enp65s0flv2
89 169.254.12.1/32 *[Local/0] 20w6d 01:48:24
90 Local via enp65s0flv2
91 169.254.12.4/30 *[Direct/0] 20w6d 01:48:24
92 > via enp65s0flv3
93 169.254.12.5/32 *[Local/0] 20w6d 01:48:24
94 Local via enp65s0flv3
95 172.29.62.0/24 *[BGP/170] 2w0d 11:52:49, localpref 100
96 AS path: 65029 65029 65021 ?, validation-state: unverified
97 > to 169.254.12.2 via enp65s0flv2
98 172.29.63.0/24 *[BGP/170] 2w0d 11:52:49, localpref 100
99 AS path: 65029 65029 65021 ?, validation-state: unverified
100 > to 169.254.12.2 via enp65s0flv2
101 172.29.64.0/24 *[BGP/170] 2w0d 11:52:49, localpref 100
102 AS path: 65029 65029 65021 ?, validation-state: unverified
103 > to 169.254.12.2 via enp65s0flv2
104 172.29.66.0/24 *[BGP/170] 2w0d 11:52:49, localpref 100
105 AS path: 65029 65029 65021 ?, validation-state: unverified
106 > to 169.254.12.2 via enp65s0flv2
107 173.29.0.0/16 *[BGP/170] 3w5d 06:14:48, localpref 100
108 AS path: 65029 I, validation-state: unverified
109 > to 169.254.12.2 via enp65s0flv2
110 192.168.33.0/24 *[Direct/0] 20w6d 01:48:24
111 > via enp65s0flv1
112 192.168.33.11/32 *[Local/0] 20w6d 01:48:24
113 Local via enp65s0flv1
114 192.168.41.0/24 *[BGP/170] 20w6d 01:48:14, localpref 100, from 10.70.35.129
115 AS path: I, validation-state: unverified
116 > via Tunnel Composite, UDP (src 10.67.121.135 dest 10.70.35.129), Push 16
117 192.168.42.0/24 *[BGP/170] 8w5d 16:01:35, localpref 100, from 10.70.35.135
118 AS path: I, validation-state: unverified
119 > via Tunnel Composite, UDP (src 10.67.121.135 dest 10.70.35.135), Push 19
120 192.168.42.21/32 *[BGP/170] 8w5d 16:01:35, localpref 100, from 10.70.35.135
121 AS path: I, validation-state: unverified
122 > via Tunnel Composite, UDP (src 10.67.121.135 dest 10.70.35.135), Push 19
123 192.168.101.0/24 *[BGP/170] 20w6d 01:03:37, localpref 100, from 10.70.35.129
124 AS path: I, validation-state: unverified
125 > via Tunnel Composite, UDP (src 10.67.121.135 dest 10.70.35.129), Push 16
126 192.168.210.0/24 *[BGP/170] 20w5d 22:48:15, localpref 100, from 10.70.35.129
127 AS path: I, validation-state: unverified
128 > via Tunnel Composite, UDP (src 10.67.121.135 dest 10.70.35.129), Push 16
129 192.168.220.0/24 *[BGP/170] 8w5d 16:05:07, localpref 100, from 10.70.35.135
130 AS path: I, validation-state: unverified
131 > via Tunnel Composite, UDP (src 10.67.121.135 dest 10.70.35.135), Push 17
132 192.168.221.0/24 *[BGP/170] 8w5d 16:05:07, localpref 100, from 10.70.35.135
133 AS path: I, validation-state: unverified
134 > via Tunnel Composite, UDP (src 10.67.121.135 dest 10.70.35.135), Push 18
135 192.168.230.0/24 *[Direct/0] 20w5d 22:39:24
136 > via jvknet1-51e5c20
137 192.168.230.21/32 *[Local/0] 20w5d 22:39:24
138 Local via jvknet1-51e5c20
139
140 [root@octo-test-jcnr-3 ~]#

```

1.5.3 Validation of end-end connectivity

1. Check for the Intra Metro traffic with ping and iperf3 TCP throughput between LAN1 and LAN2 BMS servers via JCNR-1 – VLAN2140 – JCNR-2.

```

1 root@c3-small-x86-xrd-1-LAN-1:~# ping 192.168.42.21 -I 192.168.41.21 -c 5
2 PING 192.168.42.21 (192.168.42.21) from 192.168.41.21 : 56(84) bytes of data.
3 64 bytes from 192.168.42.21: icmp_seq=1 ttl=62 time=1.93 ms
4 64 bytes from 192.168.42.21: icmp_seq=2 ttl=62 time=0.827 ms
5 64 bytes from 192.168.42.21: icmp_seq=3 ttl=62 time=0.737 ms

```

```

6 64 bytes from 192.168.42.21: icmp_seq=4 ttl=62 time=0.664 ms
7 64 bytes from 192.168.42.21: icmp_seq=5 ttl=62 time=0.750 ms
8
9 --- 192.168.42.21 ping statistics ---
10 5 packets transmitted, 5 received, 0% packet loss, time 4065ms
11 rtt min/avg/max/mdev = 0.664/0.982/1.934/0.478 ms
12 [root@octo-test-jcncr-3 ~]#
13
14 root@c3-small-x86-xrd-1-LAN-2:~# iperf3 -c 192.168.41.21 -B 192.168.42.21 -t 5
15 Connecting to host 192.168.41.21, port 5201
16 [ 4] local 192.168.42.21 port 44433 connected to 192.168.41.21 port 5201
17 [ ID] Interval Transfer Bandwidth Retr Cwnd
18 [ 4] 0.00-1.00 sec 1.05 GBytes 8.98 Gbits/sec 4 1.32 MBytes
19 [ 4] 1.00-2.00 sec 1.06 GBytes 9.08 Gbits/sec 2 4.25 MBytes
20 [ 4] 2.00-3.00 sec 1.06 GBytes 9.08 Gbits/sec 0 4.32 MBytes
21 [ 4] 3.00-4.00 sec 1.05 GBytes 9.05 Gbits/sec 0 3.25 MBytes
22 [ 4] 4.00-5.00 sec 1.05 GBytes 9.06 Gbits/sec 2 4.32 MBytes
23 - - - - -
24 [ ID] Interval Transfer Bandwidth Retr
25 [ 4] 0.00-5.00 sec 5.27 GBytes 9.05 Gbits/sec 6 sender
26 [ 4] 0.00-5.00 sec 5.26 GBytes 9.04 Gbits/sec receiver
27 [root@octo-test-jcncr-3 ~]#

```

2. Check for the Inter Metro traffic with ping and iperf3 TCP throughput between LAN2 and LAN3 BMS servers via JCNR-2 – L3VPNoMPLSoUDP – JCNR-3.

```

1 root@c3-small-x86-xrd-1-LAN-1:~# ping 192.168.33.21 -I 192.168.41.21 -c 5
2 PING 192.168.33.21 (192.168.33.21) from 192.168.41.21 : 56(84) bytes of data.
3 64 bytes from 192.168.33.21: icmp_seq=1 ttl=62 time=43.9 ms
4 64 bytes from 192.168.33.21: icmp_seq=2 ttl=62 time=42.8 ms
5 64 bytes from 192.168.33.21: icmp_seq=3 ttl=62 time=42.9 ms
6 64 bytes from 192.168.33.21: icmp_seq=4 ttl=62 time=43.0 ms
7 64 bytes from 192.168.33.21: icmp_seq=5 ttl=62 time=43.0 ms
8
9 --- 192.168.33.21 ping statistics ---
10 5 packets transmitted, 5 received, 0% packet loss, time 4007ms
11 rtt min/avg/max/mdev = 42.825/43.107/43.866/0.382 ms
12 root@c3-small-x86-xrd-1-LAN-1:~#
13
14 root@c3-small-x86-xrd-3-LAN-1:~# iperf3 -c 192.168.42.21 -B 192.168.33.21 -t 5
15 Connecting to host 192.168.42.21, port 5201
16 [ 4] local 192.168.33.21 port 52317 connected to 192.168.42.21 port 5201
17 [ ID] Interval Transfer Bandwidth Retr Cwnd
18 [ 4] 0.00-1.00 sec 130 MBytes 1.09 Gbits/sec 0 12.1 MBytes
19 [ 4] 1.00-2.00 sec 172 MBytes 1.45 Gbits/sec 0 12.1 MBytes
20 [ 4] 2.00-3.00 sec 179 MBytes 1.50 Gbits/sec 0 12.1 MBytes
21 [ 4] 3.00-4.00 sec 174 MBytes 1.46 Gbits/sec 0 12.1 MBytes
22 [ 4] 4.00-5.00 sec 174 MBytes 1.46 Gbits/sec 2 6.33 MBytes
23 - - - - -
24 [ ID] Interval Transfer Bandwidth Retr
25 [ 4] 0.00-5.00 sec 829 MBytes 1.39 Gbits/sec 2 sender
26 [ 4] 0.00-5.00 sec 825 MBytes 1.38 Gbits/sec receiver
27 [root@octo-test-jcncr-3 ~]#

```

3. Check for the iperf3 TCP throughput between AWS EC2 instance and Equinix Metal LAN-3 server via JCNR-3 over a 500Mbps AWS Direct Connect Circuit.

```

1 ubuntu-173-29-5-21:~$ iperf3 -c 192.168.33.21 -B 173.29.5.21 -M 3000
2 Connecting to host 192.168.33.21, port 5201
3 [ 5] local 173.29.5.21 port 41875 connected to 192.168.33.21 port 5201
4 [ ID] Interval Transfer Bitrate Retr Cwnd
5 [ 5] 0.00-1.00 sec 65.7 MBytes 551 Mb/s/sec 285 147 KBytes
6 [ 5] 1.00-2.00 sec 52.5 MBytes 440 Mb/s/sec 4 243 KBytes
7 [ 5] 2.00-3.00 sec 53.8 MBytes 451 Mb/s/sec 9 157 KBytes
8 [ 5] 3.00-4.00 sec 58.8 MBytes 493 Mb/s/sec 8 137 KBytes
9 [ 5] 4.00-5.00 sec 61.2 MBytes 514 Mb/s/sec 6 126 KBytes
10 [ 5] 5.00-6.00 sec 55.0 MBytes 461 Mb/s/sec 6 168 KBytes
11 [ 5] 6.00-7.00 sec 57.5 MBytes 482 Mb/s/sec 9 119 KBytes
12 [ 5] 7.00-8.00 sec 57.5 MBytes 482 Mb/s/sec 7 83.4 KBytes
13 [ 5] 8.00-9.00 sec 58.8 MBytes 493 Mb/s/sec 5 158 KBytes
14 [ 5] 9.00-10.00 sec 58.8 MBytes 493 Mb/s/sec 7 137 KBytes
15 - - - - -
16 [ ID] Interval Transfer Bitrate Retr
17 [ 5] 0.00-10.00 sec 579 MBytes 486 Mb/s/sec 346 sender
18 [ 5] 0.00-10.00 sec 576 MBytes 483 Mb/s/sec receiver

```

4. Check for the iperf3 TCP throughput between Azure VM instance and Equinix Metal LAN-3 server via JCNR-3 over a 500Mbps Azure Express Route Circuit.

```

1 pceiadmin@PCEI.11M-01:~$ iperf3 -c 192.168.33.21 -8 18.126.1.7 -M 3000
2 Connecting to host 192.168.33.21, port 5201
3 [ 4] local 10.126.1.7 port 41069 connected to 192.168.33.21 port 5281
4 [ ID] Interval Transfer Bitrate Retr Cwnd
5 [ 4] 0.00-1.00 sec 73.9 MBytes 619 Mbits/sec 1883 217 KBytes
6 [ 4] 1.00-2.00 sec 59.8 MBytes 502 Mbits/sec 100 218 KBytes
7 [ 4] 2.00-3.00 sec 55.8 MBytes 468 Mbits/sec 42 273 KBytes
8 [ 4] 3.00-4.00 sec 64.1 MBytes 537 Mbits/sec 137 180 KBytes
9 [ 4] 4.00-5.00 sec 58.1 MBytes 488 Mbits/sec 62 179 KBytes
10 [ 4] 5.00-6.00 sec 53.9 MBytes 452 Mbits/sec 51 235 KBytes
11 [ 4] 6.00-7.00 sec 68.2 MBytes 572 Mbits/sec 145 214 KBytes
12 [ 4] 7.00-8.00 sec 59.5 MBytes 499 Mbits/sec 64 216 KBytes
13 [ 4] 8.00-9.00 sec 55.4 MBytes 465 Mbits/sec 45 283 KBytes
14 [ 4] 9.00-10.00 sec 65.0 MBytes 545 Mbits/sec 156 192 KBytes
15 - - - - -
16 [ ID] Interval Transfer Bitrate Retr
17 [ 4] 0.00-10.00 sec 614 MBytes 515 Mbits/sec 2685 sender
18 [ 4] 0.00-10.00 sec 611 MBytes 513 Mbits/sec receiver

```

5. Check for iperf3 TCP throughput between AWS EC2 and Azure VM instances through JCNR-3 over a combination of Azure Express Route and AWS Direct Connect VCs with 500Mbps.

```

1 ubuntu@ip-173-29-5-21:~$ iperf3 -c 19.126.1.7 -B 173.29.5. 1 -M 1412
2 Connecting to host 10.126.1.7, port 5201
3 [ 5] local 173.29.5.21 port 37001 connected to 10.126.1.7 port 5201
4 [ ID] Interval Transfer Bitrate Retr Cwnd
5 [ 5] 8.00-1.00 sec 57.4 MBytes 481 Mbits/sec 73 776 KBytes
6 [ 5] 1.00-2.00 sec 60.9 MBytes 504 Mbits/sec 0 821 KBytes
7 [ 5] 2.00-3.00 sec 58.8 MBytes 493 Mbits/sec 2 627 KBytes
8 [ 5] 3.00-4.00 sec 45.8 MBytes 377 Mbits/sec 3 281 KBytes
9 [ 5] 4.00-5.00 sec 55.0 MBytes 461 Mbits/sec 0 397 KBytes
10 [ 5] 5.00-6.00 sec 60.0 MBytes 503 Mbits/sec 0 495 KBytes
11 [ 5] 6.00-7.00 sec 55.0 MBytes 461 Mbits/sec 0 569 KBytes
12 [ 5] 7.00-8.00 sec 60.0 MBytes 503 Mbits/sec 0 641 KBytes
13 [ 5] 8.00-9.00 sec 60.0 MBytes 503 Mbits/sec 0 706 KBytes
14 [ 5] 9.00-10.00 sec 46.2 MBytes 388 Mbits/sec 4 300 KBytes
15 - - - - -
16 [ ID] Interval Transfer Bitrate Retr
17 [ 5] 9.08-10.00 sec 557 MBytes 468 Mbits/sec 82 sender
18 [ 5] 0.09-10.00 sec 554 MBytes 465 Mbits/sec receiver

```

1.6 Conclusion

The demonstration of the above use case shows how Juniper Cloud Native Router can be used to provide seamless connectivity between various types of workloads that seamlessly cuts across cloud boundaries. In essence what this means to network operators and enterprise architects is the transformational benefits of a cloud-native networking model while maintaining operational consistency with the existing infrastructure, which means lower opex, faster speed of business, and carrier-grade reliability.

1.7 Useful links

- JCNR configurations, server inventory, and JUNOS verification logs: <https://github.com/pvs2401/Techposts/wiki>
- JCNR Documentation: <https://www.juniper.net/documentation/product/us/en/juniper-cloud-native-router/>

1.8 Glossary

- ACL: Access List
- AWS: Amazon Web Services
- BGP: Border Gateway Protocol
- CNI: Container Network Interface
- cRPD: Containerized Route Processing Daemon
- CSR: Cell Site Router
- D-RAN: Distributed Radio Access Network
- DPDK: Data Plane Development Kit
- EVPN-VxLAN: Network Attachment Definition
- JCNR: Juniper Cloud-Native Router
- NAD: Network Attachment Definition
- GRPC: Google Remote Procedure Call
- QoS: Quality of Service
- TWAMP: Two-Way Active Measurement Protocol
- VPC: Virtual Private Cloud

1.9 Acknowledgements

Many thanks to Oleg Berlin from Equinix who was instrumental in putting together this JCNR solution with Equinix Metal.

Author would like to thank Vinod Nair (Juniper Networks) for his help in putting this solution together in such a short time.

Special thanks to Julian Lucek, Guy Davies for their review of this article.



DC – AI – Routing – Switching – Security – Wireless – Automation

© Copyright 2026 Hewlett Packard Enterprise Development LP. The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty.

Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Trademark acknowledgments, if needed. All third-party marks are the property of their respective owners.

HEWLETT PACKARD ENTERPRISE

HPE.com

