



DC – AI – Routing – Switching – Security – Wireless – Automation

Link Slicing with MPLS and SRv6 Underlays

An HPE Juniper Networking Techpost

Krzysztof Szarkowicz

2023-05-12

HPE



Contents

1	Link Slicing with MPLS and SRv6 Underlays	4
1.1	Architecture Introduction	4
1.2	Link Slicing Introduction	5
1.3	Link Slicing with Slice-Aware H-QoS	6
1.4	Slice-Aware H-QoS Configuration	7
1.5	Assigning Traffic to Slices – MPLS	12
1.6	Assigning Traffic to Slices – SRv6	17
1.7	Next Steps	24
1.8	Useful links	25
1.9	Glossary	25
1.10	Acknowledgements	26

List of Figures

1	Inter-AS Topology with slicing on Inter-AS links	5
2	Link Slicing with legacy H-QoS	6
3	Link Slicing with slice-aware H-QoS	7
4	Traffic assignment to slices – MPLS	12
5	SRv6 SID Allocation Scheme	17
6	Traffic assignment to slices – SRv6	18
7	IPv6 header	20

List of Tables

1	Slice QoS profiles	10
2	Slice MPLS service (bottom) label ranges	13

1 Link Slicing with MPLS and SRv6 Underlays

Krzysztof Szarkowicz - 05/12/2023

The guaranteed link slicing feature, using MPLS and SRv6 as underlay transport.

Link slicing is a way to share physical bandwidth on links between multiple tenants, and *guaranteed* means providing minimum (guaranteed) bandwidth per tenant in case of congestion, as well as a possibility to enforce maximum transmit rates per tenant. Any leftover or unused bandwidth can be shared between tenants. Each slice can have own queue configurations for different classes of service within that slice.

This article is based on the capabilities of Junos 23.1R2 (or newer Junos release) running on MX Series routers. Configuration and operational command outputs have been collected on vMX and MX in our labs.

You can test yourself all the concepts described in this article, we created labs in JCL and vLabs: [JCL \(Junivators and partners\)](#) [vLabs \(open to all\)](#)

In this blog post, following IP (Internet Protocol) addressing is used:

Transport Infrastructure (P/PE)

- Router-ID: 198.51.100.
- Loopback: 2001:db8:bad:cafe:00::/128
- SRv6 locator: fc01:::/48
- Core Links: 2001:db8:beef:00::/112

PE-CE links:

- IPv4: .../24
- IPv6: 2001:db8:babe:face:::/112

VPN (Virtual Private Network) Loopbacks (CE/PE):

- 192.168../32
- 2001:db8:abba::/128

1.1 Architecture Introduction

The network topology and initial configuration used for this blog post is based on the network topology and configuration discussed in the [SRv6 L3VPN Inter-AS Option-C blog post](#). That is, we have multi-domain network architecture with VRFs placed on the PE routers, and L3VPN over SRv6 Inter-AS Option-C framework to provide L3VPN connectivity between VRFs on the PEs in different domains. In addition to SRv6 Inter-AS Option-C, SR-MPLS Inter-AS Option-C has been added, so that link slicing for both MPLS and SRv6 underlays can be demonstrated. This is outlined in Figure 1.




Figure 1: Inter-AS Topology with slicing on Inter-AS links

PE routers have multiple L3VPNs. Some of these VPNs use SRv6 as underlay (see the [SRv6 SID Encoding and Transposition blog post](#) for more details about L3VPN over SRv6), and some of these VPNs use SR-MPLS as underlay. Thus, both underlay types (SR-MPLS and SRv6) are used concurrently. This is a realistic scenario and might occur, for example, during migration from MPLS to SRv6 underlay.

Nevertheless, focus of this blog post is not the discussion about nuances of concurrent running of SR-MPLS or SRv6 underlays. Focus of this blog post is link slicing, and both underlays are simply used to illustrate that link slicing can work with both MPLS (any type of MPLS, not only SR-MPLS) and SRv6.

1.2 Link Slicing Introduction

Before going into configuration or operational commands details, let's discuss the use case: what is a "link slicing", where and how we could use link slicing?

Link slicing is a specific use case under overall umbrella called "network slicing". With link slicing, you slice (channelize, partition, divide, ... – use the word you like) a single link in such a way, that each slice has explicit capacity guarantee, and each slice can have multiple traffic classes and queues. For example, Flexible Ethernet (FlexE) is a technology that allows to channelize an Ethernet link, with fixed bandwidth guarantees for each channel.

Saying that, FlexE might not be the best technology for link slicing/channelization. The major drawbacks of FlexE, when used for link slicing/channelization are:

- no statistical multiplex gain/no bandwidth reuse between channels
- requires large physical links (50 Gbps and above) with large b/w increments (5 Gbps)
- requires an underlying electrical transport switching layer to support channelization

Therefore, to overcome the drawbacks associated with FlexE, this blog post uses different approach for link slicing, utilizing hierarchical QoS (H-QoS) toolset. Traditional, legacy H-QoS architecture is depicted in Figure 2.




Figure 2: Link Slicing with legacy H-QoS

In this architecture the link is divided into subinterfaces (units in Junos terms), where each subinterface is associated with a VLAN. The QoS profile (traffic control profile in Junos terms) contains QoS parameters for each subinterface, like:

- CIR (Committed Information Rate) – guaranteed minimum rate
- PIR (Peak Information Rate) – maximum (shaping) rate
- Queue parameters inside each profile, with queue priorities, queue sizes, minimum/maximum rate of each queue, etc.

We could reuse this model for link slicing – theoretically at least. So, why we don't do it? What is the problem with this model?

Well, if you look at Figure 1, the depicted use case for link slicing is to slice inter-AS link (please note, this is just an example; there could be many other use cases calling for link slicing at different locations in the network). With traditional H-QoS approach for link slicing, it would mean considerable administrative overhead:

- VLAN allocation and co-ordination between two domains
- IP addressing for each subinterface allocation and co-ordination between two domains
- In case of intra-domain link slicing, multiple IGP adjacencies would need to be started to make the subinterfaces usable for traffic
- In case of inter-domain link slicing, multiple eBGP peerings, and/or import/export BGP policies with next-hop manipulation to make the subinterfaces usable for traffic

This might become complex task, especially, if the use case demands slicing on multiple links.

1.3 Link Slicing with Slice-Aware H-QoS

To address this complexity, starting from Junos 22.2R1, Juniper successively began to introduce features allowing H-QoS deployments for link slicing in much simpler manner (slice-aware H-QoS). This blog post uses features available from Junos 23.1R2.




Figure 3: Link Slicing with slice-aware H-QoS

In essence, slice-aware H-QoS uses abstract objects called ‘slices’ to attach QoS profiles on the link. So, no subinterfaces, VLANs, multiple IP addresses, multiple IGP/BGP sessions on the link are longer required, as with traditional H-QoS model. This simplifies operations and does not affect packet forwarding in any way.

When packets enter the router, they are classified as belonging to a specific slice. Packets not classified explicitly, are associated with a default slice. Packet classification, implemented with firewall filter, can happen practically based on any existing field in the packet, like for example top MPLS label, bottom MPLS label, Src/Dst IP address, some specific bits from Src/Dst IP address, etc.

When the packet leaves the router on sliced link, packets are subject to the treatment defined in the slice-specific QoS profile, based on the slice selection performed on input.

1.4 Slice-Aware H-QoS Configuration

In the topology depicted in Figure 1, link slicing is configured on inter-AS links. As an example, configuration of P1 router will be discussed.

First, the slices must be initialized, as outlined in Configuration 1. As an example, we will be using 3 slices, called NS-A, NS-B and NS-C, to illustrate link slicing capability.

```

1 1 services {
2 2 network-slicing {
3 3 slice NS-A;
4 4 slice NS-B;
5 5 slice NS-C;
6 6 }
7 7 }

```

Configuration 1: Slice initialization

Second, hierarchical scheduler capability on the sliced interface must be enabled, as outlined in Configuration 2. Please note, up to Junos 23.1R1 enabling hierarchical scheduler is possible only on interface divided into subinterfaces (VLANs), to support classical H-QoS. Junos 23.1R2 removes that restriction, allowing enabling

hierarchical scheduler capability on plain (not divided to subinterfaces) interfaces as well, for slice-aware H-QoS support.

```
1 1 interfaces {
2 2 ge-0/0/3 {
3 3 hierarchical-scheduler;
4 4 }
5 5 }
```

Configuration 2: Enabling H-QoS capability

Now, QoS configuration must be prepared, with QoS profiles attached to the slices on the interface. This blog post does not intent to cover Junos QoS in details, as this is a huge topic. Therefore, to demonstrate slicing capability, relatively simple QoS configuration is used, as outlined in Configuration 3:

```
1 1 class-of-service {
2 2 classifiers {
3 3 dscp CL-DSCP {
4 4 forwarding-class FC-BE {
5 5 loss-priority low code-points be;
6 6 }
7 7 forwarding-class FC-EF {
8 8 loss-priority low code-points ef;
9 9 }
10 10 }
11 11 dscp-ipv6 CL-DSCP {
12 12 forwarding-class FC-BE {
13 13 loss-priority low code-points be;
14 14 }
15 15 forwarding-class FC-EF {
16 16 loss-priority low code-points ef;
17 17 }
18 18 }
19 19 exp CL-MPLS {
20 20 forwarding-class FC-BE {
21 21 loss-priority low code-points 000;
22 22 }
23 23 forwarding-class FC-EF {
24 24 loss-priority low code-points 101;
25 25 }
26 26 }
27 27 }
28 28 forwarding-classes {
29 29 class FC-BE queue-num 0 priority low;
30 30 class FC-EF queue-num 1 priority high;
31 31 }
32 32 traffic-control-profiles {
33 33 TC-1G {
34 34 shaping-rate 1g;
35 35 }
36 36 TC-NS-A {
37 37 scheduler-map SM-NS-A;
38 38 shaping-rate 100m;
39 39 guaranteed-rate 50m;
40 40 }
41 41 TC-NS-B {
42 42 scheduler-map SM-NS-B;
43 43 shaping-rate 5600000;
44 44 guaranteed-rate 5m;
45 45 }
46 46 TC-NS-C {
47 47 scheduler-map SM-NS-C;
48 48 shaping-rate 5100000;
49 49 guaranteed-rate 4m;
50 50 }
51 51 }
52 52 interfaces {
53 53 ge-* {
54 54 unit * {
55 55 classifiers {
56 56 dscp CL-DSCP;
57 57 dscp-ipv6 CL-DSCP;
58 58 exp CL-MPLS;
59 59 }
60 60 }
61 61 }
```

```

62 62 ge-0/0/3 {
63 63   output-traffic-control-profile TC-1G;
64 64   slice NS-A {
65 65     output-traffic-control-profile TC-NS-A;
66 66   }
67 67   slice NS-B {
68 68     output-traffic-control-profile TC-NS-B;
69 69   }
70 70   slice NS-C {
71 71     output-traffic-control-profile TC-NS-C;
72 72   }
73 73 }
74 74 }
75 75 scheduler-maps {
76 76   SM-NS-A {
77 77     forwarding-class FC-BE scheduler SC-BE;
78 78     forwarding-class FC-EF scheduler SC-EF;
79 79   }
80 80   SM-NS-B {
81 81     forwarding-class FC-BE scheduler SC-BE;
82 82     forwarding-class FC-EF scheduler SC-EF;
83 83   }
84 84   SM-NS-C {
85 85     forwarding-class FC-BE scheduler SC-BE;
86 86     forwarding-class FC-EF scheduler SC-EF;
87 87   }
88 88 }
89 89 schedulers {
90 90   SC-BE {
91 91     transmit-rate {
92 92       remainder;
93 93     }
94 94     priority low;
95 95   }
96 96   SC-EF {
97 97     transmit-rate {
98 98       percent 50;
99 99     rate-limit;
100 100   }
101 101   priority strict-high;
102 102 }
103 103 }
104 104 }

```

Configuration 3: Slice-aware H-QoS

The main aspects of this sample QoS configuration are as follows:

- Two forwarding classes – FC-BE and FC-EF – (lines 28-31). Please note, this is only example. Each slice can support up to 8 forwarding classes.
- DSCP and MPLS TC classifiers to classify packets into the forwarding class based on DSCP or MPLS TC values in the packet (lines 2-27). Additionally, classifiers are assigned to the interfaces (lines 53-61).
- Traffic control profiles – per port and per slice profiles – (lines 32-51). Per slice profiles have some example minimum (guaranteed) rates and maximum (shaping) rates. Additionally, traffic control profiles is attached to the sliced interface (lines 62-73), resulting in H-QoS hierarchy depicted in Figure 1.
- Traffic control profiles reference scheduler maps (lines 37, 42, 47, 75-88), which define queue parameters for each slice via schedulers (lines 89-103). In this simple example, each slice uses the same queue parameters for simplification. However, in the production deployment, each slice can be parametrized differently, based on the actual requirements.

As the result of the QoS configuration, we have three slices as follows:

Table 1: Slice QoS profiles

Slice Name	min BW	max BW	Queues
NS-A	50 Mbps	100 Mbps	FC-EF: strict-high, 50% (rate limited)FC-BE: low, remaining slice BW
NS-B	5 Mbps	5.6 Mbps	FC-EF: strict-high, 50% (rate limited)FC-BE: low, remaining slice BW
NS-C	4 Mbps	5.1 Mbps	FC-EF: strict-high, 50% (rate limited)FC-BE: low, remaining slice BW

Based on the configuration done so far, initial checks can be performed (CLI-Output 1). Please note, based on the MX line card used, the output of the operational command might be slightly different.

```

1 1 kszarkowicz@P1> show interfaces queue ge-0/0/3 slice NS-A
2 2 Slice : NS-A (Index : 1)
3 3 Anchor interface : ge-0/0/3 (Index : 152)
4 4 Forwarding classes: 16 supported, 4 in use
5 5 Egress queues: 8 supported, 4 in use
6 6 Queue: 0, Forwarding classes: FC-BE
7 7 Queued:
8 8 Packets : 0 0 pps
9 9 Bytes : 0 0 bps
10 10 Transmitted:
11 11 Packets : 0 0 pps
12 12 Bytes : 0 0 bps
13 13 Tail-dropped packets : 0 0 pps
14 14 RL-dropped packets : 0 0 pps
15 15 RL-dropped bytes : 0 0 bps
16 16 RED-dropped packets : 0 0 pps
17 17 Low : 0 0 pps
18 18 Medium-low : 0 0 pps
19 19 Medium-high : 0 0 pps
20 20 High : 0 0 pps
21 21 RED-dropped bytes : 0 0 bps
22 22 Low : 0 0 bps
23 23 Medium-low : 0 0 bps
24 24 Medium-high : 0 0 bps
25 25 High : 0 0 bps
26 26 Queue-depth bytes :
27 27 Average : 0
28 28 Current : 0
29 29 Peak : 0
30 30 Maximum : 1564672
31 31 Queue: 1, Forwarding classes: FC-EF
32 32 Queued:
33 33 Packets : 0 0 pps
34 34 Bytes : 0 0 bps
35 35 Transmitted:
36 36 Packets : 0 0 pps
37 37 Bytes : 0 0 bps
38 38 Tail-dropped packets : 0 0 pps
39 39 RL-dropped packets : 0 0 pps
40 40 RL-dropped bytes : 0 0 bps
41 41 RED-dropped packets : 0 0 pps
42 42 Low : 0 0 pps
43 43 Medium-low : 0 0 pps
44 44 Medium-high : 0 0 pps
45 45 High : 0 0 pps
46 46 RED-dropped bytes : 0 0 bps
47 47 Low : 0 0 bps
48 48 Medium-low : 0 0 bps
49 49 Medium-high : 0 0 bps

```

```

50 50 High : 0 0 bps
51 51 Queue-depth bytes :
52 52 Average : 0
53 53 Current : 0
54 54 Peak : 0
55 55 Maximum : 1564672
56 56 Queue: 2, Forwarding classes: assured-forwarding
57 57 Queued:
58 58 Packets : 0 0 pps
59 59 Bytes : 0 0 bps
60 60 Transmitted:
61 61 Packets : 0 0 pps
62 62 Bytes : 0 0 bps
63 63 Tail-dropped packets : 0 0 pps
64 64 RL-dropped packets : 0 0 pps
65 65 RL-dropped bytes : 0 0 bps
66 66 RED-dropped packets : 0 0 pps
67 67 Low : 0 0 pps
68 68 Medium-low : 0 0 pps
69 69 Medium-high : 0 0 pps
70 70 High : 0 0 pps
71 71 RED-dropped bytes : 0 0 bps
72 72 Low : 0 0 bps
73 73 Medium-low : 0 0 bps
74 74 Medium-high : 0 0 bps
75 75 High : 0 0 bps
76 76 Queue-depth bytes :
77 77 Average : 0
78 78 Current : 0
79 79 Peak : 0
80 80 Maximum : 32768
81 81 Queue: 3, Forwarding classes: network-control
82 82 Queued:
83 83 Packets : 0 0 pps
84 84 Bytes : 0 0 bps
85 85 Transmitted:
86 86 Packets : 0 0 pps
87 87 Bytes : 0 0 bps
88 88 Tail-dropped packets : 0 0 pps
89 89 RL-dropped packets : 0 0 pps
90 90 RL-dropped bytes : 0 0 bps
91 91 RED-dropped packets : 0 0 pps
92 92 Low : 0 0 pps
93 93 Medium-low : 0 0 pps
94 94 Medium-high : 0 0 pps
95 95 High : 0 0 pps
96 96 RED-dropped bytes : 0 0 bps
97 97 Low : 0 0 bps
98 98 Medium-low : 0 0 bps
99 99 Medium-high : 0 0 bps
100 100 High : 0 0 bps
101 101 Queue-depth bytes :
102 102 Average : 0
103 103 Current : 0
104 104 Peak : 0
105 105 Maximum : 32768

```

CLI-Output 1: Initial check for slice NS-A

We can see counters of queues for slice NS-A (counters of queues for slice NS-B and NS-C are not shown for brevity). At the moment, all counters are '0', despite the traffic already flows through the network (traffic generators generating traffic are attached to CE devices). The reason is, the current configuration defines slices, slice QoS profiles, and assigns slice QoS profiles to the interface. However, the current configuration doesn't assign traffic to slices, so at the moment all traffic is using default slice only.

Depending on the MX hardware used, it might happen – especially on older MX line cards – that configuration discussed so far is not sufficient to enable slice-aware H-QoS (CLI-Output 2).

```

1 1 kszarkowicz@rtme-mx-18> show interfaces queue et-3/0/2 slice NS-A
2 2 error: Get slice-id from slice-name:NS-A on et-3/0/2 failed. Error: No such file or directory.
3 3
4 4
5 5 kszarkowicz@PE23> show interfaces queue ge-0/0/3 slice NS-A
6 6 error: slice 'NS-A' not found on ge-0/0/3. Abort

```

CLI-Output 2: Failed initial check for slice NS-A

In such situation, additional configuration (enabling egress traffic manager and/or enabling flexible queueing mode) might be required (Configuration 4).

```

1 1 chassis {
2   2 fpc 0 {
3    3 pic 0 {
4     4 traffic-manager {
5      5 mode egress-only;
6     6 }
7    7 }
8   8 flexible-queueing-mode;
9   9 }
10 10 }

```

*Configuration 4: Enhanced QoS***1.5 Assigning Traffic to Slices – MPLS**

Assigning traffic to appropriate slice is a multiple step process.

First of all, there must be common agreement in the network, which fields in the packet will be used for slice identification. For MPLS based underlay the obvious choice is MPLS label. The Junos framework for link slicing is very flexible, allowing MPLS labels (or label ranges) at any position (e.g., top, bottom with/without offset from top/bottom) in the label stack to be used for slice identification purposes.

In this blog post L3VPN labels, which will be present at the bottom of the label stack (SR-TE or TI-LFA could use multiple transport labels above bottom service label) in the packet, will be used for slice identification, as depicted in Figure 4.




Figure 4: Traffic assignment to slices – MPLS

Following L3VPN label ranges are used in this blog post (again, this is just simple example – Junos link slicing framework doesn't put any restrictions here):

Table 2: Slice MPLS service (bottom) label ranges

Slice Name	Min Label	Max Label
NS-A	1000010	1000019
NS-B	1000020	1000029
NS-C	1000030	1000039

The selected label ranges are from the default Junos range used for static labels, as outlined in CLI-Output 3, lines 11 and 17.

```

1 1 kszarkowicz@PE11> show mpls label usage
2 2 Label space Total Available Applications
3 3 LSI 949984 949983 (100.00%) BGP/LDP VPLS with no-tunnel-services, BGP L3VPN with vrf-table-
   label
4 4 Block 949984 949983 (100.00%) BGP/LDP VPLS with tunnel-services, BGP L2VPN
5 5 Dynamic 949984 949983 (100.00%) RSVP, LDP, PW, L3VPN, RSVP-P2MP, LDP-P2MP, MVPN, EVPN, BGP
6 6 Static 48576 48567 (99.98%) Static LSP, Static PW
7 7 Effective Ranges
8 8 Range name Shared with Start End
9 9 Dynamic 16 99999
10 10 Dynamic 150000 999999
11 11 Static 1000000 1048575
12 12 SRGB 100000 149999 ISIS
13 13 Configured Ranges
14 14 Range name Shared with Start End
15 15 Dynamic 16 99999
16 16 Dynamic 150000 999999
17 17 Static 1000000 1048575
18 18 SRGB 100000 149999 ISIS

```

CLI-Output 3: MPLS label ranges

L3VPN service labels are of local significance, and can be re-used on every ingress PE independently. This is important for scaling. If for a particular deployment different MPLS label ranges should be used for slice identification (e.g., in multi-vendor environment, when 3rd party equipment cannot use range 1000000-1048575 for static label), the static label range could be changed with ‘set protocols mpls label-range’ command. This blog post uses default static label range.

Now, when the VRFs are orchestrated on PE routers, they must be orchestrated with appropriate VRF labels (Configuration 5).

```

1 1 routing-instances {
2 2   RI-VRF17 {
3 3     vrf-table-label static 1000017;
4 4   }
5 5   RI-VRF27 {
6 6     vrf-table-label static 1000027;
7 7   }
8 8   RI-VRF37 {
9 9     vrf-table-label static 1000037;
10 10  }
11 11 }

```

Configuration 5: Label assignment to VRFs

With this configuration, traffic of VRF17 will use service (bottom) label 1000017, traffic of VRF27 will use service (bottom) label 1000027, and traffic of VRF37 will use service (bottom) label 1000037. This corresponds to the MPLS label ranges defined in Table 2.

When the packets arrive to P1, which is agnostic to existing VPNs, but should perform slicing on inter-AS links,

packets can be classified to slices based on the bottom label ranges (Configuration 6). Multiple label ranges can be specified in a single match term.

```

1 1 groups {
2 2 GR-CORE-INTF {
3 3 interfaces {
4 4 <*> {
5 5 unit 0 {
6 6 family mpls {
7 7 filter {
8 8 input FF-MPLS-SLICE-CLASSIFIER;
9 9 }
10 10 }
11 11 }
12 12 }
13 13 }
14 14 }
15 15 }
16 16 interfaces {
17 17 ge-0/0/0 {
18 18 apply-groups GR-CORE-INTF;
19 19 }
20 20 ge-0/0/1 {
21 21 apply-groups GR-CORE-INTF;
22 22 }
23 23 ge-0/0/2 {
24 24 apply-groups GR-CORE-INTF;
25 25 }
26 26 }
27 27 firewall {
28 28 family mpls {
29 29 filter FF-MPLS-SLICE-CLASSIFIER {
30 30 term TR-SLICE-A {
31 31 from {
32 32 label 1000010-1000019 {
33 33 bottom;
34 34 }
35 35 }
36 36 then {
37 37 slice NS-A;
38 38 count CT-SLICE-A;
39 39 accept;
40 40 }
41 41 }
42 42 term TR-SLICE-B {
43 43 from {
44 44 label 1000020-1000029 {
45 45 bottom;
46 46 }
47 47 }
48 48 then {
49 49 slice NS-B;
50 50 count CT-SLICE-B;
51 51 accept;
52 52 }
53 53 }
54 54 term TR-SLICE-C {
55 55 from {
56 56 label 1000030-1000039 {
57 57 bottom;
58 58 }
59 59 }
60 60 then {
61 61 slice NS-C;
62 62 count CT-SLICE-C;
63 63 accept;
64 64 }
65 65 }
66 66 term TR-ALL {
67 67 then {
68 68 count CT-NON-SLICED;
69 69 accept;
70 70 }
71 71 }
72 72 }
73 73 }
74 74 }

```

Configuration 6: Classification to slices based on bottom MPLS label

It is a pretty simple configuration. Firewall filter matches for bottom label ranges (lines 31-35, 43-47, 55-59), and assigns packets to appropriate slice (lines 37, 49, 61). MPLS packets not matched (bottom label not within specific range) are kept in the default slice. Subsequently, the firewall filter is used as input filter on the core interfaces (lines 1-26).

Quick verification confirms that packets are now classified by the filter to slices (CLI-Output 4).

```

1 1 kszarkowicz@P1> show firewall filter FF-MPLS-SLICE-CLASSIFIER
2 2
3 3 Filter: FF-MPLS-SLICE-CLASSIFIER
4 4 Counters:
5 5 Name Bytes Packets
6 6 CT-NON-SLICED 0 0
7 7 CT-SLICE-A 16235184 12824
8 8 CT-SLICE-B 16235184 12824
9 9 CT-SLICE-C 16235184 12824

```

CLI-Output 4: Slice classification based on MPLS

When we now check the queue status of each slice, we see now some counters (CLI-Output 5):

```

1 1 kszarkowicz@P1> show interfaces queue ge-0/0/3 slice NS-A
2 2 Slice : NS-A (Index : 1)
3 3 Anchor interface : ge-0/0/3 (Index : 152)
4 4 Forwarding classes: 16 supported, 4 in use
5 5 Egress queues: 8 supported, 4 in use
6 6 Queue: 0, Forwarding classes: FC-BE
7 7 Queued:
8 8 Packets : 15686 97 pps
9 9 Bytes : 20454544 1019264 bps
10 10 Transmitted:
11 11 Packets : 15686 97 pps
12 12 Bytes : 20454544 1019264 bps
13 13 Tail-dropped packets : 0 0 pps
14 14 RL-dropped packets : 0 0 pps
15 15 RL-dropped bytes : 0 0 bps
16 16 RED-dropped packets : 0 0 pps
17 17 Low : 0 0 pps
18 18 Medium-low : 0 0 pps
19 19 Medium-high : 0 0 pps
20 20 High : 0 0 pps
21 21 RED-dropped bytes : 0 0 bps
22 22 Low : 0 0 bps
23 23 Medium-low : 0 0 bps
24 24 Medium-high : 0 0 bps
25 25 High : 0 0 bps
26 26 Queue: 1, Forwarding classes: FC-EF
27 27 Queued:
28 28 Packets : 15686 97 pps
29 29 Bytes : 20454544 1019520 bps
30 30 Transmitted:
31 31 Packets : 15686 97 pps
32 32 Bytes : 20454544 1019520 bps
33 33 Tail-dropped packets : 0 0 pps
34 34 RL-dropped packets : 0 0 pps
35 35 RL-dropped bytes : 0 0 bps
36 36 RED-dropped packets : 0 0 pps
37 37 Low : 0 0 pps
38 38 Medium-low : 0 0 pps
39 39 Medium-high : 0 0 pps
40 40 High : 0 0 pps
41 41 RED-dropped bytes : 0 0 bps
42 42 Low : 0 0 bps
43 43 Medium-low : 0 0 bps
44 44 Medium-high : 0 0 bps
45 45 High : 0 0 bps
46 46
47 47 (omitted for brevity)
48 48
49 49
50 50
51 51 kszarkowicz@P1> show interfaces queue ge-0/0/3 slice NS-B
52 52 Slice : NS-B (Index : 2)
53 53 Anchor interface : ge-0/0/3 (Index : 152)
54 54 Forwarding classes: 16 supported, 4 in use
55 55 Egress queues: 8 supported, 4 in use

```



```
56 56 Queue: 0, Forwarding classes: FC-BE
57 57 Queued:
58 58 Packets : 17193 97 pps
59 59 Bytes : 22419672 1020928 bps
60 60 Transmitted:
61 61 Packets : 17193 97 pps
62 62 Bytes : 22419672 1020928 bps
63 63 Tail-dropped packets : 0 0 pps
64 64 RL-dropped packets : 0 0 pps
65 65 RL-dropped bytes : 0 0 bps
66 66 RED-dropped packets : 0 0 pps
67 67 Low : 0 0 pps
68 68 Medium-low : 0 0 pps
69 69 Medium-high : 0 0 pps
70 70 High : 0 0 pps
71 71 RED-dropped bytes : 0 0 bps
72 72 Low : 0 0 bps
73 73 Medium-low : 0 0 bps
74 74 Medium-high : 0 0 bps
75 75 High : 0 0 bps
76 76 Queue: 1, Forwarding classes: FC-EF
77 77 Queued:
78 78 Packets : 17193 97 pps
79 79 Bytes : 22419672 1018112 bps
80 80 Transmitted:
81 81 Packets : 17193 97 pps
82 82 Bytes : 22419672 1018112 bps
83 83 Tail-dropped packets : 0 0 pps
84 84 RL-dropped packets : 0 0 pps
85 85 RL-dropped bytes : 0 0 bps
86 86 RED-dropped packets : 0 0 pps
87 87 Low : 0 0 pps
88 88 Medium-low : 0 0 pps
89 89 Medium-high : 0 0 pps
90 90 High : 0 0 pps
91 91 RED-dropped bytes : 0 0 bps
92 92 Low : 0 0 bps
93 93 Medium-low : 0 0 bps
94 94 Medium-high : 0 0 bps
95 95 High : 0 0 bps
96 96
97 97 (omitted for brevity)
98 98
99 99
100 100
101 101 ksarkowicz@P1> show interfaces queue ge-0/0/3 slice NS-C
102 102 Slice : NS-C (Index : 3)
103 103 Anchor interface : ge-0/0/3 (Index : 152)
104 104 Forwarding classes: 16 supported, 4 in use
105 105 Egress queues: 8 supported, 4 in use
106 106 Queue: 0, Forwarding classes: FC-BE
107 107 Queued:
108 108 Packets : 17618 97 pps
109 109 Bytes : 22973872 1017472 bps
110 110 Transmitted:
111 111 Packets : 17618 97 pps
112 112 Bytes : 22973872 1017472 bps
113 113 Tail-dropped packets : 0 0 pps
114 114 RL-dropped packets : 0 0 pps
115 115 RL-dropped bytes : 0 0 bps
116 116 RED-dropped packets : 0 0 pps
117 117 Low : 0 0 pps
118 118 Medium-low : 0 0 pps
119 119 Medium-high : 0 0 pps
120 120 High : 0 0 pps
121 121 RED-dropped bytes : 0 0 bps
122 122 Low : 0 0 bps
123 123 Medium-low : 0 0 bps
124 124 Medium-high : 0 0 bps
125 125 High : 0 0 bps
126 126 Queue: 1, Forwarding classes: FC-EF
127 127 Queued:
128 128 Packets : 17618 97 pps
129 129 Bytes : 22973872 1020288 bps
130 130 Transmitted:
131 131 Packets : 17618 97 pps
132 132 Bytes : 22973872 1020288 bps
133 133 Tail-dropped packets : 0 0 pps
134 134 RL-dropped packets : 0 0 pps
135 135 RL-dropped bytes : 0 0 bps
136 136 RED-dropped packets : 0 0 pps
```

```

137 137 Low : 0 0 pps
138 138 Medium-low : 0 0 pps
139 139 Medium-high : 0 0 pps
140 140 High : 0 0 pps
141 141 RED-dropped bytes : 0 0 bps
142 142 Low : 0 0 bps
143 143 Medium-low : 0 0 bps
144 144 Medium-high : 0 0 bps
145 145 High : 0 0 bps
146 146
147 147 (omitted for brevity)

```

CLI-Output 5: Verification of slice-aware H-QoS

We see around 2 Mbps of traffic in each slice, with 1 Mbps in each forwarding class in each slice (remaining two forwarding classes in each slice are not used in this example, therefore not shown for brevity). All traffic is passing through without any drops, as there is no congestion on the link (so, guaranteed rate doesn't matter), and all slices are within their maximum rate limits (thus, no shaping happens).

This looks good.

Now, let's add some more traffic to the picture.

1.6 Assigning Traffic to Slices – SRv6

This time, let's add SRv6 traffic. And yes, the same link slice can carry different traffic types. We can mix both MPLS and SRv6 flows within the same slice, and apply common QoS guarantees and constraints for such mixed slice. Junos link slicing framework is very flexible.

The process of assigning traffic to slices is similar to MPLS at the high-level. However, this time we don't have MPLS labels, but SRv6 SIDs (see [SRv6 SID Encoding and Transposition blog post](#) (and earlier SRv6 blog posts), SRv6 SID is a 128-bit long data structure, divided into Locator:Function:Argument fields. The division is flexible, and each field can be further decomposed to carry various information.

One example of the of SRv6 SID allocation scheme, used in this blog post, is presented in Figure 5.




Figure 5: SRv6 SID Allocation Scheme

4 bits in the SRv6 locator are used for AS (Domain ID), 16 bits are kept for Node ID. Function as well has designated bits for slice ID (4 bits) and VPN ID (16 bits). Please remember, it is just an example. Depending on the actual use case and requirements, Locator:Function space can be arranged in different ways. For example, slice ID could be part of the SRv6 Locator, not part of Function.

Figure 6 shows the slice selection based on slice ID encoded in SRv6 SID.




Figure 6: Traffic assignment to slices – SRv6

Base configuration for SRv6 Locator and End SID on PE11 is outlined in Configuration 7. There is nothing really new here, when compared to previous SRv6 blog posts.

```

1 1 routing-options {
2 2 source-packet-routing {
3 3 srsv6 {
4 4 locator SL-000 {
5 5 fc01:1:11::/48;
6 6 block-length 32;
7 7 function-length 32;
8 8 static-function-max-entries 1048575;
9 9 }
10 10 no-reduced-srh;
11 11 }
12 12 }
13 13 }
14 14 protocols {
15 15 isis {
16 16 source-packet-routing {
17 17 srsv6 {
18 18 locator SL-000 {
19 19 end-sid fc01:1:11:0:1234:: {
20 20 flavor {
21 21 psp;
22 22 usp;
23 23 usd;
24 24 }
25 25 }
26 26 }
27 27 }
28 28 }
29 29 }
30 30 }

```

Configuration 7: Base SRv6 Locator and End SID

Now, the interesting part is the SRv6 End.DT46 SID allocation (Configuration 8).

```

1 1 routing-instances {
2 2 RI-VRF15 {
3 3 protocols {
4 4 bgp {
5 5 source-packet-routing {
6 6 srsv6 {
7 7 locator SL-000 {
8 8 end-dt46-sid fc01:1:11:a:15::;
9 9 }
10 10 }
11 11 }

```

```

12 12 }
13 13 }
14 14 }
15 15 RI-VRF16 {
16 16 protocols {
17 17   bgp {
18 18     source-packet-routing {
19 19       srv6 {
20 20         locator SL-000 {
21 21           end-dt46-sid fc01:1:11:a:16::;
22 22         }
23 23       }
24 24     }
25 25   }
26 26 }
27 27 }
28 28 RI-VRF25 {
29 29 protocols {
30 30   bgp {
31 31     source-packet-routing {
32 32       srv6 {
33 33         locator SL-000 {
34 34           end-dt46-sid fc01:1:11:b:25::;
35 35         }
36 36       }
37 37     }
38 38   }
39 39 }
40 40 }
41 41 RI-VRF26 {
42 42 protocols {
43 43   bgp {
44 44     source-packet-routing {
45 45       srv6 {
46 46         locator SL-000 {
47 47           end-dt46-sid fc01:1:11:b:26::;
48 48         }
49 49       }
50 50     }
51 51   }
52 52 }
53 53 }
54 54 RI-VRF35 {
55 55 protocols {
56 56   bgp {
57 57     source-packet-routing {
58 58       srv6 {
59 59         locator SL-000 {
60 60           end-dt46-sid fc01:1:11:c:35::;
61 61         }
62 62       }
63 63     }
64 64   }
65 65 }
66 66 }
67 67 RI-VRF36 {
68 68 protocols {
69 69   bgp {
70 70     source-packet-routing {
71 71       srv6 {
72 72         locator SL-000 {
73 73           end-dt46-sid fc01:1:11:c:36::;
74 74         }
75 75       }
76 76     }
77 77   }
78 78 }
79 79 }
80 80 }

```

Configuration 8: SRv6 SID End.DT46

There are six VPNs defined, two for each slice – note ‘a’, ‘b’ and ‘c’ in the SID Function part (lines 8, 21, 34, 47, 60, 73), which is the agreed slice ID. Behind slice ID, you can see VPN ID (note VPN IDs: 15, 16, 25, 26, 35, 36).

Now, when the packets are sent with SRv6 encapsulation (essentially IP in IPv6 encapsulation), the destination address of the outer IPv6 header is equal to the SRv6 End.DT46 SID. Therefore, on P1 router, we need to

match for specific bits encoding slice ID in the destination address of the packet, to classify the packet to a particular slice. For this, it is helpful to understand the IPv6 header structure (Figure 7).




Figure 7: IPv6 header

Initial fields from IPv6 header occupy 64 bits, source address is 128 bits, which gives 192 bits. When looking at the Figure 5, we can observe that SRv6 Locator occupies additional 48 bits (32 bits locator block + 16 bits node ID), and most left byte (8 bits) from function is not used. This gives us in total $192 + 48 + 8 = 248$ bits, or 31 bytes. Therefore, in order to match for slice ID, we will need to match 32nd byte in the IPv6 header.

To match the slice ID on P1 router we will use [firewall filter with flexible match condition](#). This filter provides great flexibility, practically allowing to match any particular bits in the packet header (Configuration 9).

```

1 1 groups {
2 2 GR-CORE-INTF {
3 3 interfaces {
4 4 <*> {
5 5 unit 0 {
6 6 family inet6 {
7 7 filter {
8 8 input FF-SRV6-SLICE-CLASSIFIER;
9 9 }
10 10 }
11 11 }
12 12 }
13 13 }
14 14 }
15 15 }
16 16 interfaces {
17 17 ge-0/0/0 {
18 18 apply-groups GR-CORE-INTF;
19 19 }
20 20 ge-0/0/1 {
21 21 apply-groups GR-CORE-INTF;
22 22 }
23 23 ge-0/0/2 {
24 24 apply-groups GR-CORE-INTF;

```

```

25 25 }
26 26 }
27 27 firewall{
28 28 family inet6 {
29 29 filter FF-SRV6-SLICE-CLASSIFIER {
30 30 term TR-SLICE-A {
31 31 from {
32 32 flexible-match-mask {
33 33 mask-in-hex 0xf;
34 34 prefix 0xa;
35 35 flexible-mask-name FM-SLICE-ID;
36 36 }
37 37 }
38 38 then {
39 39 slice NS-A;
40 40 count CT-SLICE-A;
41 41 accept;
42 42 }
43 43 }
44 44 term TR-SLICE-B {
45 45 from {
46 46 flexible-match-mask {
47 47 mask-in-hex 0xf;
48 48 prefix 0xb;
49 49 flexible-mask-name FM-SLICE-ID;
50 50 }
51 51 }
52 52 then {
53 53 slice NS-B;
54 54 count CT-SLICE-B;
55 55 accept;
56 56 }
57 57 }
58 58 term TR-SLICE-C {
59 59 from {
60 60 flexible-match-mask {
61 61 mask-in-hex 0xf;
62 62 prefix 0xc;
63 63 flexible-mask-name FM-SLICE-ID;
64 64 }
65 65 }
66 66 then {
67 67 slice NS-C;
68 68 count CT-SLICE-C;
69 69 accept;
70 70 }
71 71 }
72 72 term TR-ALL {
73 73 then {
74 74 count CT-NON-SLICED;
75 75 accept;
76 76 }
77 77 }
78 78 }
79 79 }
80 80 flexible-match FM-SLICE-ID {
81 81 match-start layer-3;
82 82 byte-offset 31;
83 83 bit-offset 0;
84 84 bit-length 8;
85 85 }
86 86 }

```

Configuration 9: Classification to slices based on SRv6 slice ID

Lines 80-85 define the flexible match mask, which essentially is the location in the packet, where our match should be performed. In this particular case, the match will be performed in the 32nd byte (byte offset 31, i.e., we are skipping first 31 bytes) of layer 3 header. This flexible match mask is then used in the IPv6 firewall filter so, the layer 3 header becomes IPv6 header (lines 45, 49, 63). The filter has further mask to narrow down the match to last 4 bits (lines 33, 47, 61) within the byte selected by flexible match mask. And, we are looking for specific values – 0xa, 0xb and 0xc – in these 4 bits (lines 34, 48, 62) to assign packets to specific slices (lines 39, 53, 67). Similar to MPLS filter, this IPv6 filter is applied as input filter on core interfaces (lines 1-26).

Before performing any verification, let's summarize current configuration state:

- Slice-aware H-QoS (link slicing) configured on inter-AS link, with three slices having different min/max bandwidth constraints, each slice with two forwarding classes
- 9 VRFs on PE routers, 3 VRFs per slice. Within each slice, 2 VRFs use SRv6 as underlay and 1 VRF uses MPLS as underlay
- Traffic generators send traffic with traffic rate 2 Mbps per VRF with two forwarding classes (1 Mbps per forwarding class)
- So, we have in total 6 Mbps per slice (3 Mbps per forwarding class in each slice)
- Each slice on inter-AS link has mixed SRv6 and MPLS traffic

Now, let's check the statistics for both MPLS and SRv6 based slice classification (CLI-Output 6).

```

1 1 kszarkowicz@P1> show firewall
2 2
3 3 Filter: FF-SRV6-SLICE-CLASSIFIER
4 4 Counters:
5 5 Name Bytes Packets
6 6 CT-NON-SLICED 1760584 22535
7 7 CT-SLICE-A 175950390 135555
8 8 CT-SLICE-B 175949092 135554
9 9 CT-SLICE-C 175949092 135554
10 10
11 11 Filter: FF-MPLS-SLICE-CLASSIFIER
12 12 Counters:
13 13 Name Bytes Packets
14 14 CT-NON-SLICED 0 0
15 15 CT-SLICE-A 3189241368 2519148
16 16 CT-SLICE-B 3189240102 2519147
17 17 CT-SLICE-C 3189240102 2519147

```

CLI-Output 6: Slice classification based on SRv6 slice ID

We can observe that both MPLS and IPv6 (SRv6) traffic is assigned to the slices. For IPv6 traffic we have small amount not assigned to any particular slice – this is control traffic (e.g. BGP), which is not matched explicitly in the firewall filter. The default slice has by no guarantees by default, so in the real-life deployment, it is recommended to assign control traffic to a separate ‘control plane’ slice, with some guarantees (1-5% of link capacity) to avoid suppression of control plane traffic by other slices. Alternatively, some guarantees can be provided to the default slice, by attaching appropriate traffic control profile to remaining queues on the interface (‘set class-of-service interfaces output-traffic-control-profile-remaining’)

Now, checking the slice queue statistics we can make interesting observations (CLI-Output 7).

```

1 1 kszarkowicz@P1> show interfaces queue ge-0/0/3 slice NS-A
2 2 Slice : NS-A (Index : 1)
3 3 Anchor interface : ge-0/0/3 (Index : 152)
4 4 Forwarding classes: 16 supported, 4 in use
5 5 Egress queues: 8 supported, 4 in use
6 6 Queue: 0, Forwarding classes: FC-BE
7 7 Queued:
8 8 Packets : 2685763 293 pps
9 9 Bytes : 3533512904 3110656 bps
10 10 Transmitted:
11 11 Packets : 2685763 293 pps
12 12 Bytes : 3533512904 3110656 bps
13 13 Tail-dropped packets : 0 0 pps
14 14 RL-dropped packets : 0 0 pps
15 15 RL-dropped bytes : 0 0 bps
16 16 RED-dropped packets : 0 0 pps
17 17 Low : 0 0 pps
18 18 Medium-low : 0 0 pps
19 19 Medium-high : 0 0 pps
20 20 High : 0 0 pps
21 21 RED-dropped bytes : 0 0 bps
22 22 Low : 0 0 bps
23 23 Medium-low : 0 0 bps
24 24 Medium-high : 0 0 bps
25 25 High : 0 0 bps
26 26 Queue: 1, Forwarding classes: FC-EF

```

```
27 27 Queued:
28 28 Packets : 2685761 292 pps
29 29 Bytes : 3533510232 3101056 bps
30 30 Transmitted:
31 31 Packets : 2685761 292 pps
32 32 Bytes : 3533510232 3101056 bps
33 33 Tail-dropped packets : 0 0 pps
34 34 RL-dropped packets : 0 0 pps
35 35 RL-dropped bytes : 0 0 bps
36 36 RED-dropped packets : 0 0 pps
37 37 Low : 0 0 pps
38 38 Medium-low : 0 0 pps
39 39 Medium-high : 0 0 pps
40 40 High : 0 0 pps
41 41 RED-dropped bytes : 0 0 bps
42 42 Low : 0 0 bps
43 43 Medium-low : 0 0 bps
44 44 Medium-high : 0 0 bps
45 45 High : 0 0 bps
46 46 Queue: 2, Forwarding classes: assured-forwarding
47 47 Queued:
48 48 Packets : 0 0 pps
49 49 Bytes : 0 0 bps
50 50 Transmitted:
51 51 Packets : 0 0 pps
52 52 Bytes : 0 0 bps
53 53 Tail-dropped packets : 0 0 pps
54 54 RL-dropped packets : 0 0 pps
55 55 RL-dropped bytes : 0 0 bps
56 56 RED-dropped packets : 0 0 pps
57 57 Low : 0 0 pps
58 58 Medium-low : 0 0 pps
59 59 Medium-high : 0 0 pps
60 60 High : 0 0 pps
61 61 RED-dropped bytes : 0 0 bps
62 62 Low : 0 0 bps
63 63 Medium-low : 0 0 bps
64 64 Medium-high : 0 0 bps
65 65 High : 0 0 bps
66 66
67 67 (omitted for brevity)
68 68
69 69
70 70 kszarkowicz@P1> show interfaces queue ge-0/0/3 slice NS-B
71 71 Slice : NS-B (Index : 2)
72 72 Anchor interface : ge-0/0/3 (Index : 152)
73 73 Forwarding classes: 16 supported, 4 in use
74 74 Egress queues: 8 supported, 4 in use
75 75 Queue: 0, Forwarding classes: FC-BE
76 76 Queued:
77 77 Packets : 2686710 292 pps
78 78 Bytes : 3534768528 3106816 bps
79 79 Transmitted:
80 80 Packets : 2396074 235 pps
81 81 Bytes : 3155297008 2502272 bps
82 82 Tail-dropped packets : 290636 57 pps
83 83 RL-dropped packets : 0 0 pps
84 84 RL-dropped bytes : 0 0 bps
85 85 RED-dropped packets : 0 0 pps
86 86 Low : 0 0 pps
87 87 Medium-low : 0 0 pps
88 88 Medium-high : 0 0 pps
89 89 High : 0 0 pps
90 90 RED-dropped bytes : 0 0 bps
91 91 Low : 0 0 bps
92 92 Medium-low : 0 0 bps
93 93 Medium-high : 0 0 bps
94 94 High : 0 0 bps
95 95 Queue: 1, Forwarding classes: FC-EF
96 96 Queued:
97 97 Packets : 2686710 292 pps
98 98 Bytes : 3534768528 3107712 bps
99 99 Transmitted:
100 100 Packets : 2684008 292 pps
101 101 Bytes : 3531240224 3104640 bps
102 102 Tail-dropped packets : 0 0 pps
103 103 RL-dropped packets : 0 0 pps
104 104 RL-dropped bytes : 0 0 bps
105 105 RED-dropped packets : 0 0 pps
106 106 Low : 0 0 pps
107 107 Medium-low : 0 0 pps
```



```

108 108 Medium-high : 0 0 pps
109 109 High : 0 0 pps
110 110 RED-dropped bytes : 0 0 bps
111 111 Low : 0 0 bps
112 112 Medium-low : 0 0 bps
113 113 Medium-high : 0 0 bps
114 114 High : 0 0 bps
115 115
116 116 (omitted for brevity)
117 117
118 118
119 119 kszarkowicz@P1> show interfaces queue ge-0/0/3 slice NS-C
120 120 Slice : NS-C (Index : 3)
121 121 Anchor interface : ge-0/0/3 (Index : 152)
122 122 Forwarding classes: 16 supported, 4 in use
123 123 Egress queues: 8 supported, 4 in use
124 124 Queue: 0, Forwarding classes: FC-BE
125 125 Queued:
126 126 Packets : 2690769 292 pps
127 127 Bytes : 3540148056 3106176 bps
128 128 Transmitted:
129 129 Packets : 2160508 187 pps
130 130 Bytes : 2846105824 2004608 bps
131 131 Tail-dropped packets : 530261 105 pps
132 132 RL-dropped packets : 0 0 pps
133 133 RL-dropped bytes : 0 0 bps
134 134 RED-dropped packets : 0 0 pps
135 135 Low : 0 0 pps
136 136 Medium-low : 0 0 pps
137 137 Medium-high : 0 0 pps
138 138 High : 0 0 pps
139 139 RED-dropped bytes : 0 0 bps
140 140 Low : 0 0 bps
141 141 Medium-low : 0 0 bps
142 142 Medium-high : 0 0 bps
143 143 High : 0 0 bps
144 144 Queue: 1, Forwarding classes: FC-EF
145 145 Queued:
146 146 Packets : 2690768 293 pps
147 147 Bytes : 3540146752 3115648 bps
148 148 Transmitted:
149 149 Packets : 2687784 293 pps
150 150 Bytes : 3536251520 3112576 bps
151 151 Tail-dropped packets : 0 0 pps
152 152 RL-dropped packets : 0 0 pps
153 153 RL-dropped bytes : 0 0 bps
154 154 RED-dropped packets : 0 0 pps
155 155 Low : 0 0 pps
156 156 Medium-low : 0 0 pps
157 157 Medium-high : 0 0 pps
158 158 High : 0 0 pps
159 159 RED-dropped bytes : 0 0 bps
160 160 Low : 0 0 bps
161 161 Medium-low : 0 0 bps
162 162 Medium-high : 0 0 bps
163 163 High : 0 0 bps
164 164
165 165 (omitted for brevity)

```

CLI-Output 7: Verification of slice-aware H-QoS

There are no drops in slice NS-A. In slice NS-B we observe small drops in forwarding class FC-BE. In slice NS-C we observe even bigger drops, but again in forwarding class FC-BE only. The observations are in line with the expectations (please refer to Table 1 for slice rates). Also, for slices observing drops, only FE-BE class, with low priority, is affected. Strict-priority FC-EF class is not affected.

1.7 Next Steps

In the next blog post we will show another slicing use case with lookup entry assigning packets to slices, instead of firewall filter based assignment.

1.8 Useful links

- Firewall Filter Flexible Match Conditions: <https://www.juniper.net/documentation/us/en/software/junos/routing-policy/topics/concept/firewall-filter-flexible-match-conditions-overview.html>
- Slice (firewall filter action): [https://www.juniper.net/documentation/us/en/software/junos/cos-hierarchical/cos/topics/ref/statement/slice\(firewallfilteraction\).html](https://www.juniper.net/documentation/us/en/software/junos/cos-hierarchical/cos/topics/ref/statement/slice(firewallfilteraction).html)
- Slice-aware H-QoS: <https://www.juniper.net/documentation/us/en/software/junos/cos-hierarchical/cos/topics/topic-map/HierarchicalCoSqueuingforsliceper-hop-behavior.html>
- TechPost 1: SRv6 Basics Locator and End-SIDs - <https://juniper.github.io/techposts/srv6-basics-locator-and-end-sids/article>
- TechPost 2: L3VPN on SRv6 - <https://juniper.github.io/techposts/l3vpn-over-srv6/article>
- TechPost 3: SRv6 Summarisation - <https://juniper.github.io/techposts/srv6-summarization/article>
- TechPost 4: SRv6 SID Encoding and Transposition - <https://juniper.github.io/techposts/srv6-sid-encoding-and-transposition/article>
- TechPost 5: SRv6 L3VPN Inter-AS Option-C - <https://juniper.github.io/techposts/srv6-l3vpn-inter-as-option-c/article>
- TechPost 6: Link Slicing with MPLS and SRv6 Underlays - <https://juniper.github.io/techposts/link-slicing-with-mpls-and-srv6-underlays/article>
- TechPost 7: Link Slicing with MPLS and SRv6 Underlays Part 2 - <https://juniper.github.io/techposts/link-slicing-with-mpls-and-srv6-underlays-part-2/article>

1.9 Glossary

- AS: Autonomous System
- ASBR: Autonomous System Boundary Router
- BGP: Border Gateway Protocol
- BW: Bandwidth
- CE: Customer Edge
- CIR: Committed Information Rate
- CLI: Command Line Interface
- DSCP: Differentiated Services Code Point
- Dst: Destination
- eBGP: external Border Gateway Protocol
- ECN: Explicit Congestion Notification
- FlexE: Flexible Ethernet
- Gbps: Gigabits per second
- H-QoS: Hierarchical Quality of Services
- iBGP: internal Border Gateway Protocol
- ID: Identifier
- IGP: Interior Gateway Protocol
- Inter-AS: Inter Autonomous System
- IP: Internet Protocol
- IPv4: Internet Protocol version 4
- IPv6: Internet Protocol version 6
- IS-IS: Intermediate System to Intermediate System

- L2: Level 2
- L3VPN: Layer 3 Virtual Private Network
- Mbps: Megabits per second
- MPLS: Multiprotocol Label Switching
- NHS: Next Hop Self
- OSPF: Open Shortest Path First
- P: Provider
- PE: Provider Edge
- PIR: Peak Information Rate
- PSP: Penultimate Segment Pop
- QoS: Quality of Services
- RFC: Request for Comments
- RIB: Routing Information Base
- RR: Route Reflector
- SID: Segment Identifier
- SR: Segment Routing
- Src: Source
- SR-MPLS: Segment Routing with Multiprotocol Label Switching
- SR-TE: Segment Routing Traffic Engineering
- SRv6: Segment Routing version 6
- TC: Traffic Class
- TI-LFA: Topology Independent Loop Free Alternates
- TLV: Type Length Value
- USD: Ultimate Segment Decapsulation
- USP: Ultimate Segment Pop
- VLAN: Virtual Local Area Network
- VPN: Virtual Private Network
- VRF: Virtual Routing and Forwarding

1.10 Acknowledgements

Many thanks to Anton Elita for his thorough review and suggestions, and Aditya T R for preparing JCL and vLabs topologies.



DC – AI – Routing – Switching – Security – Wireless – Automation

© Copyright 2026 Hewlett Packard Enterprise Development LP. The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty.

Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Trademark acknowledgments, if needed. All third-party marks are the property of their respective owners.

HEWLETT PACKARD ENTERPRISE

HPE.com

