



DC – AI – Routing – Switching – Security – Wireless – Automation

IPv6 Observability

An HPE Juniper Networking Techpost

David Roy

2024-11-18

HPE



Contents

1	SRv6 Observability	3
1.1	Introduction	3
1.2	Any Standard Exist?	3
1.3	One Juniper Solution	4
1.4	Illustration	5
1.5	Conclusion	12
1.6	Useful links	13
1.7	Glossary	13

List of Figures

1	SRv6 Encapsulation Layers	4
2	SRv6-Enabled Network	5
3	Typical OpenSource based solution for Flows management	8
4	SRv6 DT4 and DT6 IMON Flow Exports Decoding	11
5	SRv6 TE LSP conveying IPv4 decoded by Goflow2	12

1 SRv6 Observability

David Roy - 11/18/2024

How can we monitor the SRv6 data plane, and collect statistics on the SRv6 SRH and tunnels with IPFIX option 315 / IMON?

1.1 Introduction

At Juniper, we have observed increasing customer interest in the SRv6 standard over recent years. This Techpost isn't focused on the SRv6 technology itself. It's something we've already explored in-depth in multiple previous articles:

- SRv6 Basics: Locator and End SIDs: <https://juniper.github.io/techposts/srv6-basics-locator-and-end-sids/article>
- L3VPN over SRv6: <https://juniper.github.io/techposts/l3vpn-over-srv6/article>
- SRv6 Summarization: <https://juniper.github.io/techposts/srv6-summarization/article>
- SRv6 SID Encoding and Transposition: <https://juniper.github.io/techposts/srv6-sid-encoding-and-transposition/article>
- SRv6 L3VPN Inter-AS Option-C: <https://juniper.github.io/techposts/srv6-l3vpn-inter-as-option-c/article>
- Link Slicing with MPLS and SRv6 Underlays: <https://juniper.github.io/techposts/link-slicing-with-mpls-and-srv6-underlays/article>
- Link Slicing with MPLS and SRv6 Underlays Part 2: <https://juniper.github.io/techposts/link-slicing-with-mpls-and-srv6-underlays-part-2/article>
- SRv6 in PTX Express 5: <https://juniper.github.io/techposts/srv6-in-ptx-express-5/article>

Today, we'll discuss SRv6 Observability and we'll try to answer these questions.

- How can we monitor the SRv6 data plane in SRv6-enabled networks?
- Can we collect statistics on the SRv6 Routing Header (aka. SRH) to enable comprehensive network visibility?
- Can we go even further and gather statistics within the SRv6 tunnel, specifically, from the inner frame?

1.2 Any Standard Exist?

Yes... But not entirely. Here's why: At the edge of an SRv6-enabled network, on PE/CE interfaces or peering/transit-facing interfaces, classical, standard IPFIX solutions can monitor IPv6 or IPv4 traffic before encapsulation. In other words, standard or vendor-extended IPv4 and IPv6 IPFIX templates are applicable here. But what about within the core, on P routers, or core-facing interfaces of PE routers?

A practical answer would be that, since SRv6 uses IPv6 as underlay, we could enable a standard IPv6 IPFIX template to collect statistics for the "outer" IPv6 header. This approach might be sufficient in some cases, especially with SRv6 uSID [2], which has significantly enhanced the SRv6 Network Programming Model, with uSID, the IPv6 destination address of the outer frame no longer represents a single "hop," and, for instance, the SRH isn't always required for TE. However, the SRH is still present in some cases and may carry valuable information.

The following figure simply illustrates the “outer header”; “inner frame”, “SRH” terminologies we have just discussed above:




Figure 1: SRv6 Encapsulation Layers

To address the SRH case, a recent RFC, RFC 9487 [3], has defined a new IPFIX template for exposing some additional data. The table below lists the eleven new IPFIX Information Elements (IEs) specified by this RFC:

ElementID	Name	Localization of the information
492	srhFlagsIPv6	In the packet
493	srhTagIPv6	In the packet
494	srhSegmentIPv6	In the packet
495	srhActiveSegmentIPv6	In the packet
496	srhSegmentIPv6BasicList	In the packet
497	srhSegmentIPv6ListSection	In the packet
498	srhSegmentsIPv6Left	In the packet
499	srhIPv6Section	In the packet
500	srhIPv6ActiveSegmentType	Local router information
501	srhSegmentIPv6LocatorLength	Local router information
502	srhSegmentIPv6EndpointBehavior	Local router information

1.3 One Juniper Solution

To address our goals, outlined in the introduction, Juniper offers a solution that meets all our requirements: Inline Monitoring, or IMON, also referred to as IPFIX Information Elements ID 315. This is a familiar topic, as

we've discussed this flexible and highly scalable solution several times before (<https://juniper.github.io/techposts/from-sflow-to-imon-sampling-on-mx10k-platforms/article>).

In short, IMON is a sampling solution implemented in hardware on TRIO and EXPRESS chipsets. Thanks to IPFIX, it exposes the first bytes of layer-2 frames along with some additional metadata. No aggregation or enrichment is done on the device itself; these functions are handled externally, typically by the IPFIX collector or a separate off-box software solution.

IMON on TRIO ASICs reveals the first 126 bytes of each packet, while on EXPRESS ASICs, it reveals the first 256 bytes.

So, why not use IMON as a solution for gathering data-plane statistics on SRv6-enabled networks for:

- The outer IPv6 header
- The SRH, when present
- The inner frame.

1.4 Illustration

To demonstrate the capabilities of IMON in an SRv6-only network, let's set up a simple lab, as shown in Figure 2. The topology is straightforward, featuring two CE devices connected in dual-stack mode to PE routers (in this case, two MX304s). The core network consists of PTX 10K devices that utilize the latest EXPRESS 5 ASIC generation. For this setup, we've used the PTX10002-36QDD. Each router within the SRv6 domain has been assigned one Full-SID and one uSID for testing purposes.




Figure 2: SRv6-Enabled Network

Some details about our topology:

Each PE has a basic dual-stack VRF supporting DT4 and DT6 functions. eBGP service routes received from the CEs are shared through MP-iBGP using a Route Reflector. As shown below, PE1 sent two BGP routes to the RR—one for IPv4 and another for IPv6. Both routes share the same BGP Prefix-SID, as we configured the “end-dt46” function for this VRF.

```

1 bob@pe1> show route advertising-protocol bgp <IPv6-OF-RR> detail
2
3 VRF-DT46.inet.0: 4 destinations, 4 routes (4 active, 0 holddown, 0 hidden)
4 * 10.0.0.0/24 (1 entry, 1 announced)
5   BGP group MESH-VPN type Internal
6     Route Distinguisher: 172.16.255.1:65000
7     VPN Label: 524288
8     Nexthop: Self
9     Flags: Nexthop Change
10    Localpref: 100
11    AS path: [65000] 65001 I
12    Communities: target:65000:1
13      SRv6 SID: 2001:c0c0:2:2:: Service tlv type: 5 Behavior: 20 BL: 48 NL: 16 FL: 16 AL:
14      0 TL: 16 TO: 64
15
16 VRF-DT46.inet6.0: 6 destinations, 6 routes (6 active, 0 holddown, 0 hidden)
17 * 3000:1:1:1::/64 (1 entry, 1 announced)
18   BGP group MESH-VPN type Internal
19     Route Distinguisher: 172.16.255.1:65000
20     VPN Label: 524288
21     Nexthop: Self
22     Flags: Nexthop Change
23     Localpref: 100
24     AS path: [65000] 65001 I
25     Communities: target:65000:1
26       SRv6 SID: 2001:c0c0:2:2:: Service tlv type: 5 Behavior: 20 BL: 48 NL: 16 FL: 16 AL:
27       0 TL: 16 TO: 64

```

We also set up an SRv6 TE LSP between PE1 and PE2. When this LSP is active, PE1 utilizes a constrained LSP built from a stack of Full-SIDs, which requires the use of SRH. If the TE LSP is disabled, PE1 follows the shortest path to reach PE2 (using the micro-end-sid of PE2) and doesn't need to insert an SRH between the outer header and the inner frame. This setup allows us to demonstrate both scenarios, with and without the SRH.

Finally, the configuration of the P routers is much simpler. We have just enabled SRv6 and TI-FLA. Only one DT4 admin VRF is also set up to send Inline Monitoring IPFIX flow exports to our remote collector, which is also part of the admin VRF and connected to a remote PE. Naturally, on P1, we've enabled the IMON feature for the inet6 family. For this example, we've set a sampling rate of 1 to track any packets.

Below are the key parts of P1's configuration:

```

1 lab@ptx-core> show configuration interfaces lo0 unit 1
2 family inet {
3   address 192.168.7.1/32;
4 }
5
6 lab@ptx-core> show configuration routing-instances ADMIN
7 instance-type vrf;
8 protocols {
9   bgp {
10     source-packet-routing {
11       srv6 {
12         locator SID {
13           end-dt4-sid;
14         }
15       }
16     }
17   }
18 }
19 interface lo0.1;
20 route-distinguisher 172.16.255.3:555;
21 vrf-target target:65000:555;
22
23 lab@ptx-core> show configuration services inline-monitoring
24 template {
25   srv6-template {
26     template-refresh-rate 10;
27     option-template-refresh-rate 20;
28     observation-domain-id 10;
29   }
30 }
31 instance {

```

```
32  SRv6 {
33      template-name srv6-template;
34      maximum-clip-length 256;
35      sampling-rate 1;
36      collector {
37          my-collector {
38              source-address 192.168.7.1;
39              destination-address 192.168.1.11;
40              destination-port 9999;
41              routing-instance ADMIN;
42          }
43      }
44  }
45 }
46
47 lab@ptx-core> show configuration interfaces et-0/0/0
48 description TO_PE1;
49 mtu 9200;
50 unit 0 {
51     family iso;
52     family inet6 {
53         filter {
54             input SRV6_MONITOR;
55             output SRV6_MONITOR;
56         }
57     }
58 }
59
60 lab@ptx-core> show configuration firewall family inet6
61 filter SRV6_MONITOR {
62     interface-specific;
63     term 1 {
64         then {
65             count ALL;
66             inline-monitoring-instance SRv6;
67             accept;
68         }
69     }
70 }
```

If you're still following along, you may be wondering:

Ok! With the configuration above, I should be able to export the first 256 bytes of transit packets, but how do I decode them? Precisely, how can I interpret the fields of the outer IPv6 header, the SRH when present, and also extract some specific fields from the inner frame payload?

To answer that, I recently contributed to one of the most flexible and scalable open-source flow collection tools: Goflow2 [5] <https://github.com/netsampler/goflow2>. In detail, I added support for decoding SRH and inner frame fields in SRv6 context.

Goflow2 serves as a versatile tool for flow monitoring. It supports several versions of NetFlow, sFlow, IPFIX, and, of course, IPFIX IE 315 (IMON). Among its outputs, it provides a reliable and efficient ProtoBuf export over Kafka, which facilitates resiliency, load balancing, data replication, and simplifies further enrichment or other post-processing tasks. The following figure illustrates a typical open-source architecture for decoding IPFIX/IMON with Goflow2 and ingesting high-throughput flow exports in a “flows” database such as a ClickHouse.




Figure 3: Typical OpenSource based solution for Flows management

My recent contribution has been committed and merged into the latest Goflow2 build, so let's try it out. First of all, disable the TE LSP, allowing PE1 and PE2 to use only the shortest path, then generate two pings from CE1 to CE2—one for IPv4 and one for IPv6.

Goflow2 uses a “mapping” file to specify which fields to decode from raw packets and present as “agnostic flows” (see Figure 3). For our setup, we created a YAML mapping file to decode the outer IPv6 source and destination addresses, all SRH fields (if present), and specific fields from the inner frame: IP fields (v4 or v6) and transport fields.

```

1  formatter:
2    fields:
3      - type
4      - time_received_ns
5      - sequence_num
6      - sampling_rate
7      - sampler_address
8      - bytes
9      - packets
10     - src_addr
11     - dst_addr
12     - proto
13     - in_if
14     - out_if
15     - layer_stack
16     - layer_size
17     # srv6 fields
18     - ipv6_routing_header_seg_left
19     - srhLastEntryIPv6
20     - srhFlagsIPv6
21     - srhTagIPv6
22     - ipv6_routing_header_addresses
23     # inner frame
24     - innerFrame_src_addr
25     - innerFrame_dst_addr
26     - innerFrame_proto
27     - innerFrame_src_port
28     - innerFrame_dst_port
29     - innerFrame_icmp_type
30     - innerFrame_icmp_code
31   key:
32     - sampler_address
33   protobuf:
34     # srv6 fields
35     - name: srhLastEntryIPv6
36       index: 151
37       type: varint

```

```

38 - name: srhFlagsIPv6
39   index: 152
40   type: varint
41 - name: srhTagIPv6
42   index: 153
43   type: varint
44 # inner frame
45 - name: innerFrame_src_addr
46   index: 160
47   type: string
48 - name: innerFrame_dst_addr
49   index: 161
50   type: string
51 - name: innerFrame_proto
52   index: 162
53   type: varint
54 - name: innerFrame_src_port
55   index: 163
56   type: varint
57 - name: innerFrame_dst_port
58   index: 164
59   type: varint
60 # icmp
61 - name: innerFrame_icmp_type
62   index: 172
63   type: varint
64 - name: innerFrame_icmp_code
65   index: 173
66   type: varint
67 rename:
68   ipv6_routing_header_addresses: srhSegmentIPv6BasicList
69   ipv6_routing_header_seg_left: srhSegmentsIPv6Left
70 render:
71   innerFrame_src_addr: ip
72   innerFrame_dst_addr: ip
73   innerFrame_proto: proto
74 sflow:
75   mapping:
76     # srv6
77     - layer: "ipv6eh_routing"
78       offset: 32
79       length: 8
80       destination: srhLastEntryIPv6
81     - layer: "ipv6eh_routing"
82       offset: 40
83       length: 8
84       destination: srhFlagsIPv6
85     - layer: "ipv6eh_routing"
86       offset: 48
87       length: 16
88       destination: srhTagIPv6
89     # src/dst addresses
90     - layer: "ipv6"
91       encap: true
92       offset: 64
93       length: 128
94       destination: innerFrame_src_addr
95     - layer: "ipv6"
96       encap: true
97       offset: 192
98       length: 128
99       destination: innerFrame_dst_addr
100    - layer: "ipv4"
101      encap: true
102      offset: 96
103      length: 32
104      destination: innerFrame_src_addr
105    - layer: "ipv4"
106      encap: true
107      offset: 128
108      length: 32
109      destination: innerFrame_dst_addr
110    # proto
111    - layer: "ipv6"
112      encap: true
113      offset: 48
114      length: 8
115      destination: innerFrame_proto
116    - layer: "ipv4"
117      encap: true
118      offset: 72

```

```
119     length: 8
120     destination: innerFrame_proto
121     # ports
122     - layer: "udp"
123       encap: true
124       offset: 0
125       length: 16
126       destination: innerFrame_src_port
127     - layer: "udp"
128       encap: true
129       offset: 16
130       length: 16
131       destination: innerFrame_dst_port
132     - layer: "tcp"
133       encap: true
134       offset: 0
135       length: 16
136       destination: innerFrame_src_port
137     - layer: "tcp"
138       encap: true
139       offset: 16
140       length: 16
141       destination: innerFrame_dst_port
142     # icmp
143     - layer: "icmp"
144       encap: true
145       offset: 0
146       length: 8
147       destination: innerFrame_icmp_type
148     - layer: "icmp"
149       encap: true
150       offset: 8
151       length: 8
152       destination: innerFrame_icmp_code
```

Now, let's start Goflow2 with these parameters. For simplicity, I set up Goflow2 to export "agnostic flows" in JSON format and display them directly in the terminal. Keep in mind that these "agnostic flows" are the decoded results of IPFIX/IMON export packets coming from our PTX10002.

```
1 # I use jq for decoding json in the console
2 ./goflow2 -mapping=mapping.yml -listen=netflow://:9999 | jq
```

The figure below displays the decoded fields of both the outer and inner frames. As noted earlier, we have disabled SR-TE, so in this case, DT4 and DT6 traffic does not require an SRH.




Figure 4: SRv6 DT4 and DT6 IMON Flow Exports Decoding

To illustrate once again the powerful combination of Juniper's IMON implementation and the flexibility of Goflow2, we re-enabled our SRv6 TE LSP and sent another IPv4 ping from CE1 to CE2. In this scenario, there should be an additional SRH IPv6 header between the outer and inner frame headers, which the Goflow2 export successfully confirms.

SRv6 TE DT4

Outer header's fields

Full SRH decoding

Inner's fields

```
{
  "type": "IPFIX",
  "time_received_ns": 1730883452623452000,
  "sequence_num": 219215,
  "sampling_rate": 1,
  "sampler_address": "192.168.7.1",
  "bytes": 178,
  "packets": 1,
  "src_addr": "2001:cafe::2",
  "dst_addr": "2001:c0c0:1:1::",
  "proto": "IPv6-Route",
  "in_if": 589,
  "out_if": 0,
  "layer_stack": [
    "Ethernet",
    "IPv6",
    "IPv6HeaderRouting",
    "IPv4",
    "ICMP"
  ],
  "layer_size": [
    14,
    40,
    40,
    20,
    8
  ],
  "srhSegmentsIPv6Left": 1,
  "srhLastEntryIPv6": 1,
  "srhFlagsIPv6": 0,
  "srhTagIPv6": 0,
  "srhSegmentIPv6BasicList": [
    "2001:c0c0:3:3:8000::",
    "2001:c0c0:1:1::"
  ],
  "innerFrame_src_addr": "10.0.0.1",
  "innerFrame_dst_addr": "20.0.0.1",
  "innerFrame_proto": "ICMP",
  "innerFrame_icmp_type": 8,
  "innerFrame_icmp_code": 0
}
```

Figure 5: SRv6 TE LSP conveying IPv4 decoded by Goflow2

1.5 Conclusion

In this Techpost, we've demonstrated the power of the IMON feature on Juniper routing devices with TRIO and EXPRESS chipsets. By combining the inline implementation of IPFIX IE 315 with a scalable off-box collector like Goflow2, customers can quickly adapt and expose statistics for any new encapsulations or underlays—without waiting for the implementation or specification of a new standard. This is precisely what we've demonstrated today with SRv6.

I'd also like to emphasize the flexibility of Goflow2 and its new custom mapping feature, which provides a robust and agile solution for data plane observability. In conclusion, I'd like to send my special thanks to the Goflow2 maintainer for accepting and committing my recent contribution.

1.6 Useful links

- [1] <https://juniper.github.io/techposts/srv6-in-ptx-express-5/article>
- [2] <https://juniper.github.io/techposts/introducing-ptx10002-36qdd/article>
- [3] <https://datatracker.ietf.org/doc/rfc9487/>
- [4] <https://juniper.github.io/techposts/from-sflow-to-imon-sampling-on-mx10k-platforms/article>
- [5] <https://github.com/netsampler/goflow2>

1.7 Glossary

- ASIC: Application-Specific Integrated Circuit
- DT4: Decapsulation and Transit IPv4
- DT6: Decapsulation and Transit IPv6
- IE: Information Element
- IMON: Inline MONitoring
- IPFIX: IP Flow Information Export
- JSON: JavaScript Object Notation
- LSP: Label Switching Path
- SRH: Segment Routing Header
- SID: Segment IDentifier
- SR: Segment Routing
- uSID: micro Segment IDentifier
- TE: Traffic Engineering



DC – AI – Routing – Switching – Security – Wireless – Automation

© Copyright 2026 Hewlett Packard Enterprise Development LP. The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty.

Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Trademark acknowledgments, if needed. All third-party marks are the property of their respective owners.

HEWLETT PACKARD ENTERPRISE

HPE.com

