



DC – AI – Routing – Switching – Security – Wireless – Automation

SRX MNHA - VRRP

An HPE Juniper Networking Techpost

Karel Hendrych

2026-05-27

HPE



Contents

1	SRX MNHA: VRRP	3
1.1	Abstract	3
1.2	Introduction	4
1.3	Design Principles	4
1.4	Demo Topology	5
1.5	Configuration Break-down	8
1.5.1	SRX1600-1	8
1.5.2	SRX1600-2	16
1.6	Validation	18
1.6.1	MNHA/VRRP Status	19
1.6.2	Fail-over	21
1.6.3	Stability Under Overload	25
1.7	Appendix 1: Complete Configurations	26
1.7.1	SRX1600-1	26
1.7.2	SRX1600-2	30
1.7.3	vSRX-1	34
1.8	Appendix 2: Encrypted ICL	35
1.9	Conclusion	37
1.10	Acknowledgements	37
1.11	Useful Links	37
1.12	Glossary	38

List of Figures

1	VRRP and SRX	3
2	Simplified topology diagram	6
3	MNHA ICL and ICD links topology	7
4	Different scenarios, different states	19
5	IxNetwork capture from an administrative fail-over	22
6	IxNetwork capture from PFE Failure test	22
7	IxNetwork capture from a test without process-packet-on-backup enabled	23
8	BreakingPoint capture of stateful traffic run	24
9	IxNetwork capture from an administrative fail-over under load	25
10	IxNetwork capture from PFE Failure test under load	25
11	IxNetwork capture from overload test	25

1 SRX MNHA: VRRP

Karel Hendrych - 05/27/2026



Figure 1: VRRP and SRX

1.1 Abstract

This TechPost describes an alternative solution to the native SRX Multi Node High Availability (MNHA) Switching/Hybrid Layer 2 (L2) Virtual IP address (VIP) mechanisms using Virtual Router Redundancy Protocol (VRRP). MNHA combined with VRRP can offer more flexibility and speed up convergence in some failure scenarios. The solution addresses current limitations of the native MNHA L2 VIP mechanism compared with the traditional SRX chassis-cluster technology. The additional configuration should be seen as a worthwhile investment for the flexibility and better convergence times, albeit making the configuration slightly more complex.

1.2 Introduction

Firewall clustering using failover VIP addresses within an L2 domain is the dominant high-availability mechanism in enterprise networks. SRX MNHA, originally designed for L3 clusters, later introduced an L2 VIP mechanism with Switching and Hybrid modes; the latter is becoming increasingly popular because it combines L3 and L2 redundancy which allows MNHA clusters to be enabled without wholesale network change. The SRX4700 and the emerging SRX400 series do not support the traditional SRX chassis-cluster technology, so this TechPost describes an alternative - using VRRP to provide L2 VIP services - that may be applicable in certain scenarios.

As of Junos 25.4R1, the following constraints exist for MNHA IPv4 and IPv6 L2 VIP addressing in Switching and Hybrid modes:

- Support for one IPv4 OR one IPv6 VIP address per logical interface; dualstack (one IPv4 AND one IPv6) is planned for Junos 26.2R1. This is a challenge for configurations that use multiple IP aliases, whether within the same IP prefix or across different prefixes (multiple IP prefixes within the same VLAN).
- For proxy-ARP/NDP, nodes respond individually using their interface MAC addresses. In addition to routing tricks, multiple IP aliases can partially address this.
- MNHA VIPs cannot overlap, even if a VIP is configured in a different Services Redundancy Group (SRG) with the interface bound to a different virtual router.
- MNHA 3- and 4-node formations support only the L3 mechanism; although not covered here, hypothetically VRRP could address the L2 scenario (without signal-routes-based mastership steering as of Junos 25.4R1).

Note: the maximum scale of 32 VIPs per SRG was increased to 2,000 and [release noted](#) in Junos 25.2R1-S1 for both MNHA Switching and Hybrid modes.

This article does not aim to provide any sort of qualification for specific platforms regarding scale and timer settings; that is a matter of qualification against the specific hardware, software, VIP scale, and timer configurations. In lab environments, the SRX1600/4120/4300 series has generally been observed to scale to the upper hundreds of VRRP VIPs with default timers.

Besides the official MNHA [documentation](#), related reading includes:

- [Multi-Node High Availability Basics(<https://juniper.github.io/techposts/dhcp-on-mnha-back-to-basics/article>)] by Steven Jacques
- DHCP on MNHA: [Back to Basics](#)
- [Hybrid MNHA with eBGP](#) and by James Rathbun
- SRX clustering: from [Chassis Cluster to MultiNode High Availability](#) by Laurent Paumelle

1.3 Design Principles

Before diving into the example setup, the design is essentially built on two distinct state machines - SRX-specific MNHA and standards-based Junos VRRP implementation - with the following specifics:

- VRRP tracks the MNHA active-signal route, so VRRP mastership follows the MNHA SRG1 active node during administrative and monitored-object failovers. The tracked route's priority cost is lower than

the VRRP priority, so a backup node that doesn't have the active-signal route installed remains in the VRRP backup state (it would be in the idle state if the priority cost were equal or higher).

- MNHA SRG is in Routing mode, so native MNHA VIP configuration enforced at commit is not necessarily required.
- The MNHA activeness probe (split-brain detection) uses the VRRP VIP as its source IP for ICMP echo-request. Even if ICL links are down, VRRP state influences MNHA's splitbrain detection mechanism.
- The configuration provides L2 VIP redundancy using VRRP and can use MNHA L3 signalling for upstream, making it effectively L2/L3 hybrid (L3 uplink connectivity is not covered in this article).
- MNHA can use granular SRG monitoring to transfer activeness (and VRRP mastership) to the other node if certain criteria are met (e.g., link down, ICMP/BFD monitoring).
- Multiple VRRP IPv4 or IPv6 addresses within the same prefix (IP aliases) are configured as additional VIP entries (which implies VRRP *no-accept-data* disallowing inbound traffic to the VRRP group VIPs).
- Additional VRRP VIPs can also serve for purposes of proxy-ARP/NDP. Note only 8 IPv4 addresses AND 8 IPv6 per VRRP group are supported.
- Additional prefixes are configured in their own VRRP groups, possibly mixing multiple IPs from the same prefix within those groups.
- MNHA is configured with an additional link for asymmetric operations for cases when VRRP masters are split across nodes (this can be largely mitigated by properly configured MNHA interface tracking).
- VPN services can be terminated on VRRP VIP, MNHA references the VIP(s) in configuration.
- An additional VRRP group with a single VIP in the same prefix as the other group(s) can serve for VPN termination (avoiding the mandatory *no-accept-data* setting for multiple VIPs in a single group).
- VRRPv3 is enabled for IPv6 support and fasters timers (including for IPv4).
- VRRP uses the same priority on both nodes to avoid preemption.

Note: the Validation part contains a potentially valuable table of MNHA and VRRP states reflecting above principles settings of the demo topology described later.

So, in summary VRRP which has sub-second convergence, watches MNHA signalling routes, and drives the active interface change via VRRP mechanisms. Convergence can occur in 1 second with appropriate timers. This twinning of these technologies can result in an order of magnitude faster convergence, with the option to expand to Hybrid like L2/L3 HA modes or to move entirely to L3 HA.

1.4 Demo Topology

The purpose of the demo topology is to:

- Conduct fail-over test with multiple IPv4 and IPv6 addresses on interface with running clear-text background traffic.
- Measure the elapsed time of fail-over for both clear-text and IPSEC tunneled traffic for administrative fail-over and sudden device outage.
- Test stability of VRRP when device is subjected to UDP traffic causing high PFE load

The demo topology depicted in a simplified manner below consists of:

- MNHA pair of [SRX1600s](#), each connected in a firewall on a home scenarios to a pair of switches because the SRX1600's upper forwarding capacity is just under 25 Gbps (Broadcom switching chip-to-CPU link limit).

- A vSRX for functional testing of traffic failover for an IPSEC VPN tunnel between the SRX1600 MNHA pair and the vSRX instance.
- A Keysight (IXIA) BreakingPoint tester with 4x10GE for servers and 4x10GE for clients is used to generate stateful traffic; stateless traffic uses a 2x10GE port on a Keysight (IXIA) IxNetwork tester.

From the SRX devices' perspective, the tester client and server prefixes are routed toward the tester IP addresses within the respective client and server VLAN prefixes, except for IPSEC VPN traffic, which uses tunnel interfaces as the next hops. Physically, the testers are connected to the same QFX10k switch as the SRX devices and the vSRX.




Figure 2: Simplified topology diagram

Details of SRX1600 IP addressing (IPv4 /24 and IPv6 /64):

Zone	Interface.VLAN	SRX1600-1 IP	SRX1600-2 IP	VRRP IP(s)
tester-client	ae0.1081	100.64.81.11	100.64.81.12	.10
tester-client	ae0.1081	dead:81::11	dead:81::12	::10
tester-client	ae0.1082	100.64.82.11	100.64.82.12	.10, .20
tester-client	ae0.1082	dead:82::11	dead:82::11	::10 ::20
tester-client	ae0.1083	100.64.83.11	100.64.83.12	.10
tester-client	ae0.1083	100.64.84.11	100.64.84.11	.10 .20
tester-client	ae0.1083	dead:83::11	dead:83::12	::20
tester-client	ae0.1083	dead:84::11	dead:84::12	::10 ::20
tester-client	ae0.1100	100.64.100.11	100.64.100.12	.10
tester-server	ae0.1091	100.64.91.11	100.64.91.12	.10
tester-server	ae0.1091	dead:91::11	dead:91::12	::10
tester-server	ae0.1092	100.64.92.11	100.64.92.12	.10
tester-server	ae0.1092	dead:92::11	dead:92::12	::10
tester-server	ae0.1093	100.64.93.11	100.64.93.12	.10
tester-server	ae0.1093	dead:93::11	dead:93::12	::10
tester-server	ae0.1094	100.64.94.11	100.64.94.12	.10

Zone	Interface.VLAN	SRX1600-1 IP	SRX1600-2 IP	VRRP IP(s)
tester-server	ae0.1094	dead:94::11	dead:94::12	::10
tester-server	ae0.1101	100.64.101.11	100.64.101.12	.10
tester-server-vpn	ae0.1102	100.64.102.11	100.64.102.12	.10
tester-server-vpn	st0.0	N/A	N/A	N/A

Note: the QFX10 IRB interface in VLAN 1081 uses IP address 100.64.81.20 as the destination for the MNHA activeness probe (explained later).

A brief ICL/ICD refresher before proceeding. The Inter-Chassis Link (ICL) serves two purposes: synchronizing Real Time Object (RTO) state - sessions and IPSEC SAs - between nodes, and providing a lightweight control plane for remote node status. The Inter-Chassis Datapath (ICD) handles traffic that arrives asymmetrically, forwarding it across the ICD until the local node can establish its own forwarding path (local breakout). The VRRP/MNHA setup is designed to be Active/Passive (A/P), so asymmetries should not occur, or worst case should be very brief during numerous VRRP mastership transitions; ICD is primarily present for corner cases and for migrations to Hybrid or L3 mode.

Leveraging loopbacks provides both L3 redundancy and deployment flexibility. In this scenario, connectivity for the ICL and ICD links is between loopback interfaces via directly connected 10GE interfaces. Alternatively, if the MNHA deployment for the ICL/ICD communication paths is across distributed over an L3 infrastructure, the same redundancy and flexibility exists. The *xe-0/1/0* link is the preferred path for ICL and *xe-0/1/1* for ICD. If a single link fails, both ICL and ICD merge on the remaining interface - ensured by BGP/BFD settings.

For ICL/ICD MTU requirements: when ICD is configured and traffic is asymmetric, data segments traverse ICD until all content plug-ins detach (e.g., AppID) and local breakout becomes possible. It was observed that UDP encapsulation of a 1,514byte frame containing a 1,448byte TCP segment results on ICD in a 1,606byte frame (encapsulating the TCP segment into a 1,564byte UDP packet). A 2,000byte MTU on the underlying interfaces *xe-0/1/0* and *xe-0/1/1* is a conservative measure.

The ICL/ICD topology diagram below highlights a dedicated routing instance used for the physical and loopback interfaces, both bound to the same zone.




Figure 3: MNHA ICL and ICD links topology

Details of MNHA loopback and physical interlinks addressing:

Zone	Interface/VLAN	SRX1600-1 IP	SRX1600-2 IP
mnha	xe-0/1/0.0	192.168.0.0	192.168.0.1
mnha	xe-0/1/1.0	192.168.0.2	192.168.0.3
mnha	lo0.0	192.168.1.1	192.168.2.1
mnha	lo0.0	192.168.1.2	192.168.2.2

Finally, vSRX-1 IPSEC peer addressing:

Zone	Interface/VLAN	vSRX-1 IP
tester-server-vpn	ge-0/0/0.0	100.64.102.20
tester-server-vpn	st0.0	N/A
tester-server	ge-0/0/1.0	100.64.101.20

1.5 Configuration Break-down

This section describes the configuration for SRX1600-1 in detail; SRX1600-2 is nearly identical, so only differences are highlighted. Essential system settings - management interface, root authentication, and DNS/NT-P/logging - are not covered here but are partially included in the complete configurations for both SRX1600s and vSRX in [Appendix 1](#).

1.5.1 SRX1600-1

1.5.1.1 Common Settings For configuration context: the host-name is configured; then the number of aggregated interfaces and the SRX1600's 25/10/1GE PIC switched to 10/1GE mode for 10GE HA links:

```
1 set system host-name srx1600-1
2 set chassis aggregated-devices ethernet device-count 1
3 set chassis fpc 0 pic 1 pic-mode 1G10G
```

1.5.1.2 MNHA Infrastructure Next, the loopback interface configuration: two addresses are assigned to the same interface because Junos permits only one loopback unit per routing instance. The underlying physical interfaces, with an increased MTU to carry potential ICD traffic, are configured with point-to-point addressing.

```
1 edit interfaces
2
3 set lo0 unit 0 family inet address 192.168.1.1/32
4 set lo0 unit 0 family inet address 192.168.1.2/32
5
6 set interfaces xe-0/1/0 mtu 2000
7 set interfaces xe-0/1/1 mtu 2000
8
9 set xe-0/1/0 unit 0 family inet address 192.168.0.0/31
10 set xe-0/1/1 unit 0 family inet address 192.168.0.2/31
```

The above interfaces are bound to* mnha* zone; inbound host services permit ICMP, BGP, and BFD. The loopback interface is used only for highavailability communication, ICMP for diagnostics, and BFD for ICL monitoring. BFD is theoretically redundant because traffic matches the UDP flow, but its presence is recommended as best practice:

```

1 edit security zones security-zone mnha interfaces
2
3 set xe-0/1/0.0 host-inbound-traffic system-services ping
4 set xe-0/1/0.0 host-inbound-traffic protocols bgp
5 set xe-0/1/0.0 host-inbound-traffic protocols bfd
6
7 set xe-0/1/1.0 host-inbound-traffic system-services ping
8 set xe-0/1/1.0 host-inbound-traffic protocols bgp
9 set xe-0/1/1.0 host-inbound-traffic protocols bfd
10
11 set lo0.0 host-inbound-traffic system-services high-availability
12 set lo0.0 host-inbound-traffic system-services ping
13 set lo0.0 host-inbound-traffic protocols bfd

```

Corresponding intra-zone policy permits traffic between the physical interfaces and the loopback interface:

```

1 edit security policies from-zone mnha to-zone mnha
2
3 set policy mnha match source-address any
4 set policy mnha match destination-address any
5 set policy mnha match application any
6 set policy mnha then permit

```

Note: policy controls can be tightened based on loopback IPs, or by enabling IPSEC for ICL as covered in [Appendix 2](#).

BGP export policies for the preferred ICL link xe-0/1/0 de-preference the ICD link's loopback IP by prepending the local AS number:

```

1 edit policy-options policy-statement export-mnha-icl
2
3 set term icl from interface lo0.0
4 set term icl from route-filter 192.168.1.1/32 exact
5 set term icl then accept
6
7 set term icd from interface lo0.0
8 set term icd from route-filter 192.168.1.2/32 exact
9 set term icd then as-path-prepend 65001
10 set term icd then accept
11
12 set term final then reject

```

Similarly, the preferred ICD xe-0/1/1 link's export policy de-preferences the ICL link by prepending the local AS number:

```

1 edit policy-options policy-statement export-mnha-icd
2
3 set term icl from interface lo0.0
4 set term icl from route-filter 192.168.1.1/32 exact
5 set term icl then as-path-prepend 65001
6 set term icl then accept
7
8 set term icd from interface lo0.0
9 set term icd from route-filter 192.168.1.2/32 exact
10 set term icd then accept
11
12 set term final then reject

```

Finally, as per best practice, the underlying MNHA interfaces are bound to a dedicated routing instance, with BGP peering and BFD configured for a 1.5 s detection time:

```

1 edit routing-instances vr-mnha
2
3 set instance-type virtual-router

```

```

4 set routing-options autonomous-system 65001
5
6 set protocols bgp group mnha peer-as 65002
7 set protocols bgp group mnha bfd-liveness-detection minimum-interval 500
8 set protocols bgp group mnha bfd-liveness-detection multiplier 3
9 set protocols bgp group mnha neighbor 192.168.0.1 export export-mnha-icl
10 set protocols bgp group mnha neighbor 192.168.0.3 export export-mnha-icd
11
12 set interface xe-0/1/0.0
13 set interface xe-0/1/1.0
14 set interface lo0.0

```

1.5.1.3 MNHA The base MNHA settings define the local identifier (*local-id 1*) in the MNHA formation, the local address for the ICL link, and the ICD (forwarding) IP address. Under the peer identifier (*peer-id 2*), the remote ICL and ICD IPs reachable from the local node are configured, along with the interface and routing instance used as the source for both ICL and ICD traffic. BFD is used to track reachability of the remote node (discussed in the next section) for both ICL and ICD (forwarding).

Note: that the ICL and ICD timers are slower (3 x 1000 ms) than the underlying infrastructure timers (3 x 500 ms), so the ICL/ICD links are preserved during underlay convergence:

```

1 edit chassis high-availability
2
3 set local-id 1
4 set local-id local-ip 192.168.1.1
5 set local-id local-forwarding-ip 192.168.1.2
6
7 set peer-id 2 peer-ip 192.168.2.1
8 set peer-id 2 interface lo0.0
9 set peer-id 2 routing-instance vr-mnha
10
11 set peer-id 2 peer-forwarding-ip 192.168.2.2
12 set peer-id 2 peer-forwarding-ip interface lo0.0
13 set peer-id 2 peer-forwarding-ip liveness-detection minimum-interval 1000
14 set peer-id 2 peer-forwarding-ip liveness-detection multiplier 3
15
16 set peer-id 2 liveness-detection minimum-interval 1000
17 set peer-id 2 liveness-detection multiplier 3

```

In the A/P design with VRRP, SRG-1 in **Routing mode** (explicitly set here, but the default) is a **key concept** because it **provides signal routes used by VRRP** for tracking. Also, Routing mode eliminates the commit-time enforcement to configure native MNHA VIPs (unlike Switching or Hybrid modes). Signal routes are instantiated in the routing instance that hosts the ICL and ICD links; otherwise they go to the inet.0 table by default:

```

1 edit chassis high-availability services-redundancy-group 1
2
3 set deployment-type routing
4 set peer-id 2
5
6 set active-signal-route 100.64.0.1
7 set active-signal-route routing-instance vr-mnha
8
9 set backup-signal-route 100.64.0.0
10 set backup-signal-route routing-instance vr-mnha

```

Note: for completeness - Hybrid mode provides signal routes as well; however, it also requires configuration of MNHA native VIPs, which is not wanted.

The activeness probe runs after ICL loss (tracked by multi-hop BFD session), sending ICMP echo-requests toward an IP address on adjacent routing or switching equipment. The node that receives a response remains active; the node that does not remains in backup state. **Note:** the probe source IP is the VRRP VIP, so only the

VRRP master can send it - the backup has no mechanism to initiate probes and will correctly remain backup if only the ICL has failed. Related configuration with probing effectively from *ae0.1081* VRRP VIP:

```
1 edit chassis high-availability services-redundancy-group 1
2
3 set activeness-probe dest-ip 100.64.81.20
4 set activeness-probe dest-ip src-ip 100.64.81.10
5 set activeness-probe dest-ip routing-instance vr-testers
```

The monitoring tracks *ae0* aggregate-interface link status based on *minimum-links* setting (minimum 2 of 4, described later). Loss of the aggregate interface contributes a weight of 100, which triggers the thresholds for the object and the entire SRG, rendering the MNHA node as ineligible. Finally, the activeness priority is used as a tiebreaker during boot - the higher priority becomes active (and the other node in the MNHA pair should have a lower value to enable this behaviour as shown in the SRX1600-2 section):

```
1 edit chassis high-availability services-redundancy-group 1
2
3 set monitor monitor-object interface-1 object-threshold 100
4 set monitor monitor-object interface-1 interface threshold 100
5 set monitor monitor-object interface-1 interface interface-name ae0 weight 100
6 set monitor srg-threshold 100
7
8 set activeness-priority 200
```

1.5.1.4 Interfaces and VRRP First, enable VRRP version 3 for IPv6 support and subsecond advertisements:

```
1 set protocols vrrp version-3
```

The 4-member aggregate interface *ae0* enabled for VLAN tagging, marking interface as down if less than two links are up and periodic fast set LACP:

```
1 edit interfaces
2
3 set xe-0/2/0 ether-options 802.3ad ae0
4 set xe-0/2/1 ether-options 802.3ad ae0
5 set xe-0/2/2 ether-options 802.3ad ae0
6 set xe-0/2/3 ether-options 802.3ad ae0
7
8 set ae0 vlan-tagging
9 set ae0 aggregated-ether-options minimum-links 2
10 set ae0 aggregated-ether-options lacp active
11 set ae0 aggregated-ether-options lacp periodic fast
```

Common VRRP settings for all logical units on *ae0* are defined in a Junos group. The **tracked route is the active-signal Route from the MNHA settings**; if that route is absent from the *vr-mnha* routing instance, the priority (150) is reduced by the priority cost (100), making the node a VRRP backup under normal conditions while the master remains at priority 150. A VRRP fast-interval of 200 ms enables subsecond failover in ideal conditions:

```
1 edit groups vrrp-1 interfaces ae0 unit <*>
2
3 set family inet address <*> vrrp-group <*> priority 150
4 set family inet address <*> vrrp-group <*> fast-interval 200
5 set family inet address <*> vrrp-group <*> accept-data
6 set family inet address <*> vrrp-group <*> track route 100.64.0.1/32 routing-instance vr-mnha
   priority-cost 100
7
8 set family inet6 address <*> vrrp-inet6-group <*> priority 150
9 set family inet6 address <*> vrrp-inet6-group <*> fast-interval 200
10 set family inet6 address <*> vrrp-inet6-group <*> accept-data
11 set family inet6 address <*> vrrp-inet6-group <*> track route 100.64.0.1/32 routing-instance vr-
   mnha priority-cost 100
```

Note: in the real world, LACP timers interact with VRRP intervals because only one interface on the VRRP master is chosen to send advertisements. If LACP detection is slower than an aggressively configured VRRP (for example, the 200 ms fast-interval used here), when the advertising interface fails while other interfaces remain up, the VRRP backup may briefly become master.

The first clientside VLAN unit is a dualstack interface: the VRRP addresses are `.10` and `::10`; SRX1600-1 uses `.11` and `::11`; SRX1600-2 uses `.12` and `::12`:

```
1 edit interfaces ae0
2
3 set unit 1081 vlan-id 1081
4
5 set unit 1081 family inet address 100.64.81.11/24 vrrp-group 1 virtual-address 100.64.81.10
6 set unit 1081 family inet6 address dead:81::11/64 vrrp-inet6-group 1 virtual-inet6-address dead
  :81::10
```

The second client VLAN unit adds IP aliases for VRRP VIPs within the same prefix. Using multiple IP addresses requires *no-accept-data* (for example, the node does not respond to ICMP or to locally terminated VPN traffic):

```
1 edit interfaces ae0
2
3 set unit 1082 vlan-id 1082
4
5 set unit 1082 family inet address 100.64.82.11/24 vrrp-group 1 virtual-address 100.64.82.10
6 set unit 1082 family inet address 100.64.82.11/24 vrrp-group 1 virtual-address 100.64.82.20
7 set unit 1082 family inet address 100.64.82.11/24 vrrp-group 1 no-accept-data
8
9 set unit 1082 family inet6 address dead:82::11/64 vrrp-inet6-group 1 virtual-inet6-address dead
  :82::10
10 set unit 1082 family inet6 address dead:82::11/64 vrrp-inet6-group 1 virtual-inet6-address dead
  :82::20
11 set unit 1082 family inet6 address dead:82::11/64 vrrp-inet6-group 1 no-accept-data
```

Note: a combination with an additional VRRP group using IPs in the same prefix would also work, removing the need to add *no-accept-data*; e.g., suitable for VPN termination.

The third client VLAN unit has additional IP aliases from a different prefix by using another VRRP group:

```
1 edit interfaces ae0
2
3 set unit 1083 vlan-id 1083
4
5 set unit 1083 family inet address 100.64.83.11/24 vrrp-group 1 virtual-address 100.64.83.10
6 set unit 1083 family inet address 100.64.84.11/24 vrrp-group 2 virtual-address 100.64.84.10
7 set unit 1083 family inet address 100.64.84.11/24 vrrp-group 2 virtual-address 100.64.84.20
8 set unit 1083 family inet address 100.64.84.11/24 vrrp-group 2 no-accept-data
9
10 set unit 1083 family inet6 address dead:83::11/64 vrrp-inet6-group 1 virtual-inet6-address dead
  :83::10
11 set unit 1083 family inet6 address dead:84::11/64 vrrp-inet6-group 2 virtual-inet6-address dead
  :84::10
12 set unit 1083 family inet6 address dead:84::11/64 vrrp-inet6-group 2 virtual-inet6-address dead
  :84::20
13 set unit 1083 family inet6 address dead:84::11/64 vrrp-inet6-group 2 no-accept-data
```

Serverside units are kept simple: dualstack for some interfaces and IPv4only for others:

```
1 edit interfaces ae0
2
3 set unit 1091 vlan-id 1091
4 set unit 1091 family inet address 100.64.91.11/24 vrrp-group 1 virtual-address 100.64.91.10
5 set unit 1091 family inet6 address dead:91::11/64 vrrp-inet6-group 1 virtual-inet6-address dead
  :91::10
6
7 set unit 1092 vlan-id 1092
8 set unit 1092 family inet address 100.64.92.11/24 vrrp-group 1 virtual-address 100.64.92.10
```

```

9  set unit 1092 family inet6 address dead:92::11/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :92::10
10
11  set unit 1093 vlan-id 1093
12  set unit 1093 family inet address 100.64.93.11/24 vrrp-group 1 virtual-address 100.64.93.10
13  set unit 1093 family inet6 address dead:93::11/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :93::10
14
15  set unit 1094 vlan-id 1094
16  set unit 1094 family inet address 100.64.94.11/24 vrrp-group 1 virtual-address 100.64.94.10
17  set unit 1094 family inet6 address dead:94::11/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :94::10
18
19  set unit 1100 vlan-id 1100
20  set unit 1100 family inet address 100.64.100.11/24 vrrp-group 1 virtual-address 100.64.100.10
21
22  set unit 1101 vlan-id 1101
23  set unit 1101 family inet address 100.64.101.11/24 vrrp-group 1 virtual-address 100.64.101.10
24
25  set unit 1102 vlan-id 1102
26  set unit 1102 family inet address 100.64.102.11/24 vrrp-group 1 virtual-address 100.64.102.10

```

1.5.1.5 Routing Author's preference is to enclose all traffic interfaces into separate routing instance. Static routing steers prefixes towards tester ports acting from SRX perspective as next-hops for both client and server prefixes:

```

1  edit routing-instances vr-testers
2
3  set instance-type virtual-router
4
5  set interface ae0.1081
6  set interface ae0.1082
7  set interface ae0.1083
8  set interface ae0.1091
9  set interface ae0.1092
10 set interface ae0.1093
11 set interface ae0.1094
12 set interface ae0.1100
13 set interface ae0.1101
14 set interface ae0.1102
15
16 set routing-options static route 1.0.0.0/15 next-hop 100.64.91.1
17 set routing-options static route 1.2.0.0/15 next-hop 100.64.92.1
18 set routing-options static route 1.4.0.0/15 next-hop 100.64.93.1
19 set routing-options static route 1.6.0.0/15 next-hop 100.64.94.1
20 set routing-options static route 10.0.0.0/15 next-hop 100.64.81.1
21 set routing-options static route 10.2.0.0/15 next-hop 100.64.82.1
22 set routing-options static route 10.4.0.0/15 next-hop 100.64.83.1
23 set routing-options static route 10.6.0.0/15 next-hop 100.64.84.1
24 set routing-options static route 10.100.0.0/15 next-hop 100.64.100.1
25 set routing-options static route 10.102.0.0/15 next-hop 100.64.101.1
26
27 set routing-options rib vr-testers.inet6.0 static route beef::0/64 next-hop dead:81::1
28 set routing-options rib vr-testers.inet6.0 static route beef:0:0:1::0/64 next-hop dead:82::1
29 set routing-options rib vr-testers.inet6.0 static route beef:0:0:2::0/64 next-hop dead:83::1
30 set routing-options rib vr-testers.inet6.0 static route beef:0:0:3::0/64 next-hop dead:84::1
31 set routing-options rib vr-testers.inet6.0 static route dead:0:0:0::0/64 next-hop dead:91::1
32 set routing-options rib vr-testers.inet6.0 static route dead:0:0:1::0/64 next-hop dead:92::1
33 set routing-options rib vr-testers.inet6.0 static route dead:0:0:2::0/64 next-hop dead:93::1
34 set routing-options rib vr-testers.inet6.0 static route dead:0:0:3::0/64 next-hop dead:94::1

```

1.5.1.6 Zones Besides the *mnha* infrastructure zone, only three other zones are configured - for servers, clients, and transit (for the IPSEC part). An important setting is to enable VRRP as a host-inbound protocol; otherwise the master's advertisements are not seen and both nodes become VRRP masters:

```

1  edit security zones
2
3  set security-zone tester-client host-inbound-traffic system-services ping
4  set security-zone tester-client host-inbound-traffic protocols vrrp
5  set security-zone tester-client interfaces ae0.1081
6  set security-zone tester-client interfaces ae0.1082

```

```

7 set security-zone tester-client interfaces ae0.1083
8 set security-zone tester-client interfaces ae0.1100
9
10 set security-zone tester-server host-inbound-traffic system-services ping
11 set security-zone tester-server host-inbound-traffic protocols vrrp
12 set security-zone tester-server interfaces ae0.1091
13 set security-zone tester-server interfaces ae0.1092
14 set security-zone tester-server interfaces ae0.1093
15 set security-zone tester-server interfaces ae0.1094
16 set security-zone tester-server interfaces ae0.1101
17
18 set security-zone tester-server-vpn host-inbound-traffic system-services ping
19 set security-zone tester-server-vpn host-inbound-traffic protocols vrrp
20 set security-zone tester-server-vpn interfaces ae0.1102

```

1.5.1.7 Policies Security policies are kept as simple as possible to enable this lab by permitting all cleartext traffic from the client zone to the server zone:

```

1 edit security policies
2
3 set from-zone tester-client to-zone tester-server policy tester-client-1 match source-address
  any
4 set from-zone tester-client to-zone tester-server policy tester-client-1 match destination-
  address any
5 set from-zone tester-client to-zone tester-server policy tester-client-1 match application any
6 set from-zone tester-client to-zone tester-server policy tester-client-1 then permit
7
8 set global policy final-1 match source-address any
9 set global policy final-1 match destination-address any
10 set global policy final-1 match application any
11 set global policy final-1 then deny

```

1.5.1.8 NAT In the sample setup, part of the client prefixes (*10.6.0.0/15*) is used for interfacestyle source NAT to the VRRP address:

```

1 edit security nat source
2
3 set pool pool-1 address 100.64.94.10/32
4 set rule-set source-1083 from zone tester-client
5 set rule-set source-1083 to zone tester-server
6 set rule-set source-1083 rule source-1083-1 match source-address 10.6.0.0/15
7 set rule-set source-1083 rule source-1083-1 match destination-address 0.0.0.0/0
8 set rule-set source-1083 rule source-1083-1 then source-nat pool pool-1

```

Note: the MAC address in outgoing frames is the node's MAC; response frames use the VRRP MAC as the destination. ProxyARPlike behaviour can be partially mimicked by VRRP IP aliases or by a route to the NAT pool via the VRRP IP address.

1.5.1.9 Flow Tuning Only the DNS and FTP ALGs are enabled. Disabled [drop-flow](#) to avoid excessive logging (configurable since Junos when tester traffic is denied by policies, since drop-flow inserts shortlived session records for dropped traffic and causes packets to be dropped without policy evaluation in the slow path. trict TCP handshake checks to enforce stricter TCP handling (disallowing data segments until the handshake completes):

```

1 set security alg msrpc disable
2 set security alg sunrpc disable
3 set security alg talk disable
4 set security alg tftp disable
5 set security alg pptp disable
6
7 set security flow drop-flow max-sessions 0
8 set security flow tcp-session strict-syn-check

```


1.5.1.10 IPSEC IPSEC-related settings are in the *ipsec-1* Junos group, which makes it easy to keep track of configuration across hierarchies. [Appendix 2](#) covers ICL IPSEC encryption in another Junos group, which is mandatory for real-world VPN deployments.

An unnumbered tunnel interface is mapped to a specific routing instance with a static route for the tester's serverside prefix designated to IPSEC. In the next configuration section, the referenced prefix list specifies the IP addresses where IPSEC services may terminate (effectively the *ae0.1102* VRRP address):

```
1 edit groups ipsec-1
2
3 set interfaces st0 unit 0 family inet
4 set routing-instances vr-testers routing-options static route 10.104.0.0/15 next-hop st0.0
5 set routing-instances vr-testers interface st0.0
6
7 set policy-options prefix-list ipsec_termination 100.64.102.10/32
```

MNHA configuration is expanded to enable IPSEC services and reference the previously configured prefix list. Enabling *process-packet-on-backup* allows an MNHA node in the BACKUP state to process incoming IPSEC VPN traffic before changing to ACTIVE, resulting in quicker failover (otherwise, there is a time interval when the VRRP backup becomes master while MNHA waits for the ICL BFD to go down and for the activeness probe to complete):

```
1 edit groups ipsec-1 chassis high-availability services-redundancy-group 1
2
3 set managed-services ipsec
4 set prefix-list ipsec_termination routing-instance vr-testers
5 set process-packet-on-backup
6 Allow IKE on the aggregate interface unit designated for the IPSEC VPN by expanding the tester-
  server-vpn zone settings, and for simplicity bind the tunnel interface to the same zone:
7 edit groups ipsec-1 security zones
8
9 set security-zone tester-server-vpn host-inbound-traffic system-services ike
10 set security-zone tester-server-vpn interfaces st0.0
```

Corresponding security policy:

```
1 edit groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn
2
3 set policy tester-client-vpn-1 match source-address any
4 set policy tester-client-vpn-1 match destination-address any
5 set policy tester-client-vpn-1 match application any
6 set policy tester-client-vpn-1 then permit
```

Note: in in real-world deployments, strict controls would be applied.

Sample IPSEC VPN settings binding the previously configured tunnel interface; the *local-address* is set to the VRRP address of *ae0.1102*.

```
1 edit groups ipsec-1 security
2
3 set ike proposal ike-proposal-1 authentication-method pre-shared-keys
4 set ike proposal ike-proposal-1 dh-group group20
5 set ike proposal ike-proposal-1 encryption-algorithm aes-256-gcm
6 set ike proposal ike-proposal-1 lifetime-seconds 86400
7
8 set ike policy ike-policy-1 proposals ike-proposal-1
9 set ike policy ike-policy-1 pre-shared-key ascii-text <-----PSK-HERE----->
10
11 set ike gateway ike-gw-1 ike-policy ike-policy-1
12 set ike gateway ike-gw-1 address 100.64.102.20
13 set ike gateway ike-gw-1 dead-peer-detection
14 set ike gateway ike-gw-1 external-interface ae0.1102
15 set ike gateway ike-gw-1 local-address 100.64.102.10
16 set ike gateway ike-gw-1 version v2-only
17
18 set ipsec proposal ipsec-proposal-1 encryption-algorithm aes-256-gcm
19 set ipsec proposal ipsec-proposal-1 lifetime-seconds 3600
20
```

```

21 set ipsec policy ipsec-policy-1 perfect-forward-secrecy keys group20
22 set ipsec policy ipsec-policy-1 proposals ipsec-proposal-1
23
24 set ipsec vpn vpn-1 bind-interface st0.0
25 set ipsec vpn vpn-1 ike gateway ike-gw-1
26 set ipsec vpn vpn-1 ike ipsec-policy ipsec-policy-1

```

Bonus - add-on in the form of Post-quantum Pre-shared Keys (PPK) RFC [8784](#) support, where this setup serves as a technology demonstrator using static keys:

```

1 top edit groups ipsec-1 security
2
3 set key-manager profiles quantum static key-id ascii-text <-----KEY-ID-HERE----->
4 set key-manager profiles quantum static key ascii-text <-----32-character++KEY-HERE----->
5 set ike gateway ike-gw-1 ppk-profile quantum

```

1.5.1.11 Applying Groups To apply the groups containing both the common VRRP settings and the IPSECrelated configuration:

```

1 set apply-groups vrrp-1
2 set apply-groups ipsec-1

```

Final step: reboot as prompted by the CLI after committing the MNHA configuration for the first time.

1.5.2 SRX1600-2

Appendix 1 contains the full configurations; this section captures only the differences between SRX1600-1 and SRX1600-2.

First thing, the configuration context:

```

1 set system host-name srx1600-2

```

MNHArelevant loopback and physical interconnect interface addressing:

```

1 edit interfaces
2
3 set lo0 unit 0 family inet address 192.168.2.1/32
4 set lo0 unit 0 family inet address 192.168.2.2/32
5
6 set interfaces xe-0/1/0 mtu 2000
7 set interfaces xe-0/1/1 mtu 2000
8
9 set xe-0/1/0 unit 0 family inet address 192.168.0.1/31
10 set xe-0/1/1 unit 0 family inet address 192.168.0.3/31

```

ICL and ICD export policies adjusted for SRX1600-2 addressing, following the same principles as on SRX1600-1:

```

1 edit policy-options
2
3 set policy-statement export-mnha-icl term icl from interface lo0.0
4 set policy-statement export-mnha-icl term icl from route-filter 192.168.2.1/32 exact
5 set policy-statement export-mnha-icl term icl then accept
6 set policy-statement export-mnha-icl term icd from interface lo0.0
7 set policy-statement export-mnha-icl term icd from route-filter 192.168.2.2/32 exact
8 set policy-statement export-mnha-icl term icd then as-path-prepend 65002
9 set policy-statement export-mnha-icl term icd then accept
10 set policy-statement export-mnha-icl term final then reject
11
12 set policy-statement export-mnha-icd term icl from interface lo0.0
13 set policy-statement export-mnha-icd term icl from route-filter 192.168.2.1/32 exact
14 set policy-statement export-mnha-icd term icl then as-path-prepend 65002
15 set policy-statement export-mnha-icd term icl then accept
16 set policy-statement export-mnha-icd term icd from interface lo0.0
17 set policy-statement export-mnha-icd term icd from route-filter 192.168.2.2/32 exact

```

```

18 set policy-statement export-mnha-icd term icd then accept
19 set policy-statement export-mnha-icd term final then reject

```

Reversed autonomous system numbers and IP addresses:

```

1 edit routing-instances vr-mnha
2
3 set instance-type virtual-router
4 set routing-options autonomous-system 65002
5 set protocols bgp group mnha peer-as 65001
6 set protocols bgp group mnha bfd-liveness-detection minimum-interval 500
7 set protocols bgp group mnha bfd-liveness-detection multiplier 3
8 set protocols bgp group mnha neighbor 192.168.0.0 export export-mnha-icl
9 set protocols bgp group mnha neighbor 192.168.0.2 export export-mnha-icd
10 set interface xe-0/1/0.0
11 set interface xe-0/1/1.0
12 set interface lo0.0

```

MNHA settings are reversed in terms of IP addressing and local/peer ID settings:

```

1 edit chassis high-availability
2
3 set local-id 2
4 set local-id local-ip 192.168.2.1
5 set local-id local-forwarding-ip 192.168.2.2
6
7 set peer-id 1 peer-ip 192.168.1.1
8 set peer-id 1 interface lo0.0
9 set peer-id 1 routing-instance vr-mnha
10 set peer-id 1 peer-forwarding-ip 192.168.1.2
11 set peer-id 1 peer-forwarding-ip interface lo0.0
12 set peer-id 1 peer-forwarding-ip liveness-detection minimum-interval 1000
13 set peer-id 1 peer-forwarding-ip liveness-detection multiplier 3
14 set peer-id 1 liveness-detection minimum-interval 1000
15 set peer-id 1 liveness-detection multiplier 3
16
17 set services-redundancy-group 1 deployment-type routing
18 set services-redundancy-group 1 peer-id 1
19 set services-redundancy-group 1 activeness-probe dest-ip 100.64.81.20
20 set services-redundancy-group 1 activeness-probe dest-ip src-ip 100.64.81.10
21 set services-redundancy-group 1 activeness-probe dest-ip routing-instance vr-testers
22 set services-redundancy-group 1 monitor monitor-object interface-1 object-threshold 100
23 set services-redundancy-group 1 monitor monitor-object interface-1 interface threshold 100
24 set services-redundancy-group 1 monitor monitor-object interface-1 interface interface-name ae0
   weight 100
25 set services-redundancy-group 1 monitor srg-threshold 100
26 set services-redundancy-group 1 active-signal-route 100.64.0.1
27 set services-redundancy-group 1 active-signal-route routing-instance vr-mnha
28 set services-redundancy-group 1 backup-signal-route 100.64.0.0
29 set services-redundancy-group 1 backup-signal-route routing-instance vr-mnha
30 set services-redundancy-group 1 activeness-priority 100

```

Interfaces with adjusted IP addressing:

```

1 edit interfaces ae0
2
3 set aggregated-ether-options minimum-links 2
4 set aggregated-ether-options lacp active
5 set aggregated-ether-options lacp periodic fast
6
7 set unit 1081 vlan-id 1081
8 set unit 1081 family inet address 100.64.81.12/24 vrrp-group 1 virtual-address 100.64.81.10
9 set unit 1081 family inet6 address dead:81::12/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :81::10
10
11 set unit 1082 vlan-id 1082
12 set unit 1082 family inet address 100.64.82.12/24 vrrp-group 1 virtual-address 100.64.82.10
13 set unit 1082 family inet address 100.64.82.12/24 vrrp-group 1 virtual-address 100.64.82.20
14 set unit 1082 family inet address 100.64.82.12/24 vrrp-group 1 no-accept-data
15 set unit 1082 family inet6 address dead:82::12/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :82::10
16 set unit 1082 family inet6 address dead:82::12/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :82::20
17 set unit 1082 family inet6 address dead:82::12/64 vrrp-inet6-group 1 no-accept-data
18
19 set unit 1083 vlan-id 1083
20 set unit 1083 family inet address 100.64.83.12/24 vrrp-group 1 virtual-address 100.64.83.10
21 set unit 1083 family inet address 100.64.84.12/24 vrrp-group 2 virtual-address 100.64.84.10

```

```

22 set unit 1083 family inet address 100.64.84.12/24 vrrp-group 2 virtual-address 100.64.84.20
23 set unit 1083 family inet address 100.64.84.12/24 vrrp-group 2 no-accept-data
24 set unit 1083 family inet6 address dead:83::12/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :83::10
25 set unit 1083 family inet6 address dead:84::12/64 vrrp-inet6-group 2 virtual-inet6-address dead
   :84::10
26 set unit 1083 family inet6 address dead:84::12/64 vrrp-inet6-group 2 virtual-inet6-address dead
   :84::20
27 set unit 1083 family inet6 address dead:84::12/64 vrrp-inet6-group 2 no-accept-data
28
29 set unit 1091 vlan-id 1091
30 set unit 1091 family inet address 100.64.91.12/24 vrrp-group 1 virtual-address 100.64.91.10
31 set unit 1091 family inet6 address dead:91::12/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :91::10
32
33 set unit 1092 vlan-id 1092
34 set unit 1092 family inet address 100.64.92.12/24 vrrp-group 1 virtual-address 100.64.92.10
35 set unit 1092 family inet6 address dead:92::12/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :92::10
36
37 set unit 1093 vlan-id 1093
38 set unit 1093 family inet address 100.64.93.12/24 vrrp-group 1 virtual-address 100.64.93.10
39 set unit 1093 family inet6 address dead:93::12/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :93::10
40
41 set unit 1094 vlan-id 1094
42 set unit 1094 family inet address 100.64.94.12/24 vrrp-group 1 virtual-address 100.64.94.10
43 set unit 1094 family inet6 address dead:94::12/64 vrrp-inet6-group 1 virtual-inet6-address dead
   :94::10
44
45 set unit 1100 vlan-id 1100
46 set unit 1100 family inet address 100.64.100.12/24 vrrp-group 1 virtual-address 100.64.100.10
47
48 set unit 1101 vlan-id 1101
49 set unit 1101 family inet address 100.64.101.12/24 vrrp-group 1 virtual-address 100.64.101.10
50
51 set unit 1102 vlan-id 1102
52 set unit 1102 family inet address 100.64.102.12/24 vrrp-group 1 virtual-address 100.64.102.10

```

1.6 Validation

Validation will observe MNHA and VRRP status, administrative and simulated outage-triggered failover with and without traffic load, and a test of the SRX's ability to egress VRRP advertisements when the PFE is overloaded.

The following are sample status combinations of the MNHA and VRRP state machines operating together, inspired by settings from the demo setup:

Scenario	MNHA SRG State	Active-Signal Route Present?	VRRP Priority	VRRP State	Behavior / Notes
Normal steady state of SRX1600-1 upon boot	ACTIVE	✓ Yes	150	MASTER	VRRP master aligns with MNHA active node. Traffic forwarded normally via VIPs.
Normal steady state of SRX1600-2 upon boot	BACKUP	✗ No	50 (150-100)	BACKUP	Route tracking reduces VRRP priority. SRX1600-2 has smaller priority and becomes backup.
Administrative failover initiated	ACTIVE → BACKUP	✗ Removed	50 (150-100)	BACKUP	VRRP quickly transfers mastership due to lost tracked route.
	BACKUP → ACTIVE	✓ Installed	150	MASTER	VRRP takes control almost immediately (sub-second).
ICL only failure while both nodes up (split-brain risk)	Maintained, activeness probing started.	Maintained	Maintained	Maintained	VRRP influences MNHA via probe source IP → only VRRP master can send probe.
ICL only failure + activeness probes responded on ACTIVE	ACTIVE	✓ Yes	150	MASTER	Continuous probing.
ICL only failure + activeness probes failing on BACKUP	BACKUP	✗ No	50 (150-100)	BACKUP	Continuous probing, VRRP would trigger MNHA transition.
MNHA tracked object failure on one node	ACTIVE → INELIGIBLE	✗ Removed	50 (150-100)	BACKUP	MNHA monitoring triggers failover → VRRP follows via route tracking.
	BACKUP → ACTIVE	✓ Installed	150	MASTER	VRRP becomes master immediately after MNHA transition.
Sudden failure of active node	ACTIVE lost	N/A	N/A	BACKUP → MASTER (on peer)	VRRP detects missing advertisements → takes over in ~3 intervals (~<1s).
	BACKUP → ACTIVE (delayed by MNHA probing)	✓ Installed (after activeness probing)	150	MASTER	VRRP transitions faster than MNHA; MNHA catches up later (~7s).
ICL only failure while both nodes up (split-brain risk, broken activeness probe corner case)	Maintained, activeness probing started.	Maintained	Maintained	Maintained	VRRP influences MNHA via probe source IP → only VRRP master can send probe.
ICL only failure + activeness probes failing on ACTIVE	ACTIVE → BACKUP	✗ Removed	50 (150-100)	MASTER	VRRP would not preempt due to the same priority. Continuous probing. VPNs break.
ICL only failure + activeness probes failing on BACKUP	BACKUP	✗ No	50 (150-100)	BACKUP	Continuous probing, VRRP would fail-over in case of other node failing.

Figure 4: Different scenarios, different states

1.6.1 MNHA/VRRP Status

1.6.1.1 MNHA Status Basic MNHA status on SRX1600-1 shows ICL and ICD addressing and status of local (ACTIVE) and remote node (BACKUP):

```

1  show chassis high-availability information
2
3  Node failure codes:
4      HW Hardware monitoring LB Loopback monitoring
5      MB Mbuf monitoring SP SPU monitoring
6      CS Cold Sync monitoring SU Software Upgrade
7
8  Node Status: ONLINE
9  Grid-id: 0
10 Local-id: 1
11 Local-IP: 192.168.1.1
12 Local Forwarding IP: 192.168.1.2
13 HA Peer Information:

```

```

14
15 Peer Id: 2 IP address: 192.168.2.1 Interface: lo0.0
16 Routing Instance: vr-mnha
17 Encrypted: NO Conn State: UP
18 Configured BFD Detection Time: 3 * 1000ms
19 Cold Sync Status: COMPLETE
20 Peer Forwarding IP: 192.168.2.2 Interface: lo0.0
21 Peer ICD Conn State: UP
22
23 SRG failure event codes:
24   BF BFD monitoring
25   IP IP monitoring
26   IF Interface monitoring
27   CP Control Plane monitoring
28
29 Services Redundancy Group: 1
30   Deployment Type: ROUTING
31   Status: ACTIVE
32   Activeness Priority: 200
33   Preemption: DISABLED
34   Process Packet In Backup State: YES
35   Control Plane State: READY
36   System Integrity Check: N/A
37   Failure Events: NONE
38   Peer Information:
39     Peer Id: 2
40     Status : BACKUP
41     Health Status: HEALTHY
42     Failover Readiness: READY

```

The following excerpt from the detailed view shows information about the signal routes. SRX1600-1 is in ACTIVE status and has the active-signal route installed; the mutually exclusive Backup Route is not installed. Also, because the nodes can see each other, the splitbrain probe is not running:

```

1 show chassis high-availability information detail
2
3 <SNIP>
4   Signal Route Info:
5     Active Signal Route:
6       IP: 100.64.0.1
7       Routing Instance: vr-mnha
8       Status: INSTALLED
9
10    Backup Signal Route:
11      IP: 100.64.0.0
12      Routing Instance: vr-mnha
13      Status: NOT INSTALLED
14
15    Split-brain Prevention Probe Info:
16      DST-IP: 100.64.81.20
17      SRC-IP: 100.64.81.10
18      Routing Instance: vr-testers
19      Type: ICMP Probe
20      Status: NOT RUNNING
21      Result: N/A Reason: N/A
22 <SNIP>

```

1.6.1.2 VRRP Status VRRP summary executed on the SRX1600-1 in ACTIVE state shows VRRP master status and installed VIPs:

```

1 show vrrp summary
2
3 Interface State Group VR state VR Mode Type Address
4 ae0.1081 up 1 master Active lcl 100.64.81.11
5   vip 100.64.81.10
6 ae0.1081 up 1 master Active lcl dead:81::11
7   vip fe80::200:5eff:fe00:201
8   vip dead:81::10
9 ae0.1082 up 1 master Active lcl 100.64.82.11
10  vip 100.64.82.10
11  vip 100.64.82.20
12 ae0.1082 up 1 master Active lcl dead:82::11
13  vip fe80::200:5eff:fe00:201
14  vip dead:82::10

```

```

15 vip dead:82::20
16 <SNIP>

```

The detailed view shows a priority of 150 based on the presence of the MNHA active-signal route:

```

1 show vrrp detail
2
3 Physical interface: ae0, Unit: 1081, Vlan-id: 1081, Address: 100.64.81.11/24
4 Index: 73, SNMP ifIndex: 1091, VRRP-Traps: disabled, VRRP-Version: 3
5 Interface state: up, Group: 1, State: master, VRRP Mode: Active
6 Priority: 150, Advertisement interval: 0.200, Authentication type: none
7 Advertisement threshold: 3, Computed send rate: 0
8 Preempt: yes, Accept-data mode: yes, VIP count: 1, VIP: 100.64.81.10
9 Advertisement Timer: 0.004s, Master router: 100.64.81.11
10 Virtual router uptime: 10:18:09, Master router uptime: 00:00:08
11 Virtual Mac: 00:00:5e:00:01:01
12 Preferred: yes
13 Tracking: enabled
14 Current priority: 150, Configured priority: 150
15 Priority hold time: disabled
16 Interface tracking: disabled
17 Route tracking: enabled, Route count: 1
18 Route VRF name Route state Priority cost
19 100.64.0.1/32 vr-mnha up 100
20 <SNIP>

```

The SRX1600-2 in BACKUP state, which does not have the tracked route installed, has a priority of 50:

```

1 show vrrp detail
2
3 Physical interface: ae0, Unit: 1082, Vlan-id: 1082, Address: 100.64.82.12/24
4 Index: 74, SNMP ifIndex: 1106, VRRP-Traps: disabled, VRRP-Version: 3
5 Interface state: up, Group: 1, State: backup, VRRP Mode: Active
6 Priority: 50, Advertisement interval: 0.200, Authentication type: none
7 Advertisement threshold: 3, Computed send rate: 0
8 Preempt: yes, Accept-data mode: no, VIP count: 2, VIP: 100.64.82.10, 100.64.82.20
9 Dead timer: 0.640s, Master priority: 150, Master router: 100.64.82.11
10 Virtual router uptime: 1d 08:28
11 Preferred: yes
12 Tracking: enabled
13 Current priority: 50, Configured priority: 150
14 Priority hold time: disabled
15 Interface tracking: disabled
16 Route tracking: enabled, Route count: 1
17 Route VRF name Route state Priority cost
18 100.64.0.1/32 vr-mnha unknown 100
19 <SNIP>

```

Above MNHA/VRRP status consolidated using *mnha_status* template from *template-ops* (<https://github.com/JNPRAutomate/template-ops>) tool, executed off-box:

```

1 ./template-ops.py --profile mnha_status
2
3 -----
4 | device | Status | Failure Events | Sum mon obj Weights | Sum BFD sub-obj Weights | BGP
5 | Sessions Down | VRRP masters | Remote Status | Remote Readiness | Junos Version |
6 -----
6 | srx1600-1 | ACTIVE | NONE | 0 | 0 | 0 | 22 | HEALTHY | READY | 24.4R2-S3.5 |
7 -----
8 | srx1600-2 | BACKUP | NONE | 0 | 0 | 0 | 0 | HEALTHY | N/A | 24.4R2-S3.5 |
9 -----

```

1.6.2 Fail-over

The time to fail-over differs between best-case administrator triggered fail-over and outage situations, when the Active node is suddenly lost. Methodology for measuring the time are two bi-directional UDP sessions,

one through clear-text, the other via IPSEC VPN. The packet rate is 1000 PPS meaning that every lost packet accounts for 1 millisecond.

1.6.2.1 Administrative Fail-over SRG-1 failover from SRX1600-1 to SRX1600-2:

```
1 request chassis high-availability failover services-redundancy-group 1 peer-id 2
2
3 Initiated manual failover for Services-redundancy group 1
```

The almost-too-good-to-be-true result of four lost packets in one direction (fewer in the other; see the Frames Delta column) is also influenced by both SRX devices being connected to the same QFX10k switch line card:

Tx Port	Rx Port	Traffic Item	Source/Dest Endpoint Pair	Tx Frames	Rx Frames	Frames Delta	Loss %	Tx Frame Rate	Rx Frame Rate
Ethernet - 001	Ethernet - 002	udp-single-session	10.100.0.1-10.102.0.1	255,553	255,549	4	0.002	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session	10.102.0.1-10.100.0.1	255,553	255,552	1	0.000	1,000.000	1,000.000
Ethernet - 001	Ethernet - 002	udp-single-session-ipsec	10.100.0.1-10.104.0.1	255,553	255,551	2	0.001	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session-ipsec	10.104.0.1-10.100.0.1	255,553	255,553	0	0.000	1,000.000	1,000.000

Figure 5: IxNetwork capture from an administrative fail-over

Corresponding logs from SRX1600-2 - the new MNHA SRG-1 Active node and VRRP master:

```
1 show log messages | match jsrp
2
3 srx1600-2 jsrpd[10035]: JSRPD_HA_SRG_STATE_CHANGE: SRG[1]:HA state transitioned [BACKUP -->
  ACTIVE] Reason: [Manual failover at Peer]
4 <SNIP>
5 show log messages | match vrrp
6
7 srx1600-2 vrrpd[10007]: VRRPD_NEW_MASTER: Interface ae0.1102 (local address 100.64.102.12)
  became VRRP master for group 1 with master reason Running priority high
8 <SNIP>
```

1.6.2.2 PFE Failure Fail-over The sudden outage is simulated by crashing the SRX PFE: log in to the active SRX1600 host OS from the Junos RE root shell and send *SIGKILL* (signal 9) to the *srxpfe* process:

```
1 start shell user root
2 vhcliclient -s
3 pkill -SIGKILL srxpfe
```

The stateless tester reports between 828-913ms of traffic loss, which corresponds to loss of 3x 200ms VRRP advertisements:

Tx Port	Rx Port	Traffic Item	Source/Dest Endpoint Pair	Tx Frames	Rx Frames	Frames Delta	Loss %	Tx Frame Rate	Rx Frame Rate
Ethernet - 001	Ethernet - 002	udp-single-session	10.100.0.1-10.102.0.1	75,757	74,844	913	1.205	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session	10.102.0.1-10.100.0.1	75,757	74,850	907	1.197	1,000.000	1,000.000
Ethernet - 001	Ethernet - 002	udp-single-session-ipsec	10.100.0.1-10.104.0.1	75,757	74,849	908	1.199	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session-ipsec	10.104.0.1-10.100.0.1	75,757	74,929	828	1.093	1,000.000	1,000.000

Figure 6: IxNetwork capture from PFE Failure test

Note: recovering a node from the above condition can be done by waiting for the system to recover automatically, running the *restart chassis control* command, or rebooting the device.

Corresponding logs on SRX1600-1 taking over after SRX1600-2's arbitrary PFE crash:

```
1 show log messages | match vrrp
2
3 11:38:01.843 srx1600-1 vrrpd[10007]: VRRPD_NEW_MASTER: Interface ae0.1083 (local address
  100.64.84.11) became VRRP master for group 2 with master reason Adjacency down or no
  response from master
```



```

4 show log messages | match jsrp
5
6 11:38:03.451 srx1600-1 jsrpd[10031]: JSRPD_BFD_STATE_DOWN: HA peer BFD is down
7 11:38:04.101 srx1600-1 jsrpd[10031]: JSRPD_ICD_BFD_STATE_DOWN: HA peer ICD BFD is down
8 11:38:08.495 srx1600-1 jsrpd[10031]: JSRPD_HA_SRG_STATE_CHANGE: SRG[1]:HA state transitioned [
  BACKUP -->ACTIVE] Reason: [Split Brain Prevention logic result]

```

As shown in the logs above, VRRP transitions faster than MNHA because MNHA waits for the ICL link to fail (3 x 1000 ms BFD) and for an activeness probe to complete (4-5 lost ICMP echo probes with 1 second interval). The effective transition time from BACKUP to ACTIVE can be measured by disabling *process-packet-on-backup*. According to the tester capture below from a test without *process-packet-on-backup*, the transition takes more than 7 seconds, consisting of the ICL BFD reaction time and the activeness probe completion:

Tx Port	Rx Port	Traffic Item	Source/Dest Endpoint Pair	Tx Frames	Rx Frames	Frames Delta	Loss %	Tx Frame Rate	Rx Frame Rate
Ethernet - 001	Ethernet - 002	udp-single-session	10.100.0.1-10.102.0.1	31,371	30,529	842	2.684	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session	10.102.0.1-10.100.0.1	31,371	30,536	835	2.662	1,000.000	1,000.000
Ethernet - 001	Ethernet - 002	udp-single-session-ipsec	10.100.0.1-10.104.0.1	31,371	23,883	7,488	23.869	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session-ipsec	10.104.0.1-10.100.0.1	31,371	23,883	7,488	23.869	1,000.000	1,000.000

Figure 7: IxNetwork capture from a test without *process-packet-on-backup* enabled

Note: the *process-packet-on-backup* setting is used in the real world for faster IPSEC failovers.

1.6.2.3 Fail-over Under Load Failover times are typically measured under best-case conditions with no background traffic. In this test there will be 10 Gbps of background traffic consisting of 44 KB HTTP GET transactions at a rate of 25,000 transactions per second (TPS), split 50/50 between IPv4 and IPv6; one quarter of the IPv4 traffic is subjected to source NAT/PAT, resulting in approximately 50% PFE load overall.

CPU load on SRX1600-2 acting as the active node:

```

1 show security monitoring
2
3           Flow session Flow session CP session CP session
4 FPC PIC CPU Mem current maximum current maximum
5  0 0 48 24 101140 2097152 0 0

```

CPU load on the SRX1600-1 backup node caused by session synchronization:

```

1 show security monitoring
2
3           Flow session Flow session CP session CP session
4 FPC PIC CPU Mem current maximum current maximum
5  0 0 5 24 201278 2097152 0 0

```

The stateful BreakingPoint tester shows 25,000 TPS (equal to TCP CPS), which with the given transaction length results in 10 Gbps (L2) traffic and over 1 million FPS:

Application Transaction Rate Attempted: 25,003 Successful: 24,995 Failed: 0	TCP Client Connection Rate Attempted: 25,007 Established: 25,005 Closed: 25,005	TCP Server Connection Rate Established: 25,000 Closed: 24,995	Bandwidth Tx (Mbps): 10,003.93 Rx (Mbps): 10,003.94
Application Transactions Attempted: 687,162 Successful: 687,139 Failed: 0 Aborted: 0	Cumulative Client TCP Connections Attempted: 687,108 Established: 687,104 Closed (FIN): 687,003 Closed (RST): 0	Cumulative Server TCP Connections Established: 687,135 Closed (FIN): 687,033 Closed (RST): 0	Frame Rate Tx (Fps): 1,097,946 Rx (Fps): 1,098,010 Instantaneous Average Frame Size Tx (Bytes): 1,118.94 Rx (Bytes): 1,118.88 Bidirectional (Bytes): 1,118.91

Figure 8: BreakingPoint capture of stateful traffic run

Corresponding Junos *monitor interface traffic* command views on SRX1600-2: both PPS and BPS for the *ae0* aggregate interface and its member links (*xe-0/2/**):

```

1  srx1600-2 Seconds: 14 Time: 00:09:17
2
3  Interface Link Input packets (pps) Output packets (pps)
4  <SNIP>
5  xe-0/1/0 Up 2974856 (2) 119966046 (49141)
6  xe-0/1/1 Up 624579 (7) 970524 (6)
7  xe-0/2/0 Up 1295917551 (266778) 1967372599 (267805)
8  xe-0/2/1 Up 1459834331 (270961) 2037997532 (268032)
9  xe-0/2/2 Up 1378642636 (267858) 1957593960 (266541)
10 xe-0/2/3 Up 1377443322 (268619) 1876522086 (272014)
11 ae0 Up 5511837840 (1074216) 7839486179 (1074392)
12 <SNIP>
13 srx1600-2 Seconds: 72 Time: 00:10:15
14
15 Interface Link Input bytes (bps) Output bytes (bps)
16 <SNIP>
17 xe-0/1/0 Up 3417650804 (1656) 155402471118 (499400256)
18 xe-0/1/1 Up 34375078 (2600) 82907071 (4104)
19 xe-0/2/0 Up 923295697200 (2401717984) 1583339743979 (2415735288)
20 xe-0/2/1 Up 960068287169 (2404307776) 1628903311683 (2362414416)
21 xe-0/2/2 Up 936893030827 (2427265184) 1583412086580 (2422996400)
22 xe-0/2/3 Up 951195804315 (2390637328) 1574891123349 (2422724488)
23 ae0 Up 3771452819511 (9623928272) 6370546265749 (9623870552)
24 <SNIP>

```

Note: the ICL and ICD links (*xe-0/1/0* and *xe-0/1/1*, respectively) were kept for capturing ICL packet rate and bandwidth during the 25,000 TPS background traffic.

And above using load template from [template-ops](#) tool executed off-box:

```

1  ./template-ops.py --profile load
2
3  -----
4  | device | PFE load | IPv4 sessions | IPv6 sessions | NAT sessions | IPv4 CPS | IPv6 CPS | in
5  | Gbps   | out Gbps | in MPPS       | out MPPS      | np-cache %   |          |          |
6  | srx1600-1 | 55 | 51427 | 43128 | 15601 | 12717 | 12718 | 9.62 | 9.62 | 1.08 | 1.08 | N/A |
7  |-----|
8  | srx1600-2 | 7 | 94516 | 86275 | 28470 | 12611 | 12612 | 0.00 | 0.00 | 0.00 | 0.00 | N/A |
9  |-----|

```

1.6.2.4 Administrative Fail-over Under Load During an administrative failover, a maximum outage of 2 ms was observed for encrypted traffic, showing no influence from background traffic.

Tx Port	Rx Port	Traffic Item	Source/Dest Endpoint Pair	Tx Frames	Rx Frames	Frames Delta	Loss %	Tx Frame Rate	Rx Frame Rate
Ethernet - 001	Ethernet - 002	udp-single-session	10.100.0.1-10.102.0.1	36,566	36,566	0	0.000	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session	10.102.0.1-10.100.0.1	36,566	36,566	0	0.000	1,000.000	1,000.000
Ethernet - 001	Ethernet - 002	udp-single-session-ipsec	10.100.0.1-10.104.0.1	36,566	36,564	2	0.005	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session-ipsec	10.104.0.1-10.100.0.1	36,566	36,565	1	0.003	1,000.000	1,000.000

Figure 9: IxNetwork capture from an administrative fail-over under load

1.6.2.5 PFE Failure Fail-over Under Load Also, an arbitrary PFE failure resulted in similar outage times as without background traffic:

Tx Port	Rx Port	Traffic Item	Source/Dest Endpoint Pair	Tx Frames	Rx Frames	Frames Delta	Loss %	Tx Frame Rate	Rx Frame Rate
Ethernet - 001	Ethernet - 002	udp-single-session	10.100.0.1-10.102.0.1	211,994	211,154	840	0.396	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session	10.102.0.1-10.100.0.1	211,994	211,158	836	0.394	1,000.000	1,000.000
Ethernet - 001	Ethernet - 002	udp-single-session-ipsec	10.100.0.1-10.104.0.1	211,994	211,158	836	0.394	1,000.000	1,000.000
Ethernet - 002	Ethernet - 001	udp-single-session-ipsec	10.104.0.1-10.100.0.1	211,994	211,163	831	0.392	1,000.000	1,000.000

Figure 10: IxNetwork capture from PFE Failure test under load

1.6.3 Stability Under Overload

Another aspect of highavailability mechanisms that use active advertisements on the revenue interface is stability under interface overload. In this test, IxNetwork sent 6 MPPS distributed across 16,000 sessions, of which about 1.7 MPSS per direction (3.4 Mpps total) passed:

Tx Port	Rx Port	Traffic Item	Source/Dest Endpoint Pair	Tx Frames	Rx Frames	Frames Delta	Loss %	Tx Frame Rate	Rx Frame Rate
Ethernet - 001	Ethernet - 002	udp-x-sessions	10.100.0.1-10.102.0.1	161,356,715	91,900,660	69,456,055	43.045	2,999,998.000	1,722,265.000
Ethernet - 002	Ethernet - 001	udp-x-sessions	10.102.0.1-10.100.0.1	161,356,713	91,890,073	69,466,640	43.052	2,999,997.000	1,722,810.000

Figure 11: IxNetwork capture from overload test

Resulting in 99% load across PFE threads:

```

1 show security monitoring performance spu extensive
2
3      CPU Util Wutil Status SchedCounter
4 FPC0
5   PICO
6       0 99 95 alive 245543
7       1 99 95 alive 245543
8       2 99 95 alive 245543
9       3 99 95 alive 245543

```

Junos monitor interface traffic view showing only the packet rate the SRX1600 is able to capture from the interfaces, which equals the tester's received rate:

```

1 Interface Link Input packets (pps) Output packets (pps)
2 <SNIP>
3 xe-0/2/0 Up 876527556 (856012) 1302808136 (854705)
4 xe-0/2/1 Up 1039482820 (857644) 1373632900 (854643)
5 xe-0/2/2 Up 958379506 (857998) 1293103938 (854591)
6 xe-0/2/3 Up 958078774 (856347) 1211917128 (854702)
7 ae0 Up 3832468656 (3428001) 5181462102 (3418641)
8 <SNIP>

```

During a multi-hour period, no VRRP transitions occurred, demonstrating the solution's resiliency. This is attributable to Junos prioritizing outgoing VRRP advertisements as Network Control, Junos's separation of control and data planes, and the fact that the other node monitoring VRRP advertisements was not receiving the frame rate.

Queue on interface actively transmitting VRRP advertisements:

```

1  show interfaces queue xe-0/2/0
2
3  <SNIP>
4  Queue: 3, Forwarding classes: network-control
5    Queued:
6      Packets : 577630 129 pps
7      Bytes  : 43865928 78840 bps
8    Transmitted:
9      Packets : 577630 129 pps
10     Bytes  : 43865928 78840 bps
11
12  <SNIP>

```

1.7 Appendix 1: Complete Configurations

1.7.1 SRX1600-1

```

1  set version 24.4R2-S3.5
2  set groups vrrp-1 interfaces ae0 unit <*> family inet address <*> vrrp-group <*> priority 150
3  set groups vrrp-1 interfaces ae0 unit <*> family inet address <*> vrrp-group <*> fast-interval
   200
4  set groups vrrp-1 interfaces ae0 unit <*> family inet address <*> vrrp-group <*> accept-data
5  set groups vrrp-1 interfaces ae0 unit <*> family inet address <*> vrrp-group <*> track route
   100.64.0.1/32 routing-instance vr-mnha priority-cost 100
6  set groups vrrp-1 interfaces ae0 unit <*> family inet6 address <*> vrrp-inet6-group <*> priority
   150
7  set groups vrrp-1 interfaces ae0 unit <*> family inet6 address <*> vrrp-inet6-group <*> fast-
   interval 200
8  set groups vrrp-1 interfaces ae0 unit <*> family inet6 address <*> vrrp-inet6-group <*> accept-
   data
9  set groups vrrp-1 interfaces ae0 unit <*> family inet6 address <*> vrrp-inet6-group <*> track
   route 100.64.0.1/32 routing-instance vr-mnha priority-cost 100
10 set groups ipsec-1 chassis high-availability services-redundancy-group 1 prefix-list
   ipsec_termination routing-instance vr-testers
11 set groups ipsec-1 chassis high-availability services-redundancy-group 1 managed-services ipsec
12 set groups ipsec-1 chassis high-availability services-redundancy-group 1 process-packet-on-
   backup
13 set groups ipsec-1 security ike proposal ike-proposal-1 authentication-method pre-shared-keys
14 set groups ipsec-1 security ike proposal ike-proposal-1 dh-group group20
15 set groups ipsec-1 security ike proposal ike-proposal-1 encryption-algorithm aes-256-gcm
16 set groups ipsec-1 security ike proposal ike-proposal-1 lifetime-seconds 86400
17 set groups ipsec-1 security ike policy ike-policy-1 proposals ike-proposal-1
18 set groups ipsec-1 security ike policy ike-policy-1 pre-shared-key ascii-text <-----PSK-HERE
   ----->
19 set groups ipsec-1 security ike gateway ike-gw-1 ike-policy ike-policy-1
20 set groups ipsec-1 security ike gateway ike-gw-1 address 100.64.102.20
21 set groups ipsec-1 security ike gateway ike-gw-1 dead-peer-detection
22 set groups ipsec-1 security ike gateway ike-gw-1 external-interface ae0.1102
23 set groups ipsec-1 security ike gateway ike-gw-1 local-address 100.64.102.10
24 set groups ipsec-1 security ike gateway ike-gw-1 version v2-only
25 set groups ipsec-1 security ike gateway ike-gw-1 ppk-profile quantum
26 set groups ipsec-1 security ipsec proposal ipsec-proposal-1 encryption-algorithm aes-256-gcm
27 set groups ipsec-1 security ipsec proposal ipsec-proposal-1 lifetime-seconds 3600
28 set groups ipsec-1 security ipsec policy ipsec-policy-1 perfect-forward-secrecy keys group20
29 set groups ipsec-1 security ipsec policy ipsec-policy-1 proposals ipsec-proposal-1
30 set groups ipsec-1 security ipsec vpn vpn-1 bind-interface st0.0
31 set groups ipsec-1 security ipsec vpn vpn-1 ike gateway ike-gw-1
32 set groups ipsec-1 security ipsec vpn vpn-1 ike ipsec-policy ipsec-policy-1
33 set groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn policy
   tester-client-vpn-1 match source-address any
34 set groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn policy
   tester-client-vpn-1 match destination-address any
35 set groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn policy
   tester-client-vpn-1 match application any
36 set groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn policy
   tester-client-vpn-1 then permit
37 set groups ipsec-1 security zones security-zone tester-server-vpn host-inbound-traffic system-
   services ike
38 set groups ipsec-1 security zones security-zone tester-server-vpn interfaces st0.0
39 set groups ipsec-1 security key-manager profiles quantum static key-id ascii-text <-----KEY-ID-
   HERE----->
40 set groups ipsec-1 security key-manager profiles quantum static key ascii-text <-----32-
   character+-KEY-HERE----->

```

```

41 set groups ipsec-1 interfaces st0 unit 0 family inet
42 set groups ipsec-1 policy-options prefix-list ipsec_termination 100.64.102.10/32
43 set groups ipsec-1 routing-instances vr-testers routing-options static route 10.104.0.0/15 next-
  hop st0.0
44 set groups ipsec-1 routing-instances vr-testers interface st0.0
45 set groups icl-ipsec chassis high-availability peer-id 2 vpn-profile icl
46 set groups icl-ipsec security ike proposal ike-icl-proposal authentication-method pre-shared-
  keys
47 set groups icl-ipsec security ike proposal ike-icl-proposal dh-group group20
48 set groups icl-ipsec security ike proposal ike-icl-proposal encryption-algorithm aes-256-gcm
49 set groups icl-ipsec security ike proposal ike-icl-proposal lifetime-seconds 28800
50 set groups icl-ipsec security ike policy ike-icl-policy proposals ike-icl-proposal
51 set groups icl-ipsec security ike policy ike-icl-policy pre-shared-key ascii-text <-----PSK-HERE
  ----->
52 set groups icl-ipsec security ike gateway icl ike-policy ike-icl-policy
53 set groups icl-ipsec security ike gateway icl version v2-only
54 set groups icl-ipsec security ipsec proposal ipsec-icl-proposal encryption-algorithm aes-256-gcm
55 set groups icl-ipsec security ipsec proposal ipsec-icl-proposal lifetime-seconds 3600
56 set groups icl-ipsec security ipsec policy ipsec-icl-policy perfect-forward-secrecy keys group20
57 set groups icl-ipsec security ipsec policy ipsec-icl-policy proposals ipsec-icl-proposal
58 set groups icl-ipsec security ipsec vpn icl ha-link-encryption
59 set groups icl-ipsec security ipsec vpn icl ike gateway icl
60 set groups icl-ipsec security ipsec vpn icl ike ipsec-policy ipsec-icl-policy
61 set groups icl-ipsec security zones security-zone mnha interfaces lo0.0 host-inbound-traffic
  system-services ike
62 set apply-groups vrrp-1
63 set apply-groups ipsec-1
64 set apply-groups icl-ipsec
65 set system host-name srx1600-1
66 set system services netconf ssh
67 set system services ssh root-login allow
68 set system services ssh sftp-server
69 set system services ssh client-alive-interval 120
70 set system time-zone Europe/Amsterdam
71 set system name-server 172.30.207.10
72 set system syslog file messages any any
73 set system syslog file messages archive size 5m
74 set system syslog file messages archive files 4
75 set system syslog time-format millisecond
76 set system ntp server 172.30.207.10
77 set chassis aggregated-devices ethernet device-count 1
78 set chassis fpc 0 pic 1 pic-mode 1G10G
79 set chassis high-availability local-id 1
80 set chassis high-availability local-id local-ip 192.168.1.1
81 set chassis high-availability local-id local-forwarding-ip 192.168.1.2
82 set chassis high-availability peer-id 2 peer-ip 192.168.2.1
83 set chassis high-availability peer-id 2 interface lo0.0
84 set chassis high-availability peer-id 2 routing-instance vr-mnha
85 set chassis high-availability peer-id 2 peer-forwarding-ip 192.168.2.2
86 set chassis high-availability peer-id 2 peer-forwarding-ip interface lo0.0
87 set chassis high-availability peer-id 2 peer-forwarding-ip liveness-detection minimum-interval
  1000
88 set chassis high-availability peer-id 2 peer-forwarding-ip liveness-detection multiplier 3
89 set chassis high-availability peer-id 2 liveness-detection minimum-interval 1000
90 set chassis high-availability peer-id 2 liveness-detection multiplier 3
91 set chassis high-availability services-redundancy-group 1 deployment-type routing
92 set chassis high-availability services-redundancy-group 1 peer-id 2
93 set chassis high-availability services-redundancy-group 1 activeness-probe dest-ip 100.64.81.20
94 set chassis high-availability services-redundancy-group 1 activeness-probe dest-ip src-ip
  100.64.81.10
95 set chassis high-availability services-redundancy-group 1 activeness-probe dest-ip routing-
  instance vr-testers
96 set chassis high-availability services-redundancy-group 1 monitor monitor-object interface-1
  object-threshold 100
97 set chassis high-availability services-redundancy-group 1 monitor monitor-object interface-1
  interface threshold 100
98 set chassis high-availability services-redundancy-group 1 monitor monitor-object interface-1
  interface interface-name ae0 weight 100
99 set chassis high-availability services-redundancy-group 1 monitor srg-threshold 100
100 set chassis high-availability services-redundancy-group 1 active-signal-route 100.64.0.1
101 set chassis high-availability services-redundancy-group 1 active-signal-route routing-instance
  vr-mnha
102 set chassis high-availability services-redundancy-group 1 backup-signal-route 100.64.0.0
103 set chassis high-availability services-redundancy-group 1 backup-signal-route routing-instance
  vr-mnha
104 set chassis high-availability services-redundancy-group 1 activeness-priority 200
105 set security alg msrpc disable
106 set security alg sunrpc disable
107 set security alg talk disable
108 set security alg tftp disable
109 set security alg pptp disable

```

```

110 set security flow drop-flow max-sessions 0
111 set security flow tcp-session strict-syn-check
112 set security nat source pool pool-1 address 100.64.94.10/32
113 set security nat source rule-set source-1083 from zone tester-client
114 set security nat source rule-set source-1083 to zone tester-server
115 set security nat source rule-set source-1083 rule source-1083-1 match source-address 10.6.0.0/15
116 set security nat source rule-set source-1083 rule source-1083-1 match destination-address
    0.0.0.0/0
117 set security nat source rule-set source-1083 rule source-1083-1 then source-nat pool pool-1
118 set security policies from-zone mnha to-zone mnha policy mnha match source-address any
119 set security policies from-zone mnha to-zone mnha policy mnha match destination-address any
120 set security policies from-zone mnha to-zone mnha policy mnha match application any
121 set security policies from-zone mnha to-zone mnha policy mnha then permit
122 set security policies from-zone tester-client to-zone tester-server policy tester-client-1 match
    source-address any
123 set security policies from-zone tester-client to-zone tester-server policy tester-client-1 match
    destination-address any
124 set security policies from-zone tester-client to-zone tester-server policy tester-client-1 match
    application any
125 set security policies from-zone tester-client to-zone tester-server policy tester-client-1 then
    permit
126 set security policies global policy final-1 match source-address any
127 set security policies global policy final-1 match destination-address any
128 set security policies global policy final-1 match application any
129 set security policies global policy final-1 then deny
130 set security zones security-zone mnha interfaces lo0.0 host-inbound-traffic system-services high
    -availability
131 set security zones security-zone mnha interfaces lo0.0 host-inbound-traffic system-services ping
132 set security zones security-zone mnha interfaces lo0.0 host-inbound-traffic protocols bfd
133 set security zones security-zone mnha interfaces xe-0/1/0.0 host-inbound-traffic system-services
    ping
134 set security zones security-zone mnha interfaces xe-0/1/0.0 host-inbound-traffic protocols bgp
135 set security zones security-zone mnha interfaces xe-0/1/0.0 host-inbound-traffic protocols bfd
136 set security zones security-zone mnha interfaces xe-0/1/1.0 host-inbound-traffic system-services
    ping
137 set security zones security-zone mnha interfaces xe-0/1/1.0 host-inbound-traffic protocols bgp
138 set security zones security-zone mnha interfaces xe-0/1/1.0 host-inbound-traffic protocols bfd
139 set security zones security-zone tester-client host-inbound-traffic system-services ping
140 set security zones security-zone tester-client host-inbound-traffic protocols vrrp
141 set security zones security-zone tester-client interfaces ae0.1081
142 set security zones security-zone tester-client interfaces ae0.1082
143 set security zones security-zone tester-client interfaces ae0.1083
144 set security zones security-zone tester-client interfaces ae0.1100
145 set security zones security-zone tester-server host-inbound-traffic system-services ping
146 set security zones security-zone tester-server host-inbound-traffic protocols vrrp
147 set security zones security-zone tester-server interfaces ae0.1091
148 set security zones security-zone tester-server interfaces ae0.1092
149 set security zones security-zone tester-server interfaces ae0.1093
150 set security zones security-zone tester-server interfaces ae0.1094
151 set security zones security-zone tester-server interfaces ae0.1101
152 set security zones security-zone tester-server-vpn host-inbound-traffic system-services ping
153 set security zones security-zone tester-server-vpn host-inbound-traffic protocols vrrp
154 set security zones security-zone tester-server-vpn interfaces ae0.1102
155 set interfaces xe-0/1/0 mtu 2000
156 set interfaces xe-0/1/0 unit 0 family inet address 192.168.0.0/31
157 set interfaces xe-0/1/1 mtu 2000
158 set interfaces xe-0/1/1 unit 0 family inet address 192.168.0.2/31
159 set interfaces xe-0/2/0 ether-options 802.3ad ae0
160 set interfaces xe-0/2/1 ether-options 802.3ad ae0
161 set interfaces xe-0/2/2 ether-options 802.3ad ae0
162 set interfaces xe-0/2/3 ether-options 802.3ad ae0
163 set interfaces ae0 vlan-tagging
164 set interfaces ae0 aggregated-ether-options minimum-links 2
165 set interfaces ae0 aggregated-ether-options lacp active
166 set interfaces ae0 aggregated-ether-options lacp periodic fast
167 set interfaces ae0 unit 1081 vlan-id 1081
168 set interfaces ae0 unit 1081 family inet address 100.64.81.11/24 vrrp-group 1 virtual-address
    100.64.81.10
169 set interfaces ae0 unit 1081 family inet6 address dead:81::11/64 vrrp-inet6-group 1 virtual-
    inet6-address dead:81::10
170 set interfaces ae0 unit 1082 vlan-id 1082
171 set interfaces ae0 unit 1082 family inet address 100.64.82.11/24 vrrp-group 1 virtual-address
    100.64.82.10
172 set interfaces ae0 unit 1082 family inet address 100.64.82.11/24 vrrp-group 1 virtual-address
    100.64.82.20
173 set interfaces ae0 unit 1082 family inet address 100.64.82.11/24 vrrp-group 1 no-accept-data
174 set interfaces ae0 unit 1082 family inet6 address dead:82::11/64 vrrp-inet6-group 1 virtual-
    inet6-address dead:82::10
175 set interfaces ae0 unit 1082 family inet6 address dead:82::11/64 vrrp-inet6-group 1 virtual-
    inet6-address dead:82::20

```



```

176 set interfaces ae0 unit 1082 family inet6 address dead:82::11/64 vrrp-inet6-group 1 no-accept-
177 data
177 set interfaces ae0 unit 1083 vlan-id 1083
178 set interfaces ae0 unit 1083 family inet address 100.64.83.11/24 vrrp-group 1 virtual-address
178 100.64.83.10
179 set interfaces ae0 unit 1083 family inet address 100.64.84.11/24 vrrp-group 2 virtual-address
179 100.64.84.10
180 set interfaces ae0 unit 1083 family inet address 100.64.84.11/24 vrrp-group 2 virtual-address
180 100.64.84.20
181 set interfaces ae0 unit 1083 family inet address 100.64.84.11/24 vrrp-group 2 no-accept-data
182 set interfaces ae0 unit 1083 family inet6 address dead:83::11/64 vrrp-inet6-group 1 virtual-
182 inet6-address dead:83::10
183 set interfaces ae0 unit 1083 family inet6 address dead:84::11/64 vrrp-inet6-group 2 virtual-
183 inet6-address dead:84::10
184 set interfaces ae0 unit 1083 family inet6 address dead:84::11/64 vrrp-inet6-group 2 virtual-
184 inet6-address dead:84::20
185 set interfaces ae0 unit 1083 family inet6 address dead:84::11/64 vrrp-inet6-group 2 no-accept-
185 data
186 set interfaces ae0 unit 1091 vlan-id 1091
187 set interfaces ae0 unit 1091 family inet address 100.64.91.11/24 vrrp-group 1 virtual-address
187 100.64.91.10
188 set interfaces ae0 unit 1091 family inet6 address dead:91::11/64 vrrp-inet6-group 1 virtual-
188 inet6-address dead:91::10
189 set interfaces ae0 unit 1092 vlan-id 1092
190 set interfaces ae0 unit 1092 family inet address 100.64.92.11/24 vrrp-group 1 virtual-address
190 100.64.92.10
191 set interfaces ae0 unit 1092 family inet6 address dead:92::11/64 vrrp-inet6-group 1 virtual-
191 inet6-address dead:92::10
192 set interfaces ae0 unit 1093 vlan-id 1093
193 set interfaces ae0 unit 1093 family inet address 100.64.93.11/24 vrrp-group 1 virtual-address
193 100.64.93.10
194 set interfaces ae0 unit 1093 family inet6 address dead:93::11/64 vrrp-inet6-group 1 virtual-
194 inet6-address dead:93::10
195 set interfaces ae0 unit 1094 vlan-id 1094
196 set interfaces ae0 unit 1094 family inet address 100.64.94.11/24 vrrp-group 1 virtual-address
196 100.64.94.10
197 set interfaces ae0 unit 1094 family inet6 address dead:94::11/64 vrrp-inet6-group 1 virtual-
197 inet6-address dead:94::10
198 set interfaces ae0 unit 1100 vlan-id 1100
199 set interfaces ae0 unit 1100 family inet address 100.64.100.11/24 vrrp-group 1 virtual-address
199 100.64.100.10
200 set interfaces ae0 unit 1101 vlan-id 1101
201 set interfaces ae0 unit 1101 family inet address 100.64.101.11/24 vrrp-group 1 virtual-address
201 100.64.101.10
202 set interfaces ae0 unit 1102 vlan-id 1102
203 set interfaces ae0 unit 1102 family inet address 100.64.102.11/24 vrrp-group 1 virtual-address
203 100.64.102.10
204 set interfaces lo0 unit 0 family inet address 192.168.1.1/32
205 set interfaces lo0 unit 0 family inet address 192.168.1.2/32
206 set policy-options policy-statement export-mnha-icd term icl from interface lo0.0
207 set policy-options policy-statement export-mnha-icd term icl from route-filter 192.168.1.1/32
207 exact
208 set policy-options policy-statement export-mnha-icd term icl then as-path-prepend 65001
209 set policy-options policy-statement export-mnha-icd term icl then accept
210 set policy-options policy-statement export-mnha-icd term icd from interface lo0.0
211 set policy-options policy-statement export-mnha-icd term icd from route-filter 192.168.1.2/32
211 exact
212 set policy-options policy-statement export-mnha-icd term icd then accept
213 set policy-options policy-statement export-mnha-icd term final then reject
214 set policy-options policy-statement export-mnha-icl term icl from interface lo0.0
215 set policy-options policy-statement export-mnha-icl term icl from route-filter 192.168.1.1/32
215 exact
216 set policy-options policy-statement export-mnha-icl term icl then accept
217 set policy-options policy-statement export-mnha-icl term icd from interface lo0.0
218 set policy-options policy-statement export-mnha-icl term icd from route-filter 192.168.1.2/32
218 exact
219 set policy-options policy-statement export-mnha-icl term icd then as-path-prepend 65001
220 set policy-options policy-statement export-mnha-icl term icd then accept
221 set policy-options policy-statement export-mnha-icl term final then reject
222 set routing-instances vr-mnha instance-type virtual-router
223 set routing-instances vr-mnha routing-options autonomous-system 65001
224 set routing-instances vr-mnha protocols bgp group mnha peer-as 65002
225 set routing-instances vr-mnha protocols bgp group mnha bfd-liveness-detection minimum-interval
225 500
226 set routing-instances vr-mnha protocols bgp group mnha bfd-liveness-detection multiplier 3
227 set routing-instances vr-mnha protocols bgp group mnha neighbor 192.168.0.1 export export-mnha-
227 icl
228 set routing-instances vr-mnha protocols bgp group mnha neighbor 192.168.0.3 export export-mnha-
228 icd
229 set routing-instances vr-mnha interface xe-0/1/0.0
230 set routing-instances vr-mnha interface xe-0/1/1.0

```



```

231 set routing-instances vr-mnha interface lo0.0
232 set routing-instances vr-testers instance-type virtual-router
233 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route beef::0/64
    next-hop dead:81::1
234 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route beef
    :0:0:1::0/64 next-hop dead:82::1
235 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route beef
    :0:0:2::0/64 next-hop dead:83::1
236 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route beef
    :0:0:3::0/64 next-hop dead:84::1
237 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route dead
    :0:0:0::0/64 next-hop dead:91::1
238 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route dead
    :0:0:1::0/64 next-hop dead:92::1
239 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route dead
    :0:0:2::0/64 next-hop dead:93::1
240 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route dead
    :0:0:3::0/64 next-hop dead:94::1
241 set routing-instances vr-testers routing-options static route 1.0.0.0/15 next-hop 100.64.91.1
242 set routing-instances vr-testers routing-options static route 1.2.0.0/15 next-hop 100.64.92.1
243 set routing-instances vr-testers routing-options static route 1.4.0.0/15 next-hop 100.64.93.1
244 set routing-instances vr-testers routing-options static route 1.6.0.0/15 next-hop 100.64.94.1
245 set routing-instances vr-testers routing-options static route 10.0.0.0/15 next-hop 100.64.81.1
246 set routing-instances vr-testers routing-options static route 10.2.0.0/15 next-hop 100.64.82.1
247 set routing-instances vr-testers routing-options static route 10.4.0.0/15 next-hop 100.64.83.1
248 set routing-instances vr-testers routing-options static route 10.6.0.0/15 next-hop 100.64.84.1
249 set routing-instances vr-testers routing-options static route 10.100.0.0/15 next-hop
    100.64.100.1
250 set routing-instances vr-testers routing-options static route 10.102.0.0/15 next-hop
    100.64.101.1
251 set routing-instances vr-testers interface ae0.1081
252 set routing-instances vr-testers interface ae0.1082
253 set routing-instances vr-testers interface ae0.1083
254 set routing-instances vr-testers interface ae0.1091
255 set routing-instances vr-testers interface ae0.1092
256 set routing-instances vr-testers interface ae0.1093
257 set routing-instances vr-testers interface ae0.1094
258 set routing-instances vr-testers interface ae0.1100
259 set routing-instances vr-testers interface ae0.1101
260 set routing-instances vr-testers interface ae0.1102
261 set protocols vrrp version-3

```

1.7.2 SRX1600-2

```

1 set version 24.4R2-S3.5
2 set groups vrrp-1 interfaces ae0 unit <*> family inet address <*> vrrp-group <*> priority 150
3 set groups vrrp-1 interfaces ae0 unit <*> family inet address <*> vrrp-group <*> fast-interval
    200
4 set groups vrrp-1 interfaces ae0 unit <*> family inet address <*> vrrp-group <*> accept-data
5 set groups vrrp-1 interfaces ae0 unit <*> family inet address <*> vrrp-group <*> track route
    100.64.0.1/32 routing-instance vr-mnha priority-cost 100
6 set groups vrrp-1 interfaces ae0 unit <*> family inet6 address <*> vrrp-inet6-group <*> priority
    150
7 set groups vrrp-1 interfaces ae0 unit <*> family inet6 address <*> vrrp-inet6-group <*> fast-
    interval 200
8 set groups vrrp-1 interfaces ae0 unit <*> family inet6 address <*> vrrp-inet6-group <*> accept-
    data
9 set groups vrrp-1 interfaces ae0 unit <*> family inet6 address <*> vrrp-inet6-group <*> track
    route 100.64.0.1/32 routing-instance vr-mnha priority-cost 100
10 set groups ipsec-1 chassis high-availability services-redundancy-group 1 prefix-list
    ipsec_termination routing-instance vr-testers
11 set groups ipsec-1 chassis high-availability services-redundancy-group 1 managed-services ipsec
12 set groups ipsec-1 chassis high-availability services-redundancy-group 1 process-packet-on-
    backup
13 set groups ipsec-1 security ike proposal ike-proposal-1 authentication-method pre-shared-keys
14 set groups ipsec-1 security ike proposal ike-proposal-1 dh-group group20
15 set groups ipsec-1 security ike proposal ike-proposal-1 encryption-algorithm aes-256-gcm
16 set groups ipsec-1 security ike proposal ike-proposal-1 lifetime-seconds 86400
17 set groups ipsec-1 security ike policy ike-policy-1 proposals ike-proposal-1
18 set groups ipsec-1 security ike policy ike-policy-1 pre-shared-key ascii-text <-----PSK-HERE
    ----->
19 set groups ipsec-1 security ike gateway ike-gw-1 ike-policy ike-policy-1
20 set groups ipsec-1 security ike gateway ike-gw-1 address 100.64.102.20
21 set groups ipsec-1 security ike gateway ike-gw-1 dead-peer-detection
22 set groups ipsec-1 security ike gateway ike-gw-1 external-interface ae0.1102
23 set groups ipsec-1 security ike gateway ike-gw-1 local-address 100.64.102.10
24 set groups ipsec-1 security ike gateway ike-gw-1 version v2-only

```

```

25 set groups ipsec-1 security ike gateway ike-gw-1 ppk-profile quantum
26 set groups ipsec-1 security ipsec proposal ipsec-proposal-1 encryption-algorithm aes-256-gcm
27 set groups ipsec-1 security ipsec proposal ipsec-proposal-1 lifetime-seconds 3600
28 set groups ipsec-1 security ipsec policy ipsec-policy-1 perfect-forward-secrecy keys group20
29 set groups ipsec-1 security ipsec policy ipsec-policy-1 proposals ipsec-proposal-1
30 set groups ipsec-1 security ipsec vpn vpn-1 bind-interface st0.0
31 set groups ipsec-1 security ipsec vpn vpn-1 ike gateway ike-gw-1
32 set groups ipsec-1 security ipsec vpn vpn-1 ike ipsec-policy ipsec-policy-1
33 set groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn policy
  tester-client-vpn-1 match source-address any
34 set groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn policy
  tester-client-vpn-1 match destination-address any
35 set groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn policy
  tester-client-vpn-1 match application any
36 set groups ipsec-1 security policies from-zone tester-client to-zone tester-server-vpn policy
  tester-client-vpn-1 then permit
37 set groups ipsec-1 security zones security-zone tester-server-vpn host-inbound-traffic system-
  services ike
38 set groups ipsec-1 security zones security-zone tester-server-vpn interfaces st0.0
39 set groups ipsec-1 security key-manager profiles quantum static key-id ascii-text <-----KEY-ID-
  HERE----->
40 set groups ipsec-1 security key-manager profiles quantum static key ascii-text <-----32-
  character+-KEY-HERE----->
41 set groups ipsec-1 interfaces st0 unit 0 family inet
42 set groups ipsec-1 policy-options prefix-list ipsec_termination 100.64.102.10/32
43 set groups ipsec-1 routing-instances vr-testers routing-options static route 10.104.0.0/15 next-
  hop st0.0
44 set groups ipsec-1 routing-instances vr-testers interface st0.0
45 set groups icl-ipsec chassis high-availability peer-id 1 vpn-profile icl
46 set groups icl-ipsec security ike proposal ike-icl-proposal authentication-method pre-shared-
  keys
47 set groups icl-ipsec security ike proposal ike-icl-proposal dh-group group20
48 set groups icl-ipsec security ike proposal ike-icl-proposal encryption-algorithm aes-256-gcm
49 set groups icl-ipsec security ike proposal ike-icl-proposal lifetime-seconds 28800
50 set groups icl-ipsec security ike policy ike-icl-policy proposals ike-icl-proposal
51 set groups icl-ipsec security ike policy ike-icl-policy pre-shared-key ascii-text <-----PSK-HERE
  ----->
52 set groups icl-ipsec security ike gateway icl ike-policy ike-icl-policy
53 set groups icl-ipsec security ike gateway icl version v2-only
54 set groups icl-ipsec security ipsec proposal ipsec-icl-proposal encryption-algorithm aes-256-gcm
55 set groups icl-ipsec security ipsec proposal ipsec-icl-proposal lifetime-seconds 3600
56 set groups icl-ipsec security ipsec policy ipsec-icl-policy perfect-forward-secrecy keys group20
57 set groups icl-ipsec security ipsec policy ipsec-icl-policy proposals ipsec-icl-proposal
58 set groups icl-ipsec security ipsec vpn icl ha-link-encryption
59 set groups icl-ipsec security ipsec vpn icl ike gateway icl
60 set groups icl-ipsec security ipsec vpn icl ike ipsec-policy ipsec-icl-policy
61 set groups icl-ipsec security zones security-zone mnha interfaces lo0.0 host-inbound-traffic
  system-services ike
62 set apply-groups vrrp-1
63 set apply-groups ipsec-1
64 set apply-groups icl-ipsec
65 set system host-name srx1600-2
66 set system services netconf ssh
67 set system services ssh root-login allow
68 set system services ssh sftp-server
69 set system services ssh client-alive-interval 120
70 set system time-zone Europe/Amsterdam
71 set system name-server 172.30.207.10
72 set system syslog file messages any any
73 set system syslog file messages archive size 5m
74 set system syslog file messages archive files 4
75 set system syslog time-format millisecond
76 set system ntp server 172.30.207.10
77 set chassis aggregated-devices ethernet device-count 1
78 set chassis fpc 0 pic 1 pic-mode 1G10G
79 set chassis high-availability local-id 2
80 set chassis high-availability local-id local-ip 192.168.2.1
81 set chassis high-availability local-id local-forwarding-ip 192.168.2.2
82 set chassis high-availability peer-id 1 peer-ip 192.168.1.1
83 set chassis high-availability peer-id 1 interface lo0.0
84 set chassis high-availability peer-id 1 routing-instance vr-mnha
85 set chassis high-availability peer-id 1 peer-forwarding-ip 192.168.1.2
86 set chassis high-availability peer-id 1 peer-forwarding-ip interface lo0.0
87 set chassis high-availability peer-id 1 peer-forwarding-ip liveness-detection minimum-interval
  1000
88 set chassis high-availability peer-id 1 peer-forwarding-ip liveness-detection multiplier 3
89 set chassis high-availability peer-id 1 liveness-detection minimum-interval 1000
90 set chassis high-availability peer-id 1 liveness-detection multiplier 3
91 set chassis high-availability services-redundancy-group 1 deployment-type routing
92 set chassis high-availability services-redundancy-group 1 peer-id 1
93 set chassis high-availability services-redundancy-group 1 activeness-probe dest-ip 100.64.81.20

```

```

94 set chassis high-availability services-redundancy-group 1 activeness-probe dest-ip src-ip
100.64.81.10
95 set chassis high-availability services-redundancy-group 1 activeness-probe dest-ip routing-
instance vr-testers
96 set chassis high-availability services-redundancy-group 1 monitor monitor-object interface-1
object-threshold 100
97 set chassis high-availability services-redundancy-group 1 monitor monitor-object interface-1
interface threshold 100
98 set chassis high-availability services-redundancy-group 1 monitor monitor-object interface-1
interface interface-name ae0 weight 100
99 set chassis high-availability services-redundancy-group 1 monitor srg-threshold 100
100 set chassis high-availability services-redundancy-group 1 active-signal-route 100.64.0.1
101 set chassis high-availability services-redundancy-group 1 active-signal-route routing-instance
vr-mnha
102 set chassis high-availability services-redundancy-group 1 backup-signal-route 100.64.0.0
103 set chassis high-availability services-redundancy-group 1 backup-signal-route routing-instance
vr-mnha
104 set chassis high-availability services-redundancy-group 1 activeness-priority 100
105 set security alg msrpc disable
106 set security alg sunrpc disable
107 set security alg talk disable
108 set security alg tftp disable
109 set security alg pptp disable
110 set security flow drop-flow max-sessions 0
111 set security flow tcp-session strict-syn-check
112 set security nat source pool pool-1 address 100.64.94.10/32
113 set security nat source rule-set source-1083 from zone tester-client
114 set security nat source rule-set source-1083 to zone tester-server
115 set security nat source rule-set source-1083 rule source-1083-1 match source-address 10.6.0.0/15
116 set security nat source rule-set source-1083 rule source-1083-1 match destination-address
0.0.0.0/0
117 set security nat source rule-set source-1083 rule source-1083-1 then source-nat pool pool-1
118 set security policies from-zone mnha to-zone mnha policy mnha match source-address any
119 set security policies from-zone mnha to-zone mnha policy mnha match destination-address any
120 set security policies from-zone mnha to-zone mnha policy mnha match application any
121 set security policies from-zone mnha to-zone mnha policy mnha then permit
122 set security policies from-zone tester-client to-zone tester-server policy tester-client-1 match
source-address any
123 set security policies from-zone tester-client to-zone tester-server policy tester-client-1 match
destination-address any
124 set security policies from-zone tester-client to-zone tester-server policy tester-client-1 match
application any
125 set security policies from-zone tester-client to-zone tester-server policy tester-client-1 then
permit
126 set security policies global policy final-1 match source-address any
127 set security policies global policy final-1 match destination-address any
128 set security policies global policy final-1 match application any
129 set security policies global policy final-1 then deny
130 set security zones security-zone mnha interfaces lo0.0 host-inbound-traffic system-services high
-availability
131 set security zones security-zone mnha interfaces lo0.0 host-inbound-traffic system-services ping
132 set security zones security-zone mnha interfaces lo0.0 host-inbound-traffic protocols bfd
133 set security zones security-zone mnha interfaces xe-0/1/0.0 host-inbound-traffic system-services
ping
134 set security zones security-zone mnha interfaces xe-0/1/0.0 host-inbound-traffic protocols bgp
135 set security zones security-zone mnha interfaces xe-0/1/0.0 host-inbound-traffic protocols bfd
136 set security zones security-zone mnha interfaces xe-0/1/1.0 host-inbound-traffic system-services
ping
137 set security zones security-zone mnha interfaces xe-0/1/1.0 host-inbound-traffic protocols bgp
138 set security zones security-zone mnha interfaces xe-0/1/1.0 host-inbound-traffic protocols bfd
139 set security zones security-zone tester-client host-inbound-traffic system-services ping
140 set security zones security-zone tester-client host-inbound-traffic protocols vrrp
141 set security zones security-zone tester-client interfaces ae0.1081
142 set security zones security-zone tester-client interfaces ae0.1082
143 set security zones security-zone tester-client interfaces ae0.1083
144 set security zones security-zone tester-client interfaces ae0.1100
145 set security zones security-zone tester-server host-inbound-traffic system-services ping
146 set security zones security-zone tester-server host-inbound-traffic protocols vrrp
147 set security zones security-zone tester-server interfaces ae0.1091
148 set security zones security-zone tester-server interfaces ae0.1092
149 set security zones security-zone tester-server interfaces ae0.1093
150 set security zones security-zone tester-server interfaces ae0.1094
151 set security zones security-zone tester-server interfaces ae0.1101
152 set security zones security-zone tester-server-vpn host-inbound-traffic system-services ping
153 set security zones security-zone tester-server-vpn host-inbound-traffic protocols vrrp
154 set security zones security-zone tester-server-vpn interfaces ae0.1102
155 set interfaces xe-0/1/0 mtu 2000
156 set interfaces xe-0/1/0 unit 0 family inet address 192.168.0.1/31
157 set interfaces xe-0/1/1 mtu 2000
158 set interfaces xe-0/1/1 unit 0 family inet address 192.168.0.3/31
159 set interfaces xe-0/2/0 ether-options 802.3ad ae0

```

```

160 set interfaces xe-0/2/1 ether-options 802.3ad ae0
161 set interfaces xe-0/2/2 ether-options 802.3ad ae0
162 set interfaces xe-0/2/3 ether-options 802.3ad ae0
163 set interfaces ae0 vlan-tagging
164 set interfaces ae0 aggregated-ether-options minimum-links 2
165 set interfaces ae0 aggregated-ether-options lacp active
166 set interfaces ae0 aggregated-ether-options lacp periodic fast
167 set interfaces ae0 unit 1081 vlan-id 1081
168 set interfaces ae0 unit 1081 family inet address 100.64.81.12/24 vrrp-group 1 virtual-address
100.64.81.10
169 set interfaces ae0 unit 1081 family inet6 address dead:81::12/64 vrrp-inet6-group 1 virtual-
inet6-address dead:81::10
170 set interfaces ae0 unit 1082 vlan-id 1082
171 set interfaces ae0 unit 1082 family inet address 100.64.82.12/24 vrrp-group 1 virtual-address
100.64.82.10
172 set interfaces ae0 unit 1082 family inet address 100.64.82.12/24 vrrp-group 1 virtual-address
100.64.82.20
173 set interfaces ae0 unit 1082 family inet address 100.64.82.12/24 vrrp-group 1 no-accept-data
174 set interfaces ae0 unit 1082 family inet6 address dead:82::12/64 vrrp-inet6-group 1 virtual-
inet6-address dead:82::10
175 set interfaces ae0 unit 1082 family inet6 address dead:82::12/64 vrrp-inet6-group 1 virtual-
inet6-address dead:82::20
176 set interfaces ae0 unit 1082 family inet6 address dead:82::12/64 vrrp-inet6-group 1 no-accept-
data
177 set interfaces ae0 unit 1083 vlan-id 1083
178 set interfaces ae0 unit 1083 family inet address 100.64.83.12/24 vrrp-group 1 virtual-address
100.64.83.10
179 set interfaces ae0 unit 1083 family inet address 100.64.84.12/24 vrrp-group 2 virtual-address
100.64.84.10
180 set interfaces ae0 unit 1083 family inet address 100.64.84.12/24 vrrp-group 2 virtual-address
100.64.84.20
181 set interfaces ae0 unit 1083 family inet address 100.64.84.12/24 vrrp-group 2 no-accept-data
182 set interfaces ae0 unit 1083 family inet6 address dead:83::12/64 vrrp-inet6-group 1 virtual-
inet6-address dead:83::10
183 set interfaces ae0 unit 1083 family inet6 address dead:84::12/64 vrrp-inet6-group 2 virtual-
inet6-address dead:84::10
184 set interfaces ae0 unit 1083 family inet6 address dead:84::12/64 vrrp-inet6-group 2 virtual-
inet6-address dead:84::20
185 set interfaces ae0 unit 1083 family inet6 address dead:84::12/64 vrrp-inet6-group 2 no-accept-
data
186 set interfaces ae0 unit 1091 vlan-id 1091
187 set interfaces ae0 unit 1091 family inet address 100.64.91.12/24 vrrp-group 1 virtual-address
100.64.91.10
188 set interfaces ae0 unit 1091 family inet6 address dead:91::12/64 vrrp-inet6-group 1 virtual-
inet6-address dead:91::10
189 set interfaces ae0 unit 1092 vlan-id 1092
190 set interfaces ae0 unit 1092 family inet address 100.64.92.12/24 vrrp-group 1 virtual-address
100.64.92.10
191 set interfaces ae0 unit 1092 family inet6 address dead:92::12/64 vrrp-inet6-group 1 virtual-
inet6-address dead:92::10
192 set interfaces ae0 unit 1093 vlan-id 1093
193 set interfaces ae0 unit 1093 family inet address 100.64.93.12/24 vrrp-group 1 virtual-address
100.64.93.10
194 set interfaces ae0 unit 1093 family inet6 address dead:93::12/64 vrrp-inet6-group 1 virtual-
inet6-address dead:93::10
195 set interfaces ae0 unit 1094 vlan-id 1094
196 set interfaces ae0 unit 1094 family inet address 100.64.94.12/24 vrrp-group 1 virtual-address
100.64.94.10
197 set interfaces ae0 unit 1094 family inet6 address dead:94::12/64 vrrp-inet6-group 1 virtual-
inet6-address dead:94::10
198 set interfaces ae0 unit 1100 vlan-id 1100
199 set interfaces ae0 unit 1100 family inet address 100.64.100.12/24 vrrp-group 1 virtual-address
100.64.100.10
200 set interfaces ae0 unit 1101 vlan-id 1101
201 set interfaces ae0 unit 1101 family inet address 100.64.101.12/24 vrrp-group 1 virtual-address
100.64.101.10
202 set interfaces ae0 unit 1102 vlan-id 1102
203 set interfaces ae0 unit 1102 family inet address 100.64.102.12/24 vrrp-group 1 virtual-address
100.64.102.10
204 set interfaces lo0 unit 0 family inet address 192.168.2.1/32
205 set interfaces lo0 unit 0 family inet address 192.168.2.2/32
206 set policy-options policy-statement export-mnha-icd term icl from interface lo0.0
207 set policy-options policy-statement export-mnha-icd term icl from route-filter 192.168.2.1/32
exact
208 set policy-options policy-statement export-mnha-icd term icl then as-path-prepend 65002
209 set policy-options policy-statement export-mnha-icd term icl then accept
210 set policy-options policy-statement export-mnha-icd term icd from interface lo0.0
211 set policy-options policy-statement export-mnha-icd term icd from route-filter 192.168.2.2/32
exact
212 set policy-options policy-statement export-mnha-icd term icd then accept
213 set policy-options policy-statement export-mnha-icd term final then reject

```

```

214 set policy-options policy-statement export-mnha-icl term icl from interface lo0.0
215 set policy-options policy-statement export-mnha-icl term icl from route-filter 192.168.2.1/32
    exact
216 set policy-options policy-statement export-mnha-icl term icl then accept
217 set policy-options policy-statement export-mnha-icl term icd from interface lo0.0
218 set policy-options policy-statement export-mnha-icl term icd from route-filter 192.168.2.2/32
    exact
219 set policy-options policy-statement export-mnha-icl term icd then as-path-prepend 65002
220 set policy-options policy-statement export-mnha-icl term icd then accept
221 set policy-options policy-statement export-mnha-icl term final then reject
222 set routing-instances vr-mnha instance-type virtual-router
223 set routing-instances vr-mnha routing-options autonomous-system 65002
224 set routing-instances vr-mnha protocols bgp group mnha peer-as 65001
225 set routing-instances vr-mnha protocols bgp group mnha bfd-liveness-detection minimum-interval
    500
226 set routing-instances vr-mnha protocols bgp group mnha bfd-liveness-detection multiplier 3
227 set routing-instances vr-mnha protocols bgp group mnha neighbor 192.168.0.0 export export-mnha-
    icl
228 set routing-instances vr-mnha protocols bgp group mnha neighbor 192.168.0.2 export export-mnha-
    icd
229 set routing-instances vr-mnha interface xe-0/1/0.0
230 set routing-instances vr-mnha interface xe-0/1/1.0
231 set routing-instances vr-mnha interface lo0.0
232 set routing-instances vr-testers instance-type virtual-router
233 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route beef::0/64
    next-hop dead:81::1
234 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route beef
    :0:0:1::0/64 next-hop dead:82::1
235 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route beef
    :0:0:2::0/64 next-hop dead:83::1
236 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route beef
    :0:0:3::0/64 next-hop dead:84::1
237 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route dead
    :0:0:0::0/64 next-hop dead:91::1
238 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route dead
    :0:0:1::0/64 next-hop dead:92::1
239 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route dead
    :0:0:2::0/64 next-hop dead:93::1
240 set routing-instances vr-testers routing-options rib vr-testers.inet6.0 static route dead
    :0:0:3::0/64 next-hop dead:94::1
241 set routing-instances vr-testers routing-options static route 1.0.0.0/15 next-hop 100.64.91.1
242 set routing-instances vr-testers routing-options static route 1.2.0.0/15 next-hop 100.64.92.1
243 set routing-instances vr-testers routing-options static route 1.4.0.0/15 next-hop 100.64.93.1
244 set routing-instances vr-testers routing-options static route 1.6.0.0/15 next-hop 100.64.94.1
245 set routing-instances vr-testers routing-options static route 10.0.0.0/15 next-hop 100.64.81.1
246 set routing-instances vr-testers routing-options static route 10.2.0.0/15 next-hop 100.64.82.1
247 set routing-instances vr-testers routing-options static route 10.4.0.0/15 next-hop 100.64.83.1
248 set routing-instances vr-testers routing-options static route 10.6.0.0/15 next-hop 100.64.84.1
249 set routing-instances vr-testers routing-options static route 10.100.0.0/15 next-hop
    100.64.100.1
250 set routing-instances vr-testers routing-options static route 10.102.0.0/15 next-hop
    100.64.101.1
251 set routing-instances vr-testers interface ae0.1081
252 set routing-instances vr-testers interface ae0.1082
253 set routing-instances vr-testers interface ae0.1083
254 set routing-instances vr-testers interface ae0.1091
255 set routing-instances vr-testers interface ae0.1092
256 set routing-instances vr-testers interface ae0.1093
257 set routing-instances vr-testers interface ae0.1094
258 set routing-instances vr-testers interface ae0.1100
259 set routing-instances vr-testers interface ae0.1101
260 set routing-instances vr-testers interface ae0.1102
261 set protocols vrrp version-3

```

1.7.3 vSRX-1

```

1 set version 24.4R2-S3.5
2 set groups ipsec-1 security ike proposal ike-proposal-1 authentication-method pre-shared-keys
3 set groups ipsec-1 security ike proposal ike-proposal-1 dh-group group20
4 set groups ipsec-1 security ike proposal ike-proposal-1 encryption-algorithm aes-256-gcm
5 set groups ipsec-1 security ike proposal ike-proposal-1 lifetime-seconds 86400
6 set groups ipsec-1 security ike policy ike-policy-1 proposals ike-proposal-1
7 set groups ipsec-1 security ike policy ike-policy-1 pre-shared-key ascii-text <-----PSK-HERE
    ----->
8 set groups ipsec-1 security ike gateway ike-gw-1 ike-policy ike-policy-1
9 set groups ipsec-1 security ike gateway ike-gw-1 address 100.64.102.10
10 set groups ipsec-1 security ike gateway ike-gw-1 external-interface ge-0/0/0.0

```



```

11 set groups ipsec-1 security ike gateway ike-gw-1 version v2-only
12 set groups ipsec-1 security ike gateway ike-gw-1 ppk-profile quantum
13 set groups ipsec-1 security ipsec proposal ipsec-proposal-1 encryption-algorithm aes-256-gcm
14 set groups ipsec-1 security ipsec proposal ipsec-proposal-1 lifetime-seconds 3600
15 set groups ipsec-1 security ipsec policy ipsec-policy-1 perfect-forward-secrecy keys group20
16 set groups ipsec-1 security ipsec policy ipsec-policy-1 proposals ipsec-proposal-1
17 set groups ipsec-1 security ipsec vpn vpn-1 bind-interface st0.0
18 set groups ipsec-1 security ipsec vpn vpn-1 ike gateway ike-gw-1
19 set groups ipsec-1 security ipsec vpn vpn-1 ike ipsec-policy ipsec-policy-1
20 set groups ipsec-1 security ipsec vpn vpn-1 establish-tunnels immediately
21 set groups ipsec-1 security policies from-zone tester-server to-zone tester-server-vpn policy
  tester-server_tester-server-vpn-1 match source-address any
22 set groups ipsec-1 security policies from-zone tester-server to-zone tester-server-vpn policy
  tester-server_tester-server-vpn-1 match destination-address any
23 set groups ipsec-1 security policies from-zone tester-server to-zone tester-server-vpn policy
  tester-server_tester-server-vpn-1 match application any
24 set groups ipsec-1 security policies from-zone tester-server to-zone tester-server-vpn policy
  tester-server_tester-server-vpn-1 then permit
25 set groups ipsec-1 security policies from-zone tester-server-vpn to-zone tester-server policy
  tester-server-vpn_tester-server-1 match source-address any
26 set groups ipsec-1 security policies from-zone tester-server-vpn to-zone tester-server policy
  tester-server-vpn_tester-server-1 match destination-address any
27 set groups ipsec-1 security policies from-zone tester-server-vpn to-zone tester-server policy
  tester-server-vpn_tester-server-1 match application any
28 set groups ipsec-1 security policies from-zone tester-server-vpn to-zone tester-server policy
  tester-server-vpn_tester-server-1 then permit
29 set groups ipsec-1 security zones security-zone tester-server-vpn interfaces ge-0/0/0.0 host-
  inbound-traffic system-services ike
30 set groups ipsec-1 security zones security-zone tester-server-vpn interfaces st0.0
31 set groups ipsec-1 security key-manager profiles quantum static key-id ascii-text <-----KEY-ID-
  HERE----->
32 set groups ipsec-1 security key-manager profiles quantum static key ascii-text <-----32-
  character+-KEY-HERE----->
33 set groups ipsec-1 interfaces st0 unit 0 family inet
34 set groups ipsec-1 routing-options static route 10.100.0.0/15 next-hop st0.0
35 set groups ipsec-1 routing-options static route 10.104.0.0/15 next-hop 100.64.101.1
36 set apply-groups ipsec-1
37 set system host-name vSRX-1
38 set system services ssh root-login allow
39 set system services ssh sftp-server
40 set system services ssh client-alive-interval 120
41 set system syslog file messages any any
42 set system syslog file messages archive size 5m
43 set system syslog file messages archive files 4
44 set security alg h323 disable
45 set security alg mgcp disable
46 set security alg msrpc disable
47 set security alg sunrpc disable
48 set security alg rtsp disable
49 set security alg sccp disable
50 set security alg sip disable
51 set security alg talk disable
52 set security alg tftp disable
53 set security alg pptp disable
54 set security flow drop-flow max-sessions 0
55 set security zones security-zone tester-server-vpn interfaces ge-0/0/0.0 host-inbound-traffic
  system-services ping
56 set security zones security-zone tester-server interfaces ge-0/0/1.0 host-inbound-traffic system
  -services ping
57 set interfaces ge-0/0/0 unit 0 family inet address 100.64.102.20/24
58 set interfaces ge-0/0/1 unit 0 family inet address 100.64.101.20/24

```

1.8 Appendix 2: Encrypted ICL

MNHA can secure the ICL link using IPSEC, which is mandatory in real-world VPN setups. IPSEC also helps distribute ICL traffic on the receiving side because the ICL VPN is designed to use a 1:1 IPSEC SA between local and remote PFE cores - otherwise a single CPU core on the peer can be overwhelmed by receiving clear-text RTOs due to L3 based hashing when there are many connection setups (applies mostly to larger appliances like the SRX4300).

To enable ICL IPSEC encryption, it's useful to use a Junos group to keep the settings in one place with an easy toggle to activate or deactivate (ICL communication will be temporarily disturbed). On SRX1600-1, the

noteworthy setting is the *ha-link-encryption* knob:

```

1  edit groups icl-ipsec
2
3  set security ike proposal ike-icl-proposal authentication-method pre-shared-keys
4  set security ike proposal ike-icl-proposal dh-group group20
5  set security ike proposal ike-icl-proposal encryption-algorithm aes-256-gcm
6  set security ike proposal ike-icl-proposal lifetime-seconds 28800
7
8  set security ike policy ike-icl-policy proposals ike-icl-proposal
9  set security ike policy ike-icl-policy pre-shared-key ascii-text <-----PSK-HERE----->
10
11 set security ike gateway icl ike-policy ike-icl-policy
12 set security ike gateway icl version v2-only
13
14 set security ipsec proposal ipsec-icl-proposal encryption-algorithm aes-256-gcm
15 set security ipsec proposal ipsec-icl-proposal lifetime-seconds 3600
16
17 set security ipsec policy ipsec-icl-policy perfect-forward-secrecy keys group20
18 set security ipsec policy ipsec-icl-policy proposals ipsec-icl-proposal
19
20 set security ipsec vpn icl ha-link-encryption
21 set security ipsec vpn icl ike gateway icl
22 set security ipsec vpn icl ike ipsec-policy ipsec-icl-policy
23
24 set security zones security-zone mnha interfaces lo0.0 host-inbound-traffic system-services ike

```

Note: also PKI based VPN for ICL is supported.

The only nuance between SRX1600-1 and SRX1600-2 is enabling the VPN for the peer. On SRX1600-1:

```

1  set groups icl-ipsec chassis high-availability peer-id 2 vpn-profile icl

```

And on SRX1600-2:

```

1  set groups icl-ipsec chassis high-availability peer-id 1 vpn-profile icl

```

To apply the group settings:

```

1  top set apply-groups icl-ipsec

```

After reconfiguration to an encrypted ICL, MNHA uses an internally hard-coded tunnel interface, addressing, and routing instance:

```

1  show chassis high-availability information detail
2
3  <SNIP>
4  HA Peer Information:
5
6  Peer-ID: 2 IP address: 192.168.2.1 Interface: lo0.0
7  Routing Instance: vr-mnha
8  Encrypted: YES Conn State: UP
9  Cold Sync Status: COMPLETE
10 Peer Forwarding IP: 192.168.2.2 Interface: lo0.0
11 Peer ICD Conn State: UP
12 Internal Interface: st0.16000
13 Internal Local-IP: 180.100.1.1
14 Internal Peer-IP: 180.100.1.2
15 Internal Routing-instance: __juniper_private1__
16 <SNIP>

```

Note: the internal IP addresses also appear in the BFD sessions (multi-hop); e.g., on SRX1600-1:

```

1  show bfd session
2
3                                     Detect Transmit
4 Address State Interface Time Interval Multiplier
5 180.100.1.2 Up 3.000 1.000 3
6 192.168.0.1 Up xe-0/1/0.0 1.500 0.500 3
7 192.168.0.3 Up xe-0/1/1.0 1.500 0.500 3
8 192.168.2.2 Up 3.000 1.000 3

```


The IKE and IPSEC SAs (an SA pair for every PFE flow core) can be viewed as follows:

```

1  show security ike security-associations ha-link-encryption
2
3  Index State Initiator cookie Responder cookie Mode Remote Address
4  16776192 UP 378c49ae3d96bb66 f19d71c8d92213de IKEv2 192.168.2.1
5
6  show security ipsec security-associations ha-link-encryption
7
8  Total active tunnels: 1 Total IPsec sas: 4
9  ID Algorithm SPI Life:sec/kb Mon lsys Port Gateway
10 <495001 ESP:aes-gcm-256/aes256-gcm 0x0009a68e 2128/ unlim - root 500 192.168.2.1
11 >495001 ESP:aes-gcm-256/aes256-gcm 0x000fa97b 2128/ unlim - root 500 192.168.2.1
12 <495001 ESP:aes-gcm-256/aes256-gcm 0x0409a68e 2128/ unlim - root 500 192.168.2.1
13 >495001 ESP:aes-gcm-256/aes256-gcm 0x040fa97b 2128/ unlim - root 500 192.168.2.1
14 <495001 ESP:aes-gcm-256/aes256-gcm 0x0809a68e 2128/ unlim - root 500 192.168.2.1
15 >495001 ESP:aes-gcm-256/aes256-gcm 0x080fa97b 2128/ unlim - root 500 192.168.2.1
16 <495001 ESP:aes-gcm-256/aes256-gcm 0x0c09a68e 2128/ unlim - root 500 192.168.2.1
17 >495001 ESP:aes-gcm-256/aes256-gcm 0x0c0fa97b 2128/ unlim - root 500 192.168.2.1

```

1.9 Conclusion

This TechPost shows that combining VRRP with SRX MNHA in routing mode is a practical alternative to native MNHA Switching and Hybrid L2 VIP mechanisms. By separating responsibilities - MNHA for state synchronization and activeness steering, and VRRP for flexible L2 VIP ownership - the design overcomes key MNHA limitations such as single VIP per interface, lack of dualstack support (prior to 26.2R1), nonoverlapping VIP restrictions, and multiaddress requirements.

Lab validation confirms that the MNHA-VRRP coupling delivers predictable and fast failover for both cleartext and IPSEC traffic under administrative events, sudden outages, sustained load and overload.

While more complex and requiring careful qualification of timers, scale, and LACP behavior, this approach provides a robust and highly flexible HA model for advanced L2/L3 and dualstack deployments on modern SRX platforms where traditional chassis-cluster is no longer an option.

1.10 Acknowledgements

All the brilliant people involved in exploring SRX MNHA/VRRP: Steven Jacques, Matthijs Nagel, Venu Nalladhimmu, Laurent Paumelle and others. Those who provided valuable feedback: Mark Barrett (big shout-out for his exceptional support), Henri Kalliosaari, Matthijs Nagel, Laurent Paumelle and James Rathbun. Special thanks to Nicolas Fevrier for tirelessly overseeing the Tech Posts site and handling all publishing tasks. There would be no testing without the support of the Amsterdam Proof of Concepts lab. Finally, kudos to Juniper SRX engineering and product-line management for pushing the envelope.

1.11 Useful Links

- <https://www.juniper.net/documentation/us/en/software/junos/high-availability/topics/concept/mnha-overview.html>
- <https://www.juniper.net/documentation/us/en/software/junos/high-availability/topics/topic-map/vrrp-understanding.html>
- <https://juniper.github.io/techposts/multi-node-high-availability-basics/article>
- <https://juniper.github.io/techposts/dhcp-on-mnha-back-to-basics/article>
- <https://juniper.github.io/techposts/hybrid-mnha-with-ebgp/article>

- <https://juniper.github.io/techposts/mnha-ipsec-and-multiple-routing-instances/article>
- <https://juniper.github.io/techposts/srx-from-chassis-cluster-to-mnha/article>

1.12 Glossary

- A/P:Active/Passive
- ALG:Application Layer Gateway
- ARP:Address Resolution Protocol
- AS:Autonomous System
- BFD:Bi-Directional Forwarding Detection
- BGP:Border Gateway Protocol
- CPS:Connections Per Second
- CPU:Central Processing Unit
- DNS:Domain Name System
- FTP:File Transfer Protocol
- HA:High Availability
- HTTP:Hyper Text Transfer Protocol
- ICD:Inter-Chassis-Datapath (MNHA)
- ICL:Inter-Chassis Link (MNHA)
- ICMP:Internet Control Message Protocol
- IKE:Internet Key Exchange
- IP:Internet Protocol
- IPSEC:IP Security
- IRB:Integrated Routing and Bridging (interface)
- L2:Layer 2
- LACP:Ling Aggregation Control Protocol
- MAC:Media Access Control
- MNHA:Multi Node High Availability
- MPPS:Millions Packets Per Second
- MTU:Maximum Transmission Unit
- NDP:Neighbor Discovery Protocol (IPv6)
- NTP:Network Time Protocol
- PFE:Packet Forwarding Engine
- PIC:Physical Interface Card (Junos)
- PPK:Post-quantum Pre-shared Keys
- PPS:Packets Per Second
- RTO:Real Time Object (SRX HA)
- SA:Security Association (IPSEC)
- SRG:Services Redundancy Group (MNHA)
- TCP:Transmission Control Protocol
- TPS:Transactions Per Second
- UDP:User Datagram Protocol
- VIP:Virtual IP (address)
- VLAN:Virtual Local Area Network
- VPN:Virtual Private Network

- VRRP:Virtual Router Redundancy Protocol



DC – AI – Routing – Switching – Security – Wireless – Automation

© Copyright 2026 Hewlett Packard Enterprise Development LP. The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty.

Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Trademark acknowledgments, if needed. All third-party marks are the property of their respective owners.

HEWLETT PACKARD ENTERPRISE

HPE.com

