



DC – AI – Routing – Switching – Security – Wireless – Automation

Building Virtual Fabrics with vJunos-switch and Containerlab

An HPE Juniper Networking Techpost

Aninda Chatterjee

2023-10-27

HPE



Contents

1	Building Virtual Fabrics with vJunos-switch and Containerlab	3
1.1	What is Containerlab and vJunos-switch?	3
1.2	Installing Containerlab and Building the vJunos-switch Container	3
1.3	Building a Network Topology with vJunos-switch and Containerlab	5
1.4	Integrating a vJunos-switch Containerlab Based Network Topology with Juniper Apstra . .	9
1.5	Summary	17
1.6	Useful links	17
1.7	Acknowledgments	17

List of Figures

1	Network topology deployed via Containerlab	10
2	vJunos-switch Device Profile in Apstra 4.2.0	10
3	vJunos-switch Logical Device created in Apstra	11
4	vJunos-switch Interface Map created in Apstra	11
5	vJunos-switch based Rack Type created in Apstra	12
6	vJunos-switch based Template created in Apstra	12
7	vJunos-switch based Blueprint created in Apstra	13
8	Staged Blueprint in Apstra	13
9	Resources added to a Blueprint in Apstra	14
10	Device Agent creation in Apstra for vJunos-switch virtual nodes	14
11	Device Agents created and acknowledged in Apstra	15
12	System IDs mapped in Apstra for the corresponding vJunos-switch devices	15
13	Connectivity Templates created in Apstra as part of VN creation	16

1 Building Virtual Fabrics with vJunos-switch and Containerlab

Aninda Chatterjee - 10/27/2023

How vJunos-switch is deployed as a VM, packaged within a container, on a bare metal server using the open source network emulation tool Containerlab.

We'll start with instructions on how to install Containerlab on Ubuntu 22.04 LTS, along with other dependencies that are needed to create the vJunos-switch docker image. Once done, we'll demonstrate how network topologies can be deployed using this vJunos-switch container in a simple, and declarative format using Containerlab.

Finally, we'll wrap this up by onboarding vJunos-switch devices in Juniper Apstra and building a simple Data Center fabric to show that these devices work seamlessly with Apstra.

1.1 What is Containerlab and vJunos-switch?

Containerlab is an open-source network emulation tool that offers users the ability to deploy virtual network operating systems as containers, connected together, to build small, medium or large scale network topologies for feature and functionality testing. More information can be found on their homepage - <https://containerlab.srlinux.dev/>. Containerlab supports both native containers from vendors (that provide it) and VM based qemu images, that can be packaged inside a docker container. The full list of supported images can be found in their user manual located here - <https://containerlab.dev/manual/kinds/>.

vJunos-switch is a new VM-based virtual software offering from Juniper Networks, which emulates the software functionality of Junos OS. This was integrated into Containerlab release 0.45.0, and is now a deployable node from this release onwards. All releases of vJunos-switch can be found here - <https://support.juniper.net/support/downloads/?p=vjunos>

1.2 Installing Containerlab and Building the vJunos-switch Container

Containerlab installation instructions can be found in their installation page (<https://containerlab.dev/install/>) or their quick start page (<https://containerlab.dev/quickstart/>). It can be installed as shown below, which installs the latest version of containerlab.

```

1 root@server:~# bash -c "$(curl -sL https://get.containerlab.dev)"
2 Downloading https://github.com/srl-labs/containerlab/releases/download/v0.47.0/containerlab_0
  .47.0_linux_amd64.deb
3 Preparing to install containerlab 0.47.0 from package
4 Selecting previously unselected package containerlab.
5 (Reading database ... 148116 files and directories currently installed.)
6 Preparing to unpack .../containerlab_0.47.0_linux_amd64.deb ...
7 Unpacking containerlab (0.47.0) ...
8 Setting up containerlab (0.47.0) ...
9
10
11      _ ( ) | | | |
12      / --- \ | | | | / --- \ | | | | / --- \ | | | |
13      ( ( | | | | | | | | ( ( | | | | | | | | ( ( | | | | | | | |
14      \ --- / | | | | \ --- / | | | | \ --- / | | | | \ --- /
15
16      version: 0.47.0
17      commit: d2a2ede1
18      date: 2023-10-22T19:33:30Z
19      source: https://github.com/srl-labs/containerlab
20      rel. notes: https://containerlab.dev/rn/0.47/

```

To build VM-based containers, Containerlab uses a fork of the *vrnetlab* project (<https://github.com/hellt/vrnetlab/>). In order to build such images, this repository must be cloned locally.

```
1 root@server:/home/anindac# git clone https://github.com/hellt/vrnetlab/
2 Cloning into 'vrnetlab'...
3 remote: Enumerating objects: 4133, done.
4 remote: Counting objects: 100% (1123/1123), done.
5 remote: Compressing objects: 100% (266/266), done.
6 remote: Total 4133 (delta 947), reused 927 (delta 857), pack-reused 3010
7 Receiving objects: 100% (4133/4133), 1.97 MiB | 8.31 MiB/s, done.
8 Resolving deltas: 100% (2516/2516), done.
```

Within this newly cloned repository, there will be several folders available for different vendor images, including a new *vjunosswitch* folder.

```
1 root@server:/home/anindac# cd vrnetlab/
2 root@server:/home/anindac/vrnetlab# ls -l
3 total 144
4 -rw-r--r-- 1 root root 94 Oct 23 05:36 CODE_OF_CONDUCT.md
5 -rw-r--r-- 1 root root 706 Oct 23 05:36 CONTRIBUTING.md
6 -rw-r--r-- 1 root root 1109 Oct 23 05:36 LICENSE
7 -rw-r--r-- 1 root root 342 Oct 23 05:36 Makefile
8 -rw-r--r-- 1 root root 4013 Oct 23 05:36 README.md
9 drwxr-xr-x 3 root root 4096 Oct 23 05:36 aoscx
10 drwxr-xr-x 2 root root 4096 Oct 23 05:36 ci-builder-image
11 drwxr-xr-x 2 root root 4096 Oct 23 05:36 common
12 drwxr-xr-x 3 root root 4096 Oct 23 05:36 config-engine-lite
13 drwxr-xr-x 3 root root 4096 Oct 23 05:36 csr
14 drwxr-xr-x 3 root root 4096 Oct 23 05:36 ftosv
15 -rw-r--r-- 1 root root 5210 Oct 23 05:36 git-lfs-repo.sh
16 -rw-r--r-- 1 root root 3158 Oct 23 05:36 makefile-install.include
17 -rw-r--r-- 1 root root 370 Oct 23 05:36 makefile-sanity.include
18 -rw-r--r-- 1 root root 1898 Oct 23 05:36 makefile.include
19 drwxr-xr-x 3 root root 4096 Oct 23 05:36 n9kv
20 drwxr-xr-x 3 root root 4096 Oct 23 05:36 nxos
21 drwxr-xr-x 3 root root 4096 Oct 23 05:36 ocnos
22 drwxr-xr-x 3 root root 4096 Oct 23 05:36 openwrt
23 drwxr-xr-x 3 root root 4096 Oct 23 05:36 pan
24 drwxr-xr-x 3 root root 4096 Oct 23 05:36 routers
25 drwxr-xr-x 3 root root 4096 Oct 23 05:36 sros
26 drwxr-xr-x 2 root root 4096 Oct 23 05:36 topology-machine
27 drwxr-xr-x 3 root root 4096 Oct 23 05:36 veos
28 drwxr-xr-x 3 root root 4096 Oct 23 05:36 vjunosswitch
29 drwxr-xr-x 3 root root 4096 Oct 23 05:36 vmx
30 drwxr-xr-x 3 root root 4096 Oct 23 05:36 vqfx
31 drwxr-xr-x 3 root root 4096 Oct 23 05:36 vr-bgp
32 drwxr-xr-x 2 root root 4096 Oct 23 05:36 vr-xcon
33 -rw-r--r-- 1 root root 1135 Oct 23 05:36 vrnetlab.sh
34 drwxr-xr-x 3 root root 4096 Oct 23 05:36 vrp
35 drwxr-xr-x 3 root root 4096 Oct 23 05:36 vsr1000
36 drwxr-xr-x 3 root root 4096 Oct 23 05:36 vsrx
37 drwxr-xr-x 3 root root 4096 Oct 23 05:36 xrv
38 drwxr-xr-x 3 root root 4096 Oct 23 05:36 xrv9k
```

This folder has a Makefile and a Dockerfile (within the *docker* folder), that facilitates building of the vJunos-switch docker image.

```
1 root@server:/home/anindac/vrnetlab# cd vjunosswitch/
2 root@server:/home/anindac/vrnetlab/vjunosswitch# ls -l
3 total 12
4 -rw-r--r-- 1 root root 346 Oct 23 05:36 Makefile
5 -rw-r--r-- 1 root root 513 Oct 23 05:36 README.md
6 drwxr-xr-x 2 root root 4096 Oct 23 05:36 docker
```

The user simply needs to copy the vJunos-switch images to this folder and initiate the **make** command. This will trigger the build, and once it is done, a new docker image for vJunos-switch should be available under **docker images**.

```
1 root@server:/home/anindac/vrnetlab/vjunosswitch# ls -l
2 total 3873680
3 -rw-r--r-- 1 root root 346 Oct 23 05:36 Makefile
4 -rw-r--r-- 1 root root 513 Oct 23 05:36 README.md
5 drwxr-xr-x 2 root root 4096 Oct 23 05:36 docker
```

```

6  -rwxr-xr-x 1 root root 3966631936 Oct 23 05:54 vjunos-switch-23.2R1.14.qcow2
7
8  root@dc-tme-bigtwin-02:/home/anindac/vrnetlab/vjunosswitch# make
9  for IMAGE in vjunos-switch-23.2R1.14.qcow2; do \
10 echo "Making $IMAGE"; \
11 make IMAGE=$IMAGE docker-build; \
12 done
13 Making vjunos-switch-23.2R1.14.qcow2
14 make[1]: Entering directory '/home/anindac/vrnetlab/vjunosswitch'
15 rm -f docker/*.qcow2* docker/*.tgz* docker/*.vmdk* docker/*.iso
16 Building docker image using vjunos-switch-23.2R1.14.qcow2 as vrnetlab/vr-vjunosswitch:23.2R1.14
17 cp ../common/* docker/
18 make IMAGE=$IMAGE docker-build-image-copy
19 make[2]: Entering directory '/home/anindac/vrnetlab/vjunosswitch'
20 cp vjunos-switch-23.2R1.14.qcow2* docker/
21 make[2]: Leaving directory '/home/anindac/vrnetlab/vjunosswitch'
22 (cd docker; docker build --build-arg http_proxy= --build-arg https_proxy= --build-arg IMAGE=
    vjunos-switch-23.2R1.14.qcow2 -t vrnetlab/vr-vjunosswitch:23.2R1.14 .)
23 DEPRECATED: The legacy builder is deprecated and will be removed in a future release.
24     Install the buildx component to build images with BuildKit:
25     https://docs.docker.com/go/buildx/
26
27 Sending build context to Docker daemon 3.967GB
28 Step 1/11 : FROM ubuntu:20.04
29 ----> d5447fc01ae6
30 Step 2/11 : ENV DEBIAN_FRONTEND=noninteractive
31 ----> Running in b8387480fabb
32 Removing intermediate container b8387480fabb
33 ----> 3027669a2e7b
34 Step 3/11 : RUN apt-get update -qy && apt-get upgrade -qy && apt-get install -y dosfstools
    bridge-utils iproute2 python3-ipyp socat qemu-kvm && rm -rf /var/lib/apt/lists/*
35 ----> Running in fa6033407315
36
37 <snip>
38
39 root@server:/home/anindac/vrnetlab/vjunosswitch# docker images
40 REPOSITORY TAG IMAGE ID CREATED SIZE
41 vrnetlab/vr-vjunosswitch 23.2R1.14 331fe9769a3d 33 seconds ago 4.4GB
42 sflow/sflow-rt latest cad189a26e3b 5 months ago 100MB
43 ubuntu 20.04 d5447fc01ae6 10 months ago 72.8MB

```

The image is now ready to use with Containerlab.

1.3 Building a Network Topology with vJunos-switch and Containerlab

Containerlab uses a simple, declarative format of writing a network topology in YAML. The topology is written in the form of network nodes that need to be deployed, the type and image of the node, and how these nodes are connected to each other. Containerlab provides a lot of control over the management of the network topology – this includes breaking out via the servers management interface, defining your own custom bridges and connecting the management interface of the network nodes to these bridges and so on. More details regarding network wiring can be found here - <https://containerlab.dev/manual/network/>.

For example, a network topology for a 3-stage Clos fabric is described below.

```

1  root@server:/home/anindac/# cat test-topology.yml
2  name: test-lab
3  mgmt:
4    bridge: virbr0
5    ipv4-subnet: 192.168.122.0/24
6  topology:
7    nodes:
8      spine1:
9        kind: vr-vjunosswitch
10       image: vrnetlab/vr-vjunosswitch:23.2R1.14
11       mgmt-ipv4: 192.168.122.101
12       startup-config: spine1.cfg
13      spine2:
14        kind: vr-vjunosswitch
15        image: vrnetlab/vr-vjunosswitch:23.2R1.14
16        mgmt-ipv4: 192.168.122.102
17        startup-config: spine2.cfg

```

```

18  leaf1:
19      kind: vr-vjunosswitch
20      image: vrnetlab/vr-vjunosswitch:23.2R1.14
21      mgmt-ipv4: 192.168.122.11
22      startup-config: leaf1.cfg
23  leaf2:
24      kind: vr-vjunosswitch
25      image: vrnetlab/vr-vjunosswitch:23.2R1.14
26      mgmt-ipv4: 192.168.122.12
27      startup-config: leaf2.cfg
28  h1:
29      kind: linux
30      image: aninchat/host:v1
31      mgmt-ipv4: 192.168.122.51
32      exec:
33          - sleep 5
34          - sysctl -w net.ipv4.icmp_echo_ignore_broadcasts=0
35          - ip route add 10.10.20.0/24 via 10.10.10.254
36      binds:
37          - hosts/h1_interfaces:/etc/network/interfaces
38  h2:
39      kind: linux
40      image: aninchat/host:v1
41      mgmt-ipv4: 192.168.122.52
42      exec:
43          - sleep 5
44          - sysctl -w net.ipv4.icmp_echo_ignore_broadcasts=0
45          - ip route add 10.10.10.0/24 via 10.10.20.254
46      binds:
47          - hosts/h2_interfaces:/etc/network/interfaces
48  links:
49      - endpoints: ["leaf1:eth1", "spine1:eth1"]
50      - endpoints: ["leaf1:eth2", "spine2:eth1"]
51      - endpoints: ["leaf2:eth1", "spine1:eth2"]
52      - endpoints: ["leaf2:eth2", "spine2:eth2"]
53      - endpoints: ["leaf1:eth3", "h1:eth1"]
54      - endpoints: ["leaf2:eth3", "h2:eth1"]

```

Let's understand what this topology is doing in some more detail:

- This topology describes four fabric nodes named spine1, spine2, leaf1 and leaf2.
- Each of these nodes are of type *vr:vjunosswitch*, which is the containerlab naming convention for this virtual OS.
- Each of these nodes are of image type *vrnetlab/vr-vjunosswitch:23.2R1.14*, which is the docker tag that was created for this image when it was built.
- In addition to the fabric nodes, two hosts are defined (named h1 and h2) and they are simply a Linux container of type *aninchat/host:v1*. This is a container that is hosted on docker hub within the repository 'aninchat/host' and has a tag of 'v1'.
- Each network node (fabric nodes and hosts) are connected to the internal KVM default bridge called *virbr0* using the *mgmt* hierarchy in the topology definition. Each node is also given a specific IP address from this subnet using the *mgmt-ipv4* key.
- The fabric nodes (spine1, spine2, leaf1 and leaf2) also have a startup configuration attached to them. This configuration is a base configuration to facilitate onboarding into Apstra, later in this document.
- Finally, using the *links* key, the interconnections between nodes is described.

Note: When any of these images are not found locally by docker, it will attempt to pull it from docker hub.

A Containerlab topology can be deployed using the **containerlab deploy** command, as shown below. This spins up the containers, creates the virtual wires to connect nodes as described in the topology, and executes any instructions tied to a node.

```

1  root@dc-tme-bigtwin-01:/home/anindac/jvd# containerlab deploy -t test-topology.yml
2  INFO[0000] Containerlab v0.47.0 started

```

```

3 INFO[0000] Parsing & checking topology file: test-topology.yml
4 INFO[0000] Creating docker network: Name="clab", IPv4Subnet="192.168.122.0/24", IPv6Subnet="",
  MTU='1'
5 INFO[0000] Creating lab directory: /home/anindac/jvd/clab-jvd
6 INFO[0000] Creating container: "h1"
7 INFO[0000] Creating container: "h2"
8 INFO[0000] Creating container: "spine2"
9 INFO[0000] Creating container: "leaf1"
10 INFO[0000] Creating container: "leaf2"
11 INFO[0000] Creating container: "spine1"
12 INFO[0001] Creating link: leaf2:eth1 <--> spine1:eth2
13 INFO[0001] Creating link: leaf1:eth1 <--> spine1:eth1
14 INFO[0001] Creating link: leaf2:eth3 <--> h2:eth1
15 INFO[0001] Creating link: leaf1:eth2 <--> spine2:eth1
16 INFO[0001] Creating link: leaf2:eth2 <--> spine2:eth2
17 INFO[0001] Creating link: leaf1:eth3 <--> h1:eth1
18 INFO[0002] Adding containerlab host entries to /etc/hosts file
19 INFO[0002] Adding ssh config for containerlab nodes
20 INFO[0013] Executed command "sleep 5" on the node "h2". stdout:
21 INFO[0013] Executed command "sysctl -w net.ipv4.icmp_echo_ignore_broadcasts=0" on the node "h2".
  stdout:
22 net.ipv4.icmp_echo_ignore_broadcasts = 0
23 INFO[0013] Executed command "ip route add 10.10.10.0/24 via 10.10.20.254" on the node "h2".
  stdout:
24 INFO[0013] Executed command "sleep 5" on the node "h1". stdout:
25 INFO[0013] Executed command "sysctl -w net.ipv4.icmp_echo_ignore_broadcasts=0" on the node "h1".
  stdout:
26 net.ipv4.icmp_echo_ignore_broadcasts = 0
27 INFO[0013] Executed command "ip route add 10.10.20.0/24 via 10.10.10.254" on the node "h1".
  stdout:
28 +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
29 | # | Name | Container ID | Image | Kind | State | IPv4 Address | IPv6 Address |
30 +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
31 | 1 | clab-jvd-h1 | 6e3e7b99b71e | aninchat/host:v1 | linux | running | 192.168.122.51/24 | N/A
32 | 2 | clab-jvd-h2 | 591bd1b9ed1d | aninchat/host:v1 | linux | running | 192.168.122.52/24 | N/A
33 | 3 | clab-jvd-leaf1 | 47bd8120ac60 | vrnetlab/vr-vjunosswitch:23.2R1.14 | vr-vjunosswitch |
  running | 192.168.122.11/24 | N/A |
34 | 4 | clab-jvd-leaf2 | a5c2e5e85315 | vrnetlab/vr-vjunosswitch:23.2R1.14 | vr-vjunosswitch |
  running | 192.168.122.12/24 | N/A |
35 | 5 | clab-jvd-spine1 | c492d76a54b5 | vrnetlab/vr-vjunosswitch:23.2R1.14 | vr-vjunosswitch |
  running | 192.168.122.101/24 | N/A |
36 | 6 | clab-jvd-spine2 | 2d6cc3b2a6ee | vrnetlab/vr-vjunosswitch:23.2R1.14 | vr-vjunosswitch |
  running | 192.168.122.102/24 | N/A |
37 +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+

```

The virbr0 bridge interface is created on the host itself, and has the following IP address assigned to it by default.

```

1 root@dc-tme-bigtwin-01:~# ifconfig virbr0
2 virbr0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
3     inet 192.168.122.1 netmask 255.255.255.0 broadcast 192.168.122.255
4     ether 52:54:00:d2:94:37 txqueuelen 1000 (Ethernet)
5     RX packets 1641995 bytes 559839364 (559.8 MB)
6     RX errors 0 dropped 0 overruns 0 frame 0
7     TX packets 1627134 bytes 244901034 (244.9 MB)
8     TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

Once the nodes are deployed via Containerlab, they should be reachable from the host server, via the virbr0 interface. In the example below, all nodes are reachable via the **ping** utility from the host server.

```

1 root@server:~# ping 192.168.122.101
2 PING 192.168.122.101 (192.168.122.101) 56(84) bytes of data.
3 64 bytes from 192.168.122.101: icmp_seq=1 ttl=64 time=0.081 ms
4 64 bytes from 192.168.122.101: icmp_seq=2 ttl=64 time=0.059 ms
5 ^C
6 --- 192.168.122.101 ping statistics ---
7 2 packets transmitted, 2 received, 0% packet loss, time 1009ms
8 rtt min/avg/max/mdev = 0.059/0.070/0.081/0.011 ms
9
10 root@server:~# ping 192.168.122.102
11 PING 192.168.122.102 (192.168.122.102) 56(84) bytes of data.
12 64 bytes from 192.168.122.102: icmp_seq=1 ttl=64 time=0.078 ms

```

```

13 64 bytes from 192.168.122.102: icmp_seq=2 ttl=64 time=0.034 ms
14 ^C
15 --- 192.168.122.102 ping statistics ---
16 2 packets transmitted, 2 received, 0% packet loss, time 1029ms
17 rtt min/avg/max/mdev = 0.034/0.056/0.078/0.022 ms
18
19 root@server:~# ping 192.168.122.11
20 PING 192.168.122.11 (192.168.122.11) 56(84) bytes of data.
21 64 bytes from 192.168.122.11: icmp_seq=1 ttl=64 time=0.111 ms
22 64 bytes from 192.168.122.11: icmp_seq=2 ttl=64 time=0.056 ms
23 ^C
24 --- 192.168.122.11 ping statistics ---
25 2 packets transmitted, 2 received, 0% packet loss, time 1032ms
26 rtt min/avg/max/mdev = 0.056/0.083/0.111/0.027 ms
27
28 root@server:~# ping 192.168.122.12
29 PING 192.168.122.12 (192.168.122.12) 56(84) bytes of data.
30 64 bytes from 192.168.122.12: icmp_seq=1 ttl=64 time=0.109 ms
31 64 bytes from 192.168.122.12: icmp_seq=2 ttl=64 time=0.054 ms
32 ^C
33 --- 192.168.122.12 ping statistics ---
34 2 packets transmitted, 2 received, 0% packet loss, time 1003ms
35 rtt min/avg/max/mdev = 0.054/0.081/0.109/0.027 ms
36
37 root@server:~# ping 192.168.122.51
38 PING 192.168.122.51 (192.168.122.51) 56(84) bytes of data.
39 64 bytes from 192.168.122.51: icmp_seq=1 ttl=64 time=0.115 ms
40 64 bytes from 192.168.122.51: icmp_seq=2 ttl=64 time=0.027 ms
41 ^C
42 --- 192.168.122.51 ping statistics ---
43 2 packets transmitted, 2 received, 0% packet loss, time 1017ms
44 rtt min/avg/max/mdev = 0.027/0.071/0.115/0.044 ms
45
46 root@server:~# ping 192.168.122.52
47 PING 192.168.122.52 (192.168.122.52) 56(84) bytes of data.
48 64 bytes from 192.168.122.52: icmp_seq=1 ttl=64 time=0.152 ms
49 64 bytes from 192.168.122.52: icmp_seq=2 ttl=64 time=0.057 ms
50 ^C
51 --- 192.168.122.52 ping statistics ---
52 2 packets transmitted, 2 received, 0% packet loss, time 1013ms
53 rtt min/avg/max/mdev = 0.057/0.104/0.152/0.047 ms

```

We can login to one of the devices (spine1, as an example) and confirm that it only has the base configuration that we had mapped to the startup-config key in the Containerlab topology.

```

1 root@server:~# ssh -l admin 192.168.122.101
2 (admin@192.168.122.101) Password:
3 Last login: Mon Oct 23 07:36:19 2023
4 --- JUNOS 23.2R1.14 Kernel 64-bit JNPR-12.1-20230613.7723847_buil
5 admin@spine1> show configuration
6 ## Last commit: 2023-10-23 07:36:12 UTC by root
7 version 23.2R1.14;
8 system {
9   host-name spine1;
10   root-authentication {
11     encrypted-password "$6$RAqyBEFe$ZeuTiT94b4nN7WCPs7eNBtnK6oX.
12       nqKEZLjdV7ckY0Nddh9wfYGPwKslqD2hUKtJFNle5Lt2LlD36FiK0Lo701"; ## SECRET-DATA
13   }
14   commit synchronize;
15   login {
16     user admin {
17       uid 2000;
18       class super-user;
19       authentication {
20         encrypted-password "$6$p5WEDsOX$sg.MSJal/
21           bHvaaoUOB6Vq3Htar9Y4HE4aFn8uQC85TlvQ0e8GqQEMhhOQWtmLjwB1ybSb4jX6AV1TTQAdl6b1";
22         ## SECRET-DATA
23       }
24     }
25   }
26   services {
27     ssh {
28       root-login allow;
29     }
30     netconf {
31       ssh;
32     }
33   }
34   management-instance;

```

```
32 }
33 interfaces {
34     fxp0 {
35         unit 0 {
36             family inet {
37                 address 10.0.0.15/24;
38                 address 192.168.122.101/24;
39             }
40         }
41     }
42 }
43 routing-instances {
44     mgmt_junos {
45         routing-options {
46             static {
47                 route 0.0.0.0/0 next-hop 10.0.0.2;
48             }
49         }
50     }
51 }
52 protocols {
53     lldp {
54         interface all;
55     }
56 }
```

1.4 Integrating a vJunos-switch Containerlab Based Network Topology with Juniper Apstra

The deployed Containerlab network topology is orchestrated and managed by Juniper Apstra, to demonstrate how virtual fabrics can be brought to life with Apstra as well for feature and functionality testing.

The Apstra version used for the purposes of this demonstration is Apstra 4.2.0. This version comes with a Device Profile for vJunos-switch. Before we begin, the topology that has been deployed with Containerlab is as shown below.




Figure 1: Network topology deployed via Containerlab

The Device Profile for vJunos-switch, in Apstra, is shown below.

The screenshot shows the Apstra 4.2.0 interface. The left sidebar contains navigation icons for Blueprints, Devices, Design, Resources, External Systems, Platform, and Favorites. The main content area displays the 'vJunos-switch' device profile configuration. The 'Summary' tab is active, showing a table with the following details:

Name	vJunos-switch
Device Profile Type	monolithic
Modular?	no
Slot count	0
Physical Device	no
Ports preview	

Below the summary, the 'Selector' section shows a table with the following details:

Manufacturer [®]	Juniper
Model [®]	vJunos-switch
OS family [®]	Junos
Version [®]	.*

The 'Capabilities' section is visible at the bottom but contains no data.

Figure 2: vJunos-switch Device Profile in Apstra 4.2.0

A Logical Device is created to match this Device Profile, with 96x1G ports as shown below.




Figure 3: vJunos-switch Logical Device created in Apstra

To tie the Device Profile and the Logical Device together, an Interface Map is created as shown below.




Figure 4: vJunos-switch Interface Map created in Apstra

We can start to create the actual fabric now. This process includes building a rack with vJunos-switch Logical Devices, building a 3-stage Clos Template based on this rack, and finally deploying a Blueprint using this Template.

First, a new Rack Type is created in Apstra to match our network topology.




Figure 5: vJunos-switch based Rack Type created in Apstra

This new Rack Type can now be used to create a Template, which defines the overall schema of the fabric, as shown below.




Figure 6: vJunos-switch based Template created in Apstra

Finally, once this Template is created, it is used as the only input to the Blueprint to create the initial fabric for this virtual Data Center.




Figure 7: vJunos-switch based Blueprint created in Apstra

The Blueprint itself needs several other basic resources. This includes:

- An IP pool for the point-to-point links between the leafs and the spines.
- An IP pool for loopbacks of all VTEPs in the fabric.
- An ASN pool for the leafs and the spines.

These can be created from the Resources tab in Apstra. For the sake of brevity, this is already done for the purposes of this demonstration.

At this point, the Blueprint is simply *staged*, and no configuration has been generated or deployed to the network devices.




Figure 8: Staged Blueprint in Apstra

We need to add all required sources into this Blueprint and map an appropriated Interface Map against every fabric node. Once this is done, you should see everything start to turn green from red.




Figure 9: Resources added to a Blueprint in Apstra

This can now be committed as the first revision of this Data Center deployment. So far, no devices have actually been mapped to this fabric. In order to do that, our vJunos-switch virtual devices must be onboarded into Apstra by creating offline Device Agents for them, as shown below.




Figure 10: Device Agent creation in Apstra for vJunos-switch virtual nodes

Once the agent is created, it can be acknowledged and these devices should now be available to be deployed in a Blueprint.




Figure 11: Device Agents created and acknowledged in Apstra

To deploy these in the Blueprint we just built, their respective System IDs need to be assigned in the Blueprint to the corresponding fabric nodes.




Figure 12: System IDs mapped in Apstra for the corresponding vJunos-switch devices

Now, all relevant configuration that was generated for the respective nodes will be pushed to the fabric devices. From both spines, we can confirm that BGP peering for the underlay and overlay is in an *Established* state.

```

1 admin@spine1> show bgp summary
2
3 Threading mode: BGP I/O
4 Default eBGP mode: advertise - accept, receive - accept
5 Groups: 2 Peers: 4 Down peers: 0
6 Table Tot Paths Act Paths Suppressed History Damp State Pending
7 inet.0

```

```

8      6 4 0 0 0 0
9  bgp.evpn.0
10     0 0 0 0 0 0
11  Peer AS InPkt OutPkt OutQ Flaps Last Up/Dwn State|#Active/Received/Accepted/Damped...
12  192.0.2.2 65423 4 2 0 0 57 Establ
13    bgp.evpn.0: 0/0/0/0
14  192.0.2.3 65424 5 3 0 0 58 Establ
15    bgp.evpn.0: 0/0/0/0
16  198.51.100.1 65423 6 5 0 0 1:05 Establ
17    inet.0: 2/3/3/0
18  198.51.100.3 65424 6 5 0 0 1:06 Establ
19    inet.0: 2/3/3/0
20
21  admin@spine2> show bgp summary
22
23  Threading mode: BGP I/O
24  Default eBGP mode: advertise - accept, receive - accept
25  Groups: 2 Peers: 4 Down peers: 0
26  Table Tot Paths Act Paths Suppressed History Damp State Pending
27  inet.0
28     6 4 0 0 0 0
29  bgp.evpn.0
30     0 0 0 0 0 0
31  Peer AS InPkt OutPkt OutQ Flaps Last Up/Dwn State|#Active/Received/Accepted/Damped...
32  192.0.2.2 65423 7 5 0 0 1:54 Establ
33    bgp.evpn.0: 0/0/0/0
34  192.0.2.3 65424 7 6 0 0 1:50 Establ
35    bgp.evpn.0: 0/0/0/0
36  198.51.100.5 65423 8 7 0 0 2:02 Establ
37    inet.0: 2/3/3/0
38  198.51.100.7 65424 8 8 0 0 1:58 Establ
39    inet.0: 2/3/3/0

```

This builds our core infrastructure for the virtual DC, however, the connectivity to the hosts is still not established. For this, we will create a Routing Zone and two Virtual Networks within this Routing Zone (corresponding to h1 and h2).

This also creates corresponding Connectivity Templates to provide Layer-2, untagged connectivity down to the hosts.




Figure 13: Connectivity Templates created in Apstra as part of VN creation

Once this is deployed, the two hosts can communicate with each other, as shown below.

```

1  root@h1:~# ping 10.10.20.2
2  PING 10.10.20.2 (10.10.20.2) 56(84) bytes of data:
3  64 bytes from 10.10.20.2: icmp_seq=1 ttl=62 time=2.33 ms

```

```
4 64 bytes from 10.10.20.2: icmp_seq=2 ttl=62 time=2.35 ms
5 64 bytes from 10.10.20.2: icmp_seq=3 ttl=62 time=2.58 ms
6 64 bytes from 10.10.20.2: icmp_seq=4 ttl=62 time=2.32 ms
7 ^C
8 --- 10.10.20.2 ping statistics ---
9 4 packets transmitted, 4 received, 0% packet loss, time 8ms
10 rtt min/avg/max/mdev = 2.321/2.396/2.580/0.112 ms
```

1.5 Summary

Through this post, we were able to demonstrate a working deployment of Juniper's new virtual offering, vJunos-switch, using Containerlab and integrate it with Juniper Apstra to deploy a Data Center fabric.

1.6 Useful links

- Juniper Networks software downloads - <https://www.juniper.net/us/en/dm/vjunos-labs.html>
- Containerlab home page - <https://containerlab.dev/>
- Juniper Apstra documentation - <https://www.juniper.net/documentation/product/us/en/apstra/>

1.7 Acknowledgments

- Ridha Hamidi – TME, CRDC
- Vivek V – TME, CRDC
- Nick Davey – Director, Product Management, CRDC
- Cathy Gadecki - Senior Director, Product Management, CRDC



DC – AI – Routing – Switching – Security – Wireless – Automation

© Copyright 2026 Hewlett Packard Enterprise Development LP. The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty.

Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Trademark acknowledgments, if needed. All third-party marks are the property of their respective owners.

HEWLETT PACKARD ENTERPRISE

HPE.com

HPE

