



DC – AI – Routing – Switching – Security – Wireless – Automation

vSRX on mini-PC with Linux/KVM

An HPE Juniper Networking Techpost

Karel Hendrych

2024-04-21

HPE



Contents

1	vSRX on mini-PC with Linux/KVM	3
1.1	Introduction	3
1.1.1	Pre-requisite for Reading and Applying this TechPost Practically	3
1.1.2	Motivating Factors for Using vSRX on Mini-PC	4
1.1.3	Everything Comes at a Price	4
1.2	Sample Mini-PC Setup	5
1.2.1	Time Proven Approach	5
1.3	Linux/KVM host	6
1.3.1	Install OS on LVM	6
1.3.2	Relevant software packages	6
1.3.3	Linux Bridge Networking Layout	7
1.3.4	Revealing CPU layout details	9
1.3.5	Tuning Host OS	10
1.3.6	Startup Script	11
1.3.7	Break-down of the setup script	11
1.3.8	Startup Script at Boot	11
1.3.9	Host OS Firewall	11
1.4	vSRX VM	13
1.4.1	vSRX VM Settings	13
1.4.2	Basic Performance Benchmark	15
1.4.3	Performance test conclusions	15
1.5	Lessons Learned	15
1.6	Appendix 1: WireGuard on host OS note	15
1.6.1	Server Configuration	17
1.6.2	Client Configuration	17
1.6.3	Sample setup script	17
1.7	Useful links	19
1.8	Glossary	19
1.9	Acknowledgements	20

List of Figures

1	Logo	3
2	SRXMON	4
3	Sample layout of host OS and vSRX networking bindings	8
4	Sample split of CPU resources between tasks	10
5	Results of performance test in given setup	15
6	Sample layout of vSRX WireGuard network topology	16

1 vSRX on mini-PC with Linux/KVM

Karel Hendrych - 04/21/2024

Using Juniper vSRX on hardware with constrained resources, typically a mini-PC serving as flexible Internet gateway. Those are lately very popular due to low footprint yet with capabilities making them suitable for running virtual machines.



Figure 1: Logo

1.1 Introduction

At very minimum, a mini-PC is tool for engineers to have x86 based SRX ([vSRX](#)) at home in small form factor, with low power consumption and no noise on fan-less units. For example, Juniper and J-partner Engineers can use it for the purpose of always on learning/demo system passing real life traffic patterns, especially interesting along with additional components like Juniper [Security Director Cloud](#), [Cloud ATP](#) and [Juniper Secure Connect](#). Complete professional Enterprise mini-setup can be formed together with the [EX switch](#) and [MIST APs/software stack](#).

Production use for small enterprise Firewall or distributed WAN solution is perfectly possible too, there are many suitable Whitebox appliances including environmentally hardened for industrial applications.

1.1.1 Pre-requisite for Reading and Applying this TechPost Practically

- Basic Linux/KVM skills, however not necessarily oriented on specific Linux distribution. Things are made partially in non-distribution specific way to avoid dependency on Linux flavor specifics by using generic tools like crond for startup script
- Knowledge of Junos, SRX, vSRX
- Modern x86 machine of choice, advisable are two or more Ethernet interfaces. Lately reasonably priced mini-PCs even with 2.5GE Ethernet and new generations of CPUs are available
- vSRX license, however every fresh install comes with 60 days trial, Advanced L7 trial licenses (including ATP) are provided on [Juniper web](#)

1.1.2 Motivating Factors for Using vSRX on Mini-PC

- Control plane speed: Junos commit times taking only second or two
- Decent boot times: packets can be passing in about two minutes after host power-on
- x86 PFE excels for crypto and heavy L7 security tasks: Anti-malware, IDP, SSL proxy, ...
- Operational factors like
 - easy to snapshot/rollback entire SRX “device” in a form of single file
 - out of band style of work with the SRX during maintenance
 - tcpdump of traffic on Linux bridge before and after entering the vSRX
- HW flexibility
 - in-build WiFi in some of the appliances can be turned into an AP/AP client
 - USBs can be used to connect additional NICs
- Possible to use any other virtual or container applications together with the vSRX VM
- Last but not least research and learning - it could have been too easy to use some out of the box virtualization solutions
- Meeting pre-requisites for planned follow-up article focusing on L7 security setup of vSRX on mini-PC along with details about SRXMON: great tool for Application, Network and System visibility created by Mark Barrett




Figure 2: SRXMON

1.1.3 Everything Comes at a Price

- Mini-PC with hypervisor and vSRX is not out of the box solution like SRX appliances
- Constrained HW resources of a mini-PC require initially attention
- Performance of software bringing can't match direct PCI resource access
- Unlike server HW, consumer grade mini-PC do not usually have

- out of band iLo/DRAC like management (on some could be substituted by vPro or similar)
 - SR-IOV support in chipset/BIOS
 - hardware watchdog known from servers
 - NICs may not have Open vSwitch/DPDK support for accelerated networking
- Mini-PCs are usually coming as barebones and need to be equipped with RAM and storage
 - Secure boot/TPM might not be present
 - Custom host OS install needs be maintained, although that might be just fine with unattended-upgrades like tooling

1.2 Sample Mini-PC Setup

Tested hardware specification:

- Intel i3-10110U CPU (2 Cores, 4 Threads)
- 2x 1GE Realtek NIC port
- 16GB DDR4 RAM and 1TB SATA SSD

Reasonable minimum on RAM/storage side would have been 8GB /128GB respectively. This specific setup has been running flawlessly for 3 years including a period of more than a year of host OS uptime. Greater SSD capacity is better for write intensive applications due to increased TBW: essentially write endurance in TB Written.

Note: at the time of writing 2.96 x 2.96 x 2.05” sized Intel N95 based mini-PC with 4x2.5GE Ethernet, 8GB RAM and 128GB SSD is priced for less than 300 USD.

1.2.1 Time Proven Approach

(Effectively spoiler, more details in next parts of the TechPost)

Using 2C/4T CPU Resources

- 1st core and HT sibling assigned to vSRX RE, NIC interrupt handling and Linux soft bridge processing
- 2nd core assigned to vSRX PFE core, both core and HT sibling excluded from OS scheduler
- Turbo CPU clock disabled couple of minutes after boot
 - Sustainable turbo clock is between 3.2-3.4GHz (at about 85deg/185F) when looked after thermal
 - Nominal clock is 2.1GHz (55deg/131F) when vSRX PFE thread is running, other threads idling
 - When running Turbo clock vSRX PFE load seem to be 20% lower with the same traffic pattern

Note: lately to achieve better vSRX performance on Hyperthreading capable CPU, it might be worth experimenting with vSRX PFE cores spawned across core(s) and HT siblings(s). Newer CPUs may not introduce jitter/latency for parallel performance of DPDK applications.

Two Ethernet ports

- 1st port dedicated to external bridge hosting vSRX untrust interface, usually Internet uplink
- 2nd port used for internal bridge passing 802.1Q VLAN tags to vSRX, avoids ge-0/0/x interfaces on vSRX for every network (possible scaling issue).

- VirtIO NICs bound to regular Linux bridge(s) are probably the best choice as there are no capabilities and resources for any accelerated networking option with consumer grade Realtek NIC. Might be different with higher end mini-PCs using NICs suitable for SR-IOV or Open vSwitch/DPDK.
- VLAN filtering where applicable, firewall to protect host
- Security considered: secured Linux bridge, disabled IPv6 except specific management

Host OS

- Linux has all the tooling and features to achieve above
- One of the choices is Debian Linux, known for stability and rich package offering
- In the specific setup Debian 11 had been in use for almost 3 years, now upgraded to Debian 12 with newer kernel and user space (Debian strength is usually very seamless upgrade)
- Generally, any modern Linux distribution should do, just matter of preference

1.3 Linux/KVM host

For Debian Linux, lightweight [net install](#) launched in text mode with only SSH server selected on software selection page makes the installation minimal. Disabling SSH password authentication (default for root) and use of only SSH keys is highly recommended.

1.3.1 Install OS on LVM

- Logical Volume Management offers partitioning layout flexibility and has block-layer snapshots if there is available space in Volume Group
- In case of boot issues /boot can be a separate partition (e.g. sda2), then 2nd partition as Physical Volume for LVM mapped to Volume Group. But lately booting off LVM is not an issue anymore.
- In the specific HW/SW setup, VG has 850GiB containing / root Logical Volume of 10GiB, /var LV sized 60GiB, swap LV 8GiB, the rest is spare for snapshots and future use
- Alternatively, BTRFS or similar filesystems can be considered (nothing wrong about using even on an LV) for brilliant `cp -reflink=always` to quickly copy files, compression, checksums, subvolumes, etc.)
- Ext4 filesystem is solid and reliable, without LVM is the easiest option, less flexible though

1.3.2 Relevant software packages

Hint: for non-GUI Debian system, `apt install --no-install-recommends [package(s)]` keeps the installed dependencies at minimum

- `qemu-system libvirt-daemon-system libvirt-clients bridge-utils`: KVM hypervisor base
- `haveged` - enhances Linux /dev/random entropy
- `thermald`: prevents overheating of Intel CPU by monitoring system and taking measures
- `smartmontools`: delivers `smartd` for monitoring SSD/HDD using S.M.A.R.T. attributes
- `chrony`: might be better option than `systemd NTP` capabilities (`timedatectl set-ntp 0`)
- `nftables`: manages, saves and restores IPv4/IPv6 Linux firewall for host OS protection (enabled upon installation by `systemctl enable nftables`)
- `postfix`: powerful yet relatively simple and lightweight MTA software, handy for sending out reports and notifications from host

- logwatch: might help to reveal issues based on log analytics
- apticron: monitors and reveals packages eligible for upgrade
- unattended-upgrades: automatically upgrades packages
- htop: ultimate alternative to top, monitors temperature and network throughput as well
- rfkill: software kill switch for WiFi/BT
- inxi: detailed system information (inxi -F -v7)
- intel-microcode: Intel CPU microcode, has AMD variant as well
- firmware-linux* - firmware for various hardware components (NIC, WiFi, BT, GPU, ...)
- iptraf-ng: real time traffic monitoring tool
- tcpdump: ultimate traffic capture tool
- wpa_supplicant: wireless client tooling
- hostapd: software for turning wireless card into an AP
- mate-desktop-environment, virt-manager: lightweight Linux desktop environment and desktop application for managing VMs
- xrdp, xorgxrdp: access to Linux GUI using RDP protocol, works well with above

Note: desktop environment brings lots of dependencies, inflates the installation and opens potentially new attack surfaces. Ideal is just to use SSH for management of a small setup.

List of core packages for copy/pasting:

```
1  qemu-system libvirt-daemon-system libvirt-clients bridge-utils haveged thermald smartmontools
    chrony nftables postfix logwatch apticron unattended-upgrades htop rfkill inxi firmware-
    linux iptraf-ng tcpdump
```

1.3.3 Linux Bridge Networking Layout

The schematics below describes more in detail the idea for networking layout described earlier, proposed vSRX zones are green colored.

- enp1s0 host interface connects to br-untrust bridge with no VLAN tagging
- enp2s0.14 is bridging VLAN 14 to br-mgmt for purposes of host OS and vSRX management (fxp0 + ge-0/0/1, latter one acting as for default GW for management segment), br-mgmt binds host OS management IP
- enp2s0 is passing VLAN tags to the vSRX VM (like VLAN 10 and VLAN 11) via bridge br-trunk, with the exception of VLAN 14
- With more than two NICs, better practice to use separate NIC for management




Figure 3: Sample layout of host OS and vSRX networking bindings

- Important: [Vincent Bernat's blog post about Linux bridge isolation](#)

On Debian networking configuration takes place in `/etc/network/interfaces` file. Sample setup of bridges for trust, untrust and management:

br-untrust

external Linux bridge connecting vSRX untrust interface `ge-0/0/0` with host physical interface `enp1s0`

```
1 auto enp1s0 br-untrust
2 iface enp1s0 inet manual
3 iface br-untrust inet manual
4 bridge_ports enp1s0
5 bridge_stp off
6 bridge_fd 0
7 bridge_maxwait 0
8 post-up echo 1 > /sys/class/net/br-untrust/bridge/vlan_filtering
9 post-up bridge vlan del dev br-untrust vid 1 self
```

- Above creates br-untrust bridge for vSRX untrust interface
- Bridges host `enp1s0` physical NIC
- Disables STP, no delay prior bringing up
- Isolates host from IP processing according to the blog-post above

br-trunk

internal Linux bridge for passing VLAN tagged traffic from switch to vSRX `ge-0/0/2` (no explicit filtering, switch taking care)

```
1 auto enp2s0 br-trunk
2 iface enp2s0 inet manual
3 iface br-trunk inet manual
4 bridge_ports enp2s0
5 bridge_stp off
6 bridge_fd 0
7 bridge_maxwait 0
8 post-up bridge vlan del dev br-trunk vid 1 self
```

- Creates br-trunk bridge for vSRX interface using tagging from within VM

- Bridges host enp2s0 physical NIC
- Disables STP, no delay prior bringing up
- Isolates host from possible IP processing
- No explicit VLAN filtering, OK if adjacent switch is under control

br-mgmt

br-mgmt Linux bridge is an exception to br-trunk VLAN tagging, instead of passing VLAN 14 via br-trunk, VLAN 14 is bridged to br-mgmt for management purposes

```
1 auto enp2s0.14 br-mgmt
2 iface enp2s0.14 inet manual
3 iface br-mgmt inet static
4 address 192.168.6.136/27
5 gateway 192.168.6.129
6 dns-nameservers 1.1.1.1 8.8.8.8
7 bridge_ports enp2s0.14
8 bridge_stp off
9 bridge_fd 0
10 bridge_maxwait 0
11 post-up echo 0 > /proc/sys/net/ipv6/conf/br-mgmt/disable_ipv6
```

- Creates br-mgmt bridge for hosting fpx0 and vSRX interface ge-0/0/1 (acting as default gateway for fpx0 and host OS)
- Assigns host IP address, default gateway (vSRX) and name servers
- Bridges host VLAN 14 from enp2s0 (removes VLAN tag)
- Disables STP, no delay prior bringing up
- Enables IPv6 (disabled globally otherwise, link local address without any other config)
- VLAN filtering not applicable here

1.3.4 Revealing CPU layout details

Based on `lscpu -e` command output below:

```
1  CPU NODE SOCKET CORE L1d:L1i:L2:L3 ONLINE MAXMHZ MINMHZ
2    0  0  0  0  0:0:0:0 yes 4100.0000 400.0000
3    1  0  0  1  1:1:1:0 yes 4100.0000 400.0000
4    2  0  0  0  0:0:0:0 yes 4100.0000 400.0000
5    3  0  0  1  1:1:1:0 yes 4100.0000 400.0000
```

CPU topology is following:

```
1 CPU core 0: CPU0 + HT sibling CPU2
2 CPU core 1: CPU1 + HT sibling CPU3
```

Proposed CPU resources split in detail:




Figure 4: Sample split of CPU resources between tasks

1.3.5 Tuning Host OS

To achieve above resource split, vSRX PFE CPU core and its HT sibling needs to be isolated from scheduling by tuning kernel parameters in `/etc/default/grub` (or equivalent), followed by `update-grub` to reflect changes for next boot:

```
1 GRUB_CMDLINE_LINUX_DEFAULT="isolcpus=1,3 nohz_full=1,3 irqaffinity=0,2 mitigations=off
  default_hugepagesz=1GB hugepagesz=1G hugepages=8 transparent_hugepage=never"
```

Break-down of kernel parameters:

- `isolcpus` - remove CPU 1 and 3 from OS balancing and scheduling
- `nohz_full`

removes as much as possible kernel noise from vSRX CPU (LOC interrupts)

not supported in every kernel image, indicated in `dmesg` output as: Housekeeping: `nohz` unsupported (kernel not built with `CONFIG_NO_HZ_FULL`) - `irqaffinity`: configures default IRQ affinity to CPU 0 and 2 (can be configured also run-time, in `/proc/irq/...`) - `mitigations` - disables all the performance impacting in-kernel mitigations for CPU vulnerabilities, - not that relevant for single tenant network appliance - as side channel attacks are possible only if someone compromised the host/vSRX and can run an application - for those concerned, keep enabled by removing the parameter along with up-to-date Linux kernel and Intel/AMD CPU microcode packages - `hugepages`: prevents memory fragmentation and increases performance for DPDK applications, 8x 1GB hugepages on host with 16GB RAM - Analyze output below or similar on given HW to identify devices causing most interrupts on CPU core/sibling before and after reboot (complete picture may appear over time, NICs are a good start). Prior adjustments, interrupts would have interfered with vSRX PFE core:

```
1 cat /proc/interrupts
2
3          CPU0 CPU1 CPU2 CPU3
4 126: 0 7231440 0 0 enp1s0
5 128: 0 0 0 5687501 enp2s0
```

- Classic ext4 filesystem SSD related settings in `/etc/fstab` (adjust accordingly for filesystem of choice to treat SSDs well)

```
1 /dev/mapper/vg01-root / ext4 errors=remount-ro,noatime,nodiratime,discard 0 1
```

- `discard` informs SSD device which block can be trimmed internally

- noatime, nodiratime relieve SSD from some more writes

1.3.6 Startup Script

```

1  #!/bin/bash
2  export PATH=/usr/local/sbin:/usr/local/bin:/sbin:/bin:/usr/sbin:/usr/bin
3  # disable IPv6, enable per interfaces in network settings
4  echo 1 > /proc/sys/net/ipv6/conf/all/disable_ipv6
5  echo 1 > /proc/sys/net/ipv6/conf/default/disable_ipv6
6  # block wifi and BT
7  rfkill block 0
8  rfkill block 1
9  # minimize swap use on SSD
10 echo 1 > /proc/sys/vm/swappiness
11 # sample pointer to WireGuard setup script
12 #/root/scripts/./wg0.sh
13 # disable turbo after 3 minutes
14 sleep 180; echo 1 > /sys/devices/system/cpu/intel_pstate/no_turbo

```

1.3.7 Break-down of the setup script

- Disables IPv6 globally (enabled explicitly per interface in network config)
- Disables BT/WiFi if needed
- Adjusts host OS swap behavior to reduce SSD wear out
- Proposal how to initiate WireGuard settings from Appendix 1
- Disables turbo clock 3 minutes after boot (leverages full steam during vSRX boot)

1.3.8 Startup Script at Boot

- Systematic way is to create systemd unit or any form of init script used by specific init system
- Similarly kernel run-time tunables are systematically set at boot in /etc/sysctl.conf
- Or use distribution specific rc.local like alternatives
- The easy way is to place the script (+ chmod 755 [file]) for example root /root/scripts with @reboot crontab event. That keeps the things in one place while the setup is being tuned and potentially helps to navigate. For example:

```

1  crontab --e
2
3  @reboot /root/scripts/startup.sh

```

- When doing non-systematic things, message in /etc/motd should remind about that during every login

1.3.9 Host OS Firewall

Good practice is to engage Linux packet filtering to protect host OS management interface. Linux firewall can be used for both stateless and stateful filtering on bridge if certain traffic patterns would need to be dropped prior reaching vSRX. Sample nftables settings for host OS management protection stored in /etc/nftables.conf:

```

1  #!/usr/sbin/nft -f
2
3  flush ruleset
4
5  table inet filter {

```

```

6
7     set MGMT {
8         #permits SSH
9         type ipv4_addr
10        counter
11        flags interval
12        elements = {
13            192.168.6.128/27,
14            192.168.6.2,
15        }
16    }
17
18    set MGMT6 {
19        #permits IPv6 SSH
20        type ipv6_addr
21        counter
22        flags interval
23        elements = {
24            2a02:cbe0:dead::/64,
25        }
26    }
27
28    chain INPUT {
29        type filter hook input priority filter; policy accept;
30        ct state invalid counter drop
31        ip protocol icmp counter limit rate 10/second accept
32        ip protocol icmp counter drop
33        ip6 nexthdr icmpv6 counter limit rate 10/second accept
34        ip6 nexthdr icmpv6 counter drop
35        ct state related,established counter accept
36        ip saddr @MGMT ct state new tcp dport 22 counter log flags all prefix "MGMT " accept
37        ip6 saddr @MGMT6 ct state new tcp dport 22 counter log flags all prefix "MGMT "
38        accept
39        ip saddr 192.168.6.0/24 udp dport 123 counter accept
40        iifname "lo" counter accept
41        counter limit rate 10/second reject with icmpx type admin-prohibited
42        counter drop
43    }
44
45    chain FORWARD {
46        type filter hook forward priority filter; policy drop;
47        counter drop
48    }
49
50    chain OUTPUT {
51        type filter hook output priority filter; policy accept;
52        counter accept
53    }

```

Break-down of above nftables settings:

- Both ICMPv4 and v6 are throttled to 10 packets/second,
- Established connections are permitted,
- SSH can be established from specific IPv4 and IPv6 addresses stored in sets only, those can be expanded runtime using nft add/delete element syntax
- NTP permitted from specific /24 prefix
- All outbound traffic is permitted (sub-par)

Validation, application and listing of changes done in /etc/nftables.conf:

```

1 nft -c -f /etc/nftables.conf
2 systemctl reload nftables
3 nft list ruleset

```

1.4 vSRX VM

1.4.1 vSRX VM Settings

- Sample QEMU/KVM VM config for vSRX3 in text box below
- Suitable for CTRL+A, CTRL+C and placement to favorite text editor, then to VM configuration file /etc/libvirt/qemu/vsrx.xml
- Important part is section with CPU pinning, this effectively pins 2nd vCPU to host CPU1 (numbering starts with 0)

```
1 <cpupin>
2   <vcpupin vcpu='1' cpuset='1' />
3 </cpupin>
```

- VM memory size must be aligned ($\leq$) to configured hugepages (extra OS settings part)
- vNICs are configured and mapped as follows:

```
1 1st fxp0 (br-mgmt)
2 2nd ge-0/0/0 (br-untrust)
3 3rd ge-0/0/1 (br-mgmt)
4 4th ge-0/0/2 (br-trunk)
```

- vHDD image placed in /var/lib/libvirt/images/vsrx.qcow2
- Minimal amount of vRAM (4GiB) and vCPUs (2)

```
1 <domain type='kvm'>
2   <name>vsrx</name>
3   <memory unit='KiB'>4194304</memory>
4   <currentMemory unit='KiB'>4194304</currentMemory>
5   <memoryBacking>
6     <hugepages>
7       <page size='1048576' unit='KiB' />
8     </hugepages>
9     <nosharepages />
10    <locked />
11    <allocation mode='immediate' />
12  </memoryBacking>
13  <vcpu placement='static'>2</vcpu>
14  <cpupin>
15    <vcpupin vcpu='1' cpuset='1' />
16  </cpupin>
17  <os>
18    <type arch='x86_64' machine='pc-i440fx-3.1'>hvm</type>
19    <boot dev='hd' />
20  </os>
21  <features>
22    <acpi />
23    <apic />
24  </features>
25  <cpu mode='host-model' check='partial' />
26  <clock offset='utc'>
27    <timer name='rtc' tickpolicy='catchup' />
28    <timer name='pit' tickpolicy='delay' />
29    <timer name='hpet' present='no' />
30  </clock>
31  <on_poweroff>destroy</on_poweroff>
32  <on_reboot>restart</on_reboot>
33  <on_crash>restart</on_crash>
34  <pm>
35    <suspend-to-mem enabled='no' />
36    <suspend-to-disk enabled='no' />
37  </pm>
38  <devices>
39    <emulator>/usr/bin/kvm</emulator>
40    <disk type='file' device='disk'>
41      <driver name='qemu' type='qcow2' />
42      <source file='/var/lib/libvirt/images/vsrx.qcow2' />
43      <target dev='vda' bus='virtio' />
44      <address type='pci' domain='0x0000' bus='0x00' slot='0x06' function='0x0' />
```

```

45 </disk>
46 <controller type='usb' index='0' model='ich9-ehci1'>
47   <address type='pci' domain='0x0000' bus='0x00' slot='0x05' function='0x7' />
48 </controller>
49 <controller type='usb' index='0' model='ich9-uhci1'>
50   <master startport='0' />
51   <address type='pci' domain='0x0000' bus='0x00' slot='0x05' function='0x0' multifunction='on' />
52 </controller>
53 <controller type='usb' index='0' model='ich9-uhci2'>
54   <master startport='2' />
55   <address type='pci' domain='0x0000' bus='0x00' slot='0x05' function='0x1' />
56 </controller>
57 <controller type='usb' index='0' model='ich9-uhci3'>
58   <master startport='4' />
59   <address type='pci' domain='0x0000' bus='0x00' slot='0x05' function='0x2' />
60 </controller>
61 <controller type='pci' index='0' model='pci-root' />
62 <controller type='virtio-serial' index='0'>
63   <address type='pci' domain='0x0000' bus='0x00' slot='0x04' function='0x0' />
64 </controller>
65 <interface type='bridge'>
66   <source bridge='br-mgmt' />
67   <model type='virtio' />
68   <address type='pci' domain='0x0000' bus='0x00' slot='0x03' function='0x0' />
69 </interface>
70 <interface type='bridge'>
71   <source bridge='br-untrust' />
72   <model type='virtio' />
73   <address type='pci' domain='0x0000' bus='0x00' slot='0x09' function='0x0' />
74 </interface>
75 <interface type='bridge'>
76   <source bridge='br-mgmt' />
77   <model type='virtio' />
78   <address type='pci' domain='0x0000' bus='0x00' slot='0x0a' function='0x0' />
79 </interface>
80 <interface type='bridge'>
81   <source bridge='br-trunk' />
82   <model type='virtio' />
83   <address type='pci' domain='0x0000' bus='0x00' slot='0x0b' function='0x0' />
84 </interface>
85 <serial type='pty'>
86   <target type='isa-serial' port='0'>
87     <model name='isa-serial' />
88   </target>
89 </serial>
90 <console type='pty'>
91   <target type='serial' port='0' />
92 </console>
93 <input type='mouse' bus='ps2' />
94 <input type='keyboard' bus='ps2' />
95 <graphics type='vnc' port='-1' autoport='yes'>
96   <listen type='address' />
97 </graphics>
98 <video>
99   <model type='qxl' ram='65536' vram='65536' vgamem='16384' heads='1' primary='yes' />
100   <address type='pci' domain='0x0000' bus='0x00' slot='0x02' function='0x0' />
101 </video>
102 <memballoon model='virtio'>
103   <address type='pci' domain='0x0000' bus='0x00' slot='0x07' function='0x0' />
104 </memballoon>
105 <rng model='virtio'>
106   <backend model='random'>/dev/random</backend>
107   <address type='pci' domain='0x0000' bus='0x00' slot='0x08' function='0x0' />
108 </rng>
109 </devices>
110 </domain>

```

Finally define VM from XML, set automatic start with host and start with serial console attached.

```

1 virsh define vsrx.xml
2 virsh autostart vsrx
3 virsh start vsrx --console

```

Note: use virsh edit vsrx after vSRX VM has been defined, otherwise changes won't be reflected.

1.4.2 Basic Performance Benchmark

- Simple L4 FW+NAT test using two nuttcp TCP streams across two physical NICs (theoretical limit is 1Gbps full-duplex)
- 4 test cases:
 - w/wo effective isolcpus, irqaffinity
 - w/wo turbo CPU clock (nominal 2.1GHz, turbo in range between 3.2-3.4GHz)
 - PPS and BPS figures are average over 60s run, it's summary of ingress on both enp1s0 and enp2s0

IRQ affinity	Turbo clock	Ingress [Mbps]	Ingress [kPPS]	vSRX PFE load [%]	CPU Temp [°C]
Y	Y	1682	166	22	80
Y	N	1590	156	28	69
N	Y	1691	166	45	81
N	N	1680	165	61	67

Figure 5: Results of performance test in given setup

1.4.3 Performance test conclusions

- Noticeable difference is lower vSRX PFE load when CPU1 is dedicated purely to vSRX
- In optimized cases still lots of headroom on vSRX PFE core, e.g., for heavy L7 services
- Turbo clock not worth it in the given setup, probably only more cores would overcome software bridge processing bottleneck
- Performance should do the trick on 1Gbps Internet link
- Professional equipment would be needed for proper testing, like to see where packet loss starts occurring

1.5 Lessons Learned

- Having alternative access (VGA, USB-serial tty, WiFi,...) is a good idea when tuning host OS networking and firewall
- Quite handy are Junos config interface descriptions containing name of adjacent Linux bridge and bound physical NIC / VLAN
- Tools for E2E stability and performance test if the stack is stable and well behaved: iperf, nuttcp
- When using nuttcp like tools, taskset -c [CPU] [command] prevents jitter, otherwise OS scheduler re-balances continuously the test tool task to another CPU
- Bright WiFi status LED diodes can be potentially disabled using kernel module parameters
- Devices with slow commit are after vSRX experience challenging to deal with

1.6 Appendix 1: WireGuard on host OS note

- Following is an idea to complement SRX IPSEC capabilities for both site to site and remote access VPNs

- Main purpose is to get a practical grasp of Linux network namespaces and potentially have traffic passing through the vSRX all the time
- [WireGuard](#) is simple modern just works like VPN technology with in-build roaming capability, available out of the box starting with Ubuntu 20.04 and Debian 11
- Linux, Android and Windows clients works nicely in existing setup, recently also BSD flavors got their ports too
- Easy way to integrate with vSRX would be next-hop VM running WireGuard, however that's not needed when host OS supports WireGuard directly
- Proposed lightweight like namespace approach has a flip side of
- WireGuard being possibly processed also on vSRX PFE core sibling (kernel threads, needs investigation), but practically no issue has been seen
- Potential attack surface against host OS (great reputation so far though)
- WireGuard tunnel termination IPv4 must be reachable via vSRX (NAT, routing)
- Works with IPv6 addresses, both inner and outer
- Decapsulated traffic is controlled on vSRX between ravpn zone and other zones, encapsulated needs to be controlled towards ravpn zone




Figure 6: Sample layout of vSRX WireGuard network topology

WireGuard namespace setup description:

- vSRX zone ravpn, ge-0/0/3.0 (configured as another exception to tagging like br-mgmt)
- Host OS bridge br-ravpn connecting namespace and vSRX
- WireGuard namespace connected using veth pair to br-ravpn
- veth is like Junos lt- interface, works in a pair
- Namespace having default gateway via vSRX 100.64.0.1
- vSRX routing 192.168.1.0/27 (client tunnel address pool) via namespace IP (+IPv6 if in use)
- wg0 tunnel adapter has the 1st IP from the tunnel prefix.
- Assumes DNAT to tunnel termination port from outside towards 100.64.0.10

1.6.1 Server Configuration

- Following is Sample WireGuard config in “remote access” headend style, including IPv6
- Placed in /etc/wireguard/wg0.conf on host OS, effectively used in wg0 namespace

```

1 [Interface]
2 Address = 192.168.6.1/27, 2a02:cbe0:dead:f::1/64
3 ListenPort = 51820
4 PostUp = wg set %i private-key /etc/wireguard/private.key
5
6 [Peer]
7 AllowedIPs = 192.168.6.2, 2a02:cbe0:dead:f::2
8 PublicKey = [remote public key]
9 PresharedKey = [pre-shared key]
10
11 [Peer]
12 AllowedIPs = 192.168.6.3, 2a02:cbe0:dead:f::3
13 PublicKey = [remote public key]
14 PresharedKey = [pre-shared key]

```

- Addresses WireGuard tunnel interface with IPv4 and IPv6
- Assigns listening port
- Loads private key from a file (not revealed in configuration)
- Two peers, peer’s traffic can’t be sourced from any other IPs than allowed
- Pre-shared key is optional however best practice for post-quantum security
- At the same time route traffic towards peers

1.6.2 Client Configuration

- Sample WireGuard config, “remote access” client style matching one of the clients defined in server side configuration

```

1 [Interface]
2 Address = 192.168.6.2, 2a02:cbe0:dead:f::3
3 PrivateKey = [private key]
4 DNS = [DNS1, DNS2, search-prefix1.tld, search-prefix2.tld]
5
6 [Peer]
7 EndPoint = [hostname]:51820
8 PublicKey = [remote public key]
9 PresharedKey = [pre-shared key]
10 AllowedIPs = 0.0.0.0/0, ::/0
11 PersistentKeepalive = 25

```

- Defines local tunnel interface IPv4 and IPv6 IPs
- Private key
- Pre-shared key if enabled on server side as best practice
- DNS servers and prefix search list
- Remote peer IP/hostname and port
- Tunnels all IPv4 and IPv6 traffic
- Optional keepalives, preserves the same UDP session

1.6.3 Sample setup script

- Namespace and WireGuard sample start/restart script, includes IPv6
- Would be called probably from startup.sh (by uncommenting specific line in the general startup script if placed in /root/scripts/wg0.sh) or some distribution specific init style way

```

1  #/bin/bash
2  IFS=$'\n'
3
4  function ns () {
5      ip netns exec wg0 $@
6  }
7
8  test -e /run/netns/wg0 && ip netns del wg0 && sleep 0.5
9  ip netns add wg0
10 ip link add veth_wg0_host type veth peer veth_wg0_ns
11 ip link set veth_wg0_host master br-wgvpn up
12 ip link set dev veth_wg0_ns netns wg0
13
14 # ip netns exec wg0 nft list ruleset
15 ns nft add table filter
16 ns nft add chain filter input { type filter hook input priority 0\; policy drop \; }
17 ns nft add chain filter forward { type filter hook forward priority 0\; policy drop \; }
18 ns nft add chain filter output { type filter hook output priority 0\; policy accept \; }
19
20 ns nft add rule filter input udp dport 51820 counter accept
21 ns nft add rule filter input meta l4proto icmp icmp type echo-request counter accept
22 ns nft add rule filter input meta l4proto ipv6-icmp counter accept
23 ns nft add rule filter input counter reject
24
25 ns nft add rule filter output counter accept
26
27 ns nft add rule filter forward iifname "wg0" oifname "wg0" ct state related,established counter
   accept
28 ns nft add rule filter forward iifname "wg0" oifname "wg0" ip saddr 192.168.6.2 counter accept
29 ns nft add rule filter forward iifname "wg0" oifname "wg0" counter reject
30 ns nft add rule filter forward counter accept
31
32 ns ip link set veth_wg0_ns up
33
34 ns ip addr add 100.64.0.10/24 dev veth_wg0_ns
35 ns ip rou add default via 100.64.0.1
36 ns sysctl -q -w net.ipv4.ip_forward=1
37 ns ping -c1 -w1 100.64.0.1 &>/dev/null
38
39 ns ip -6 addr add 2a02:cbe0:dead:e::10/64 dev veth_wg0_ns
40 ns ip -6 route add ::/0 via 2a02:cbe0:dead:e::1
41 ns sysctl -q -w net.ipv6.conf.all.forwarding=1
42
43 #echo module wireguard +p > /sys/kernel/debug/dynamic_debug/control
44 ns wg-quick up wg0 &>/dev/null

```

Break-down of the init script:

- Removes wg0 namespace if present
- Creates wg0 namespace, veth pair
- veth_wg0_host is host side veth, veth_wg0_ns becomes wg0 namespace side
- veth_wg0_host bound to br-ravpn bridge
- nftables in namespace allowing on input only ICMP and WireGuard port
- nftables rejecting VPN endpoints talking each other with exception of 192.168.6.2
- Enables forwarding for both IPv4 and IPv6 in namespace
- Adds IPv4/IPv6 addressing and routing in namespace
- vSRX ARP record refresh by running ping
- Optional WireGuard kernel module debug (deactivated)
- Initializes WireGuard using /etc/wireguard/wg0.conf

To execute commands in namespace, for example to retrieve WireGuard status:

```
1  ip netns exec wg0 wg
```

1.7 Useful links

- <https://www.debian.org/distrib/>
- <https://wiki.archlinux.org/index.php/LVM>
- <https://wiki.debian.org/SSDOptimization>
- <https://manpages.debian.org/buster/ifupdown/interfaces.5.en.html>
- <https://vincent.bernat.ch/en/blog/2017-linux-bridge-isolation>
- <https://manpages.debian.org/testing/bridge-utils/bridge-utils-interfaces.5.en.html>
- <https://www.kernel.org/doc/Documentation/admin-guide/kernel-parameters.txt>
- https://www.suse.com/c/cpu-isolation-nohz_full-part-3/
- <https://wiki.debian.org/KVM>
- https://wiki.nftables.org/wiki-nftables/index.php/Main_Page
- <https://libvirt.org/formatdomain.html>
- <https://www.juniper.net/documentation/us/en/software/vsrx/vsrx-kvm/topics/concept/security-vsrx-kvm-understanding.html>
- <https://www.wireguard.com/>
- <https://btrfs.readthedocs.io/en/latest/>

1.8 Glossary

- AP Access Point
- BPS Bits Per Second
- DPDK Data Plane Development Kit
- DRAC Dell Remote Access Controller
- HT Hyper Threading
- IDP Intrusion Detection Prevention
- iLo Integrated Lights Out
- KVM Kernel Virtual Machine
- LV Logical Volume
- NIC Network Interface Card
- OVS Open vSwitch
- PFE Packet Forwarding Engine
- PPS Packets Per Second
- RDP Remote Desktop Protocol
- RE Routing Engine
- SR-IOV Single Root Input Output Virtualization
- SSD Solid State Disk
- STP Spanning Tree Protocol
- TBW Terabytes Written
- TPM Trusted Platform Module
- VG Volume Group

1.9 Acknowledgements

Subbiah Kandasamy for looking at the x86 aspects in the initial PPT form of this document. Then all the people who provided feedback - Matthijs Nagel, Mark Barrett, Steven Jacques, Kelly Brazil and others from team JNPR! Of course, no way to make things happen without all the brilliant Open-Source Software. Finally, vSRX/SRX dev and product teams for delivering the Swiss army knife for security and networking.



DC – AI – Routing – Switching – Security – Wireless – Automation

© Copyright 2026 Hewlett Packard Enterprise Development LP. The information contained herein is subject to change without notice. The only warranties for Hewlett Packard Enterprise products and services are set forth in the express warranty statements accompanying such products and services. Nothing herein should be construed as constituting an additional warranty.

Hewlett Packard Enterprise shall not be liable for technical or editorial errors or omissions contained herein.

Trademark acknowledgments, if needed. All third-party marks are the property of their respective owners.

HEWLETT PACKARD ENTERPRISE

HPE.com

